Correctness and Performance of Apache Spark SQL

Spark + AI Summit, London

October 4, 2018

databricks

About us

BOGDAN GHIT

Databricks, Software Engineer

- SQL performance optimizations

IBM T.J. Watson, Research Intern

- Bid advisor for cloud spot markets

Delft University of Technology, PhD in Computer Science

- Resource management in datacenters
- Performance of Spark, Hadoop

NICOLAS POGGI

Databricks, Performance Engineer

- Spark benchmarking

Barcelona Supercomputing - Microsoft Research Centre

- Lead researcher ALOJA project
- New architectures for Big Data

BarcelonaTech (UPC), PhD in Computer Architecture

- Autonomic resource manager for the cloud
- Web customer modeling

Databricks ecosystem

databricks

Databricks runtime (DBR) releases

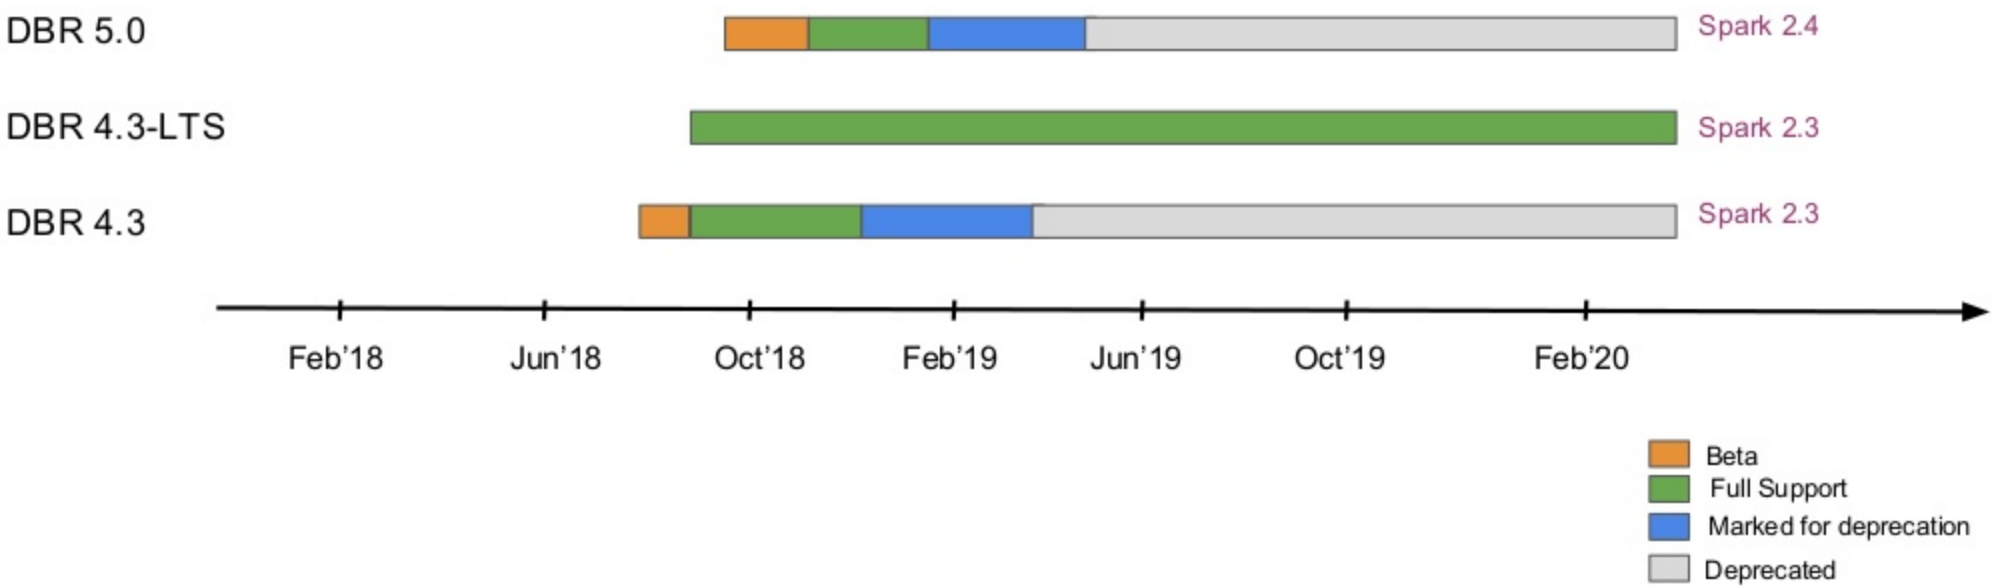

Our goal is to make releases **automatic** and **frequent**

Apache Spark contributions

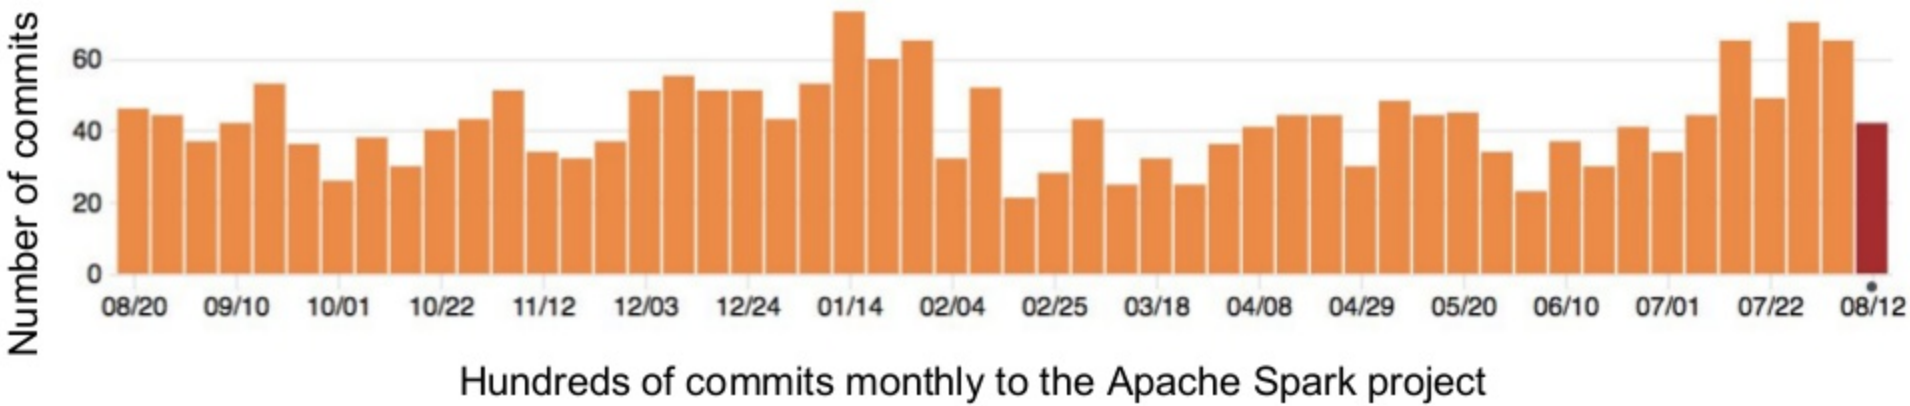

At this pace of development, **mistakes** are bound to happen

Where do these contributions go?

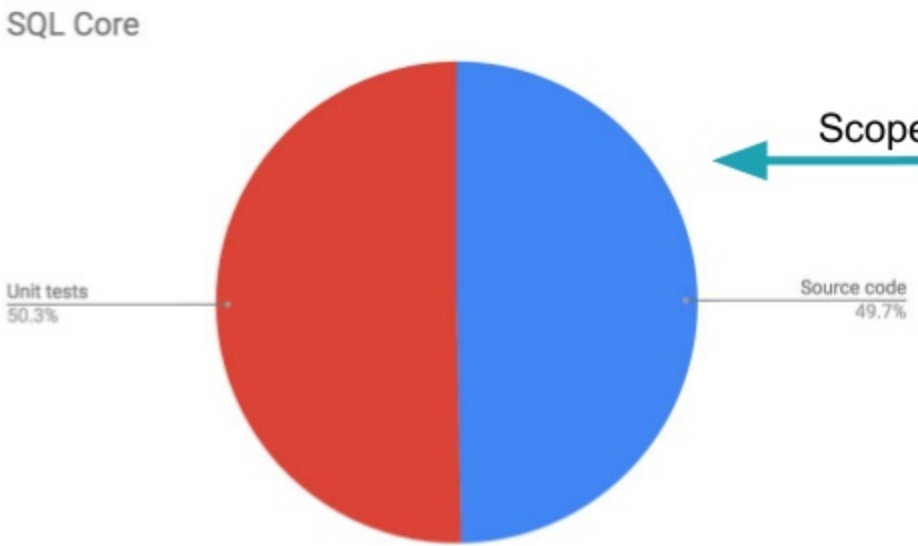

Scope of the testing

Over 200 built-in functions

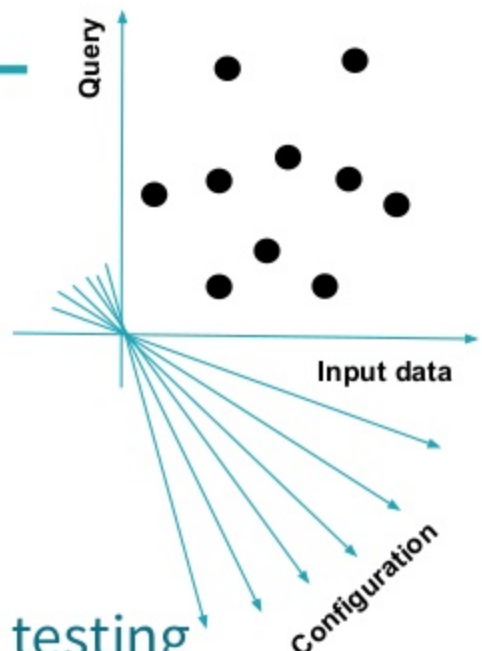

Developers put a significant **engineering effort** in testing

Yet another brick in the wall

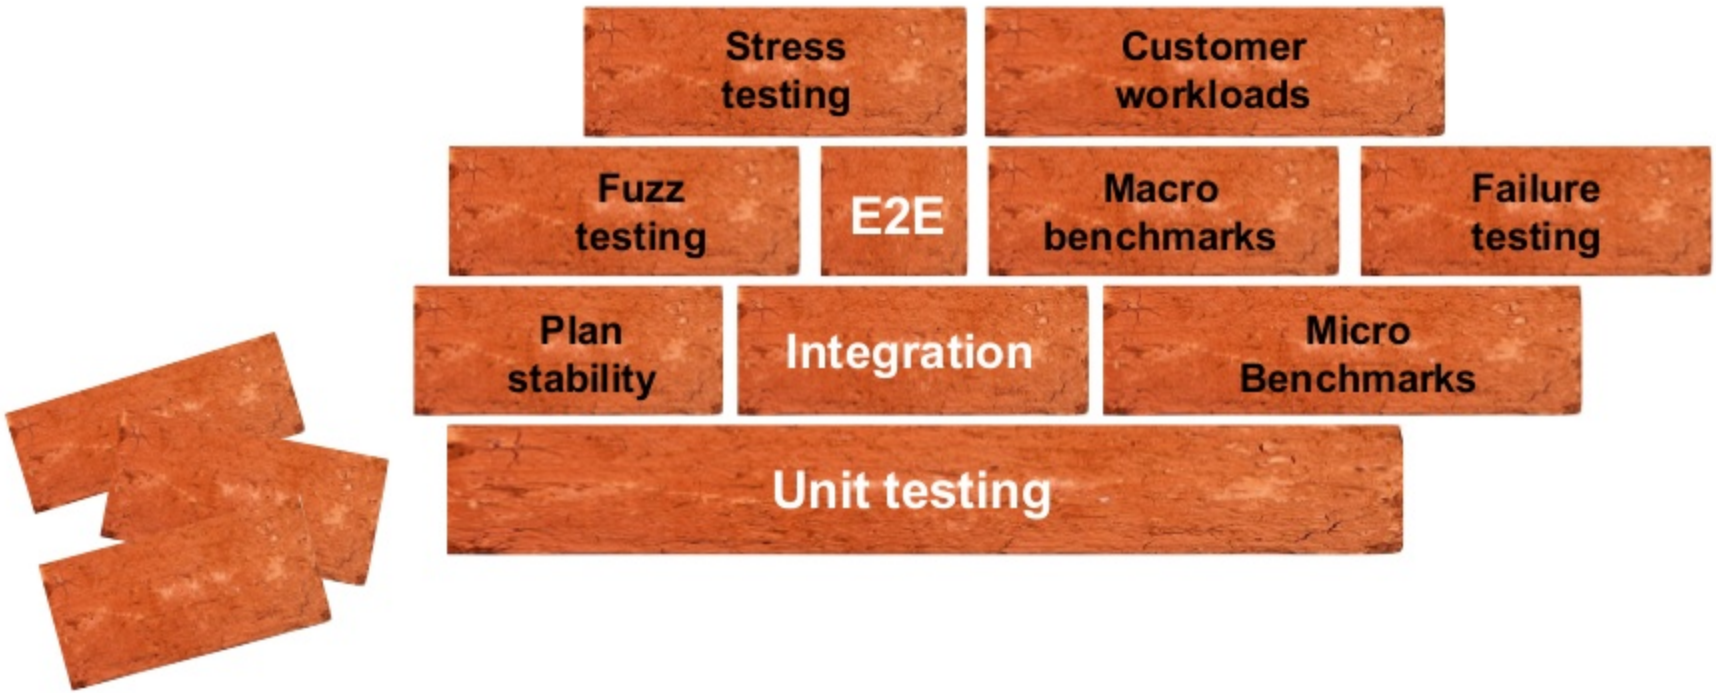

Unit testing ***is not enough*** to guarantee correctness and performance

Continuous Integration pipeline

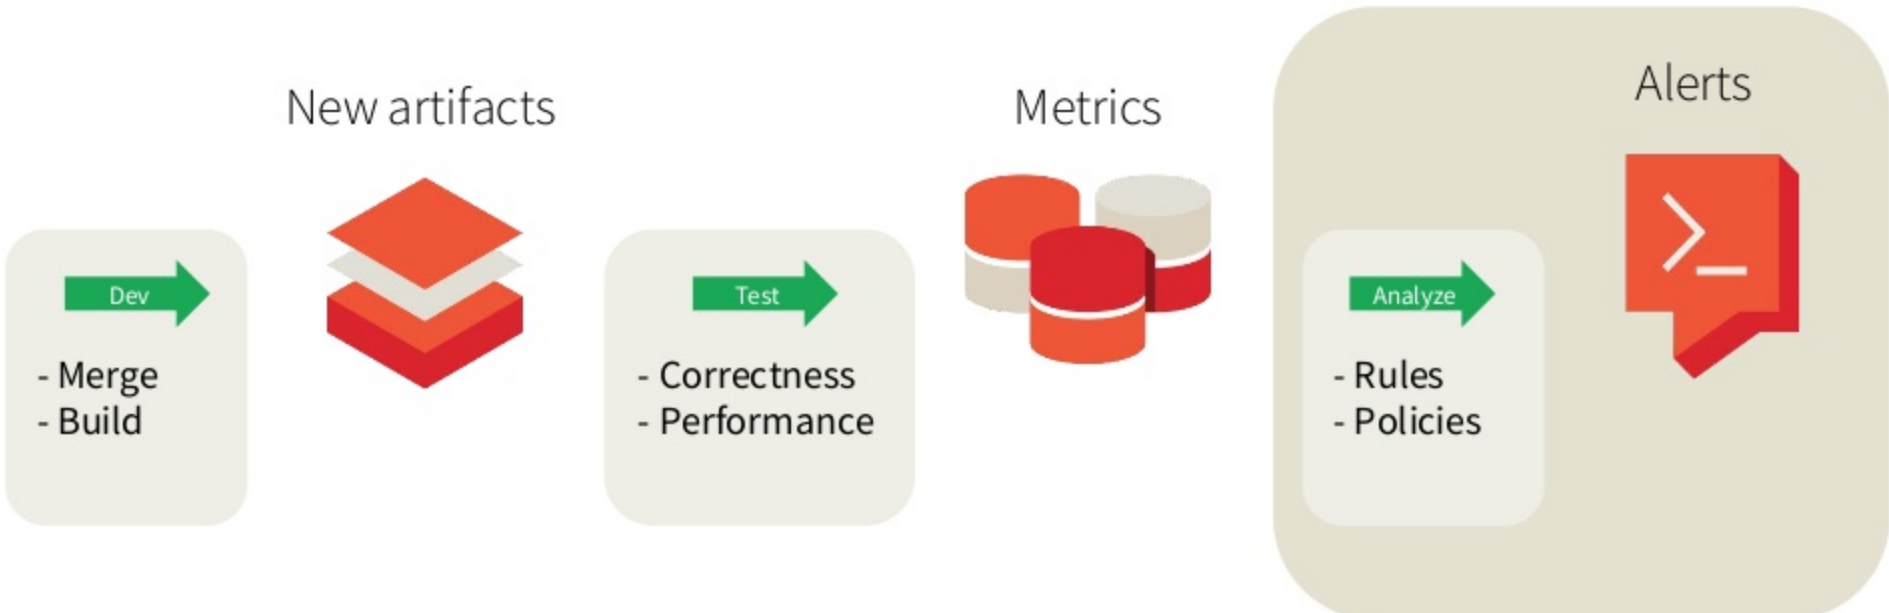

Classification and alerting

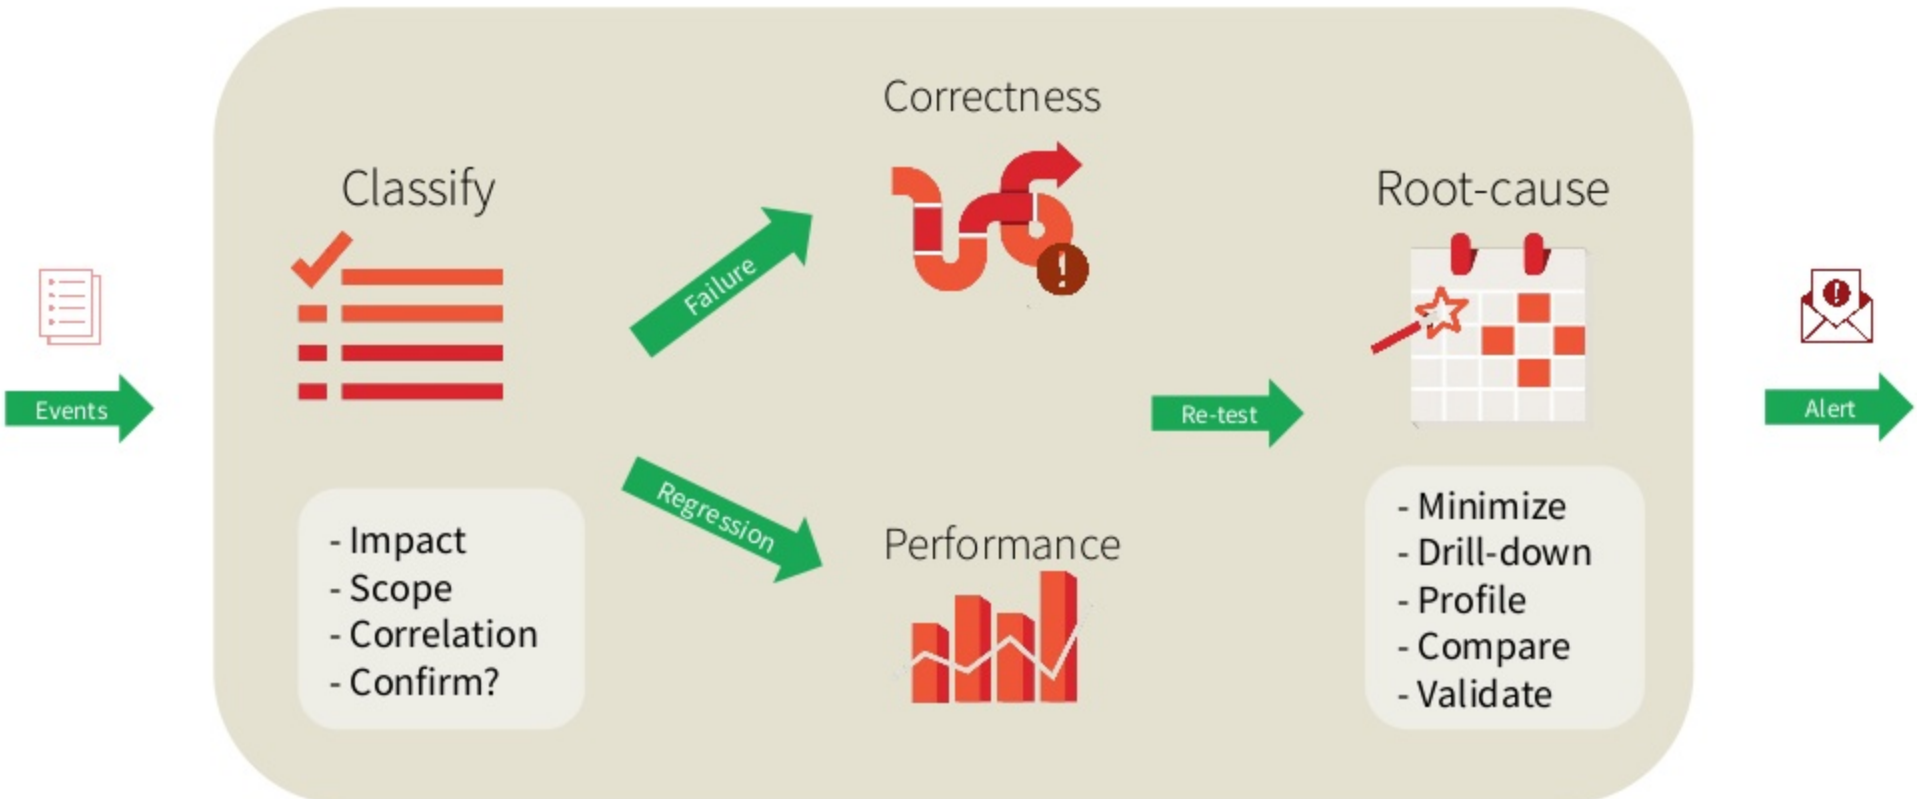

Correctness

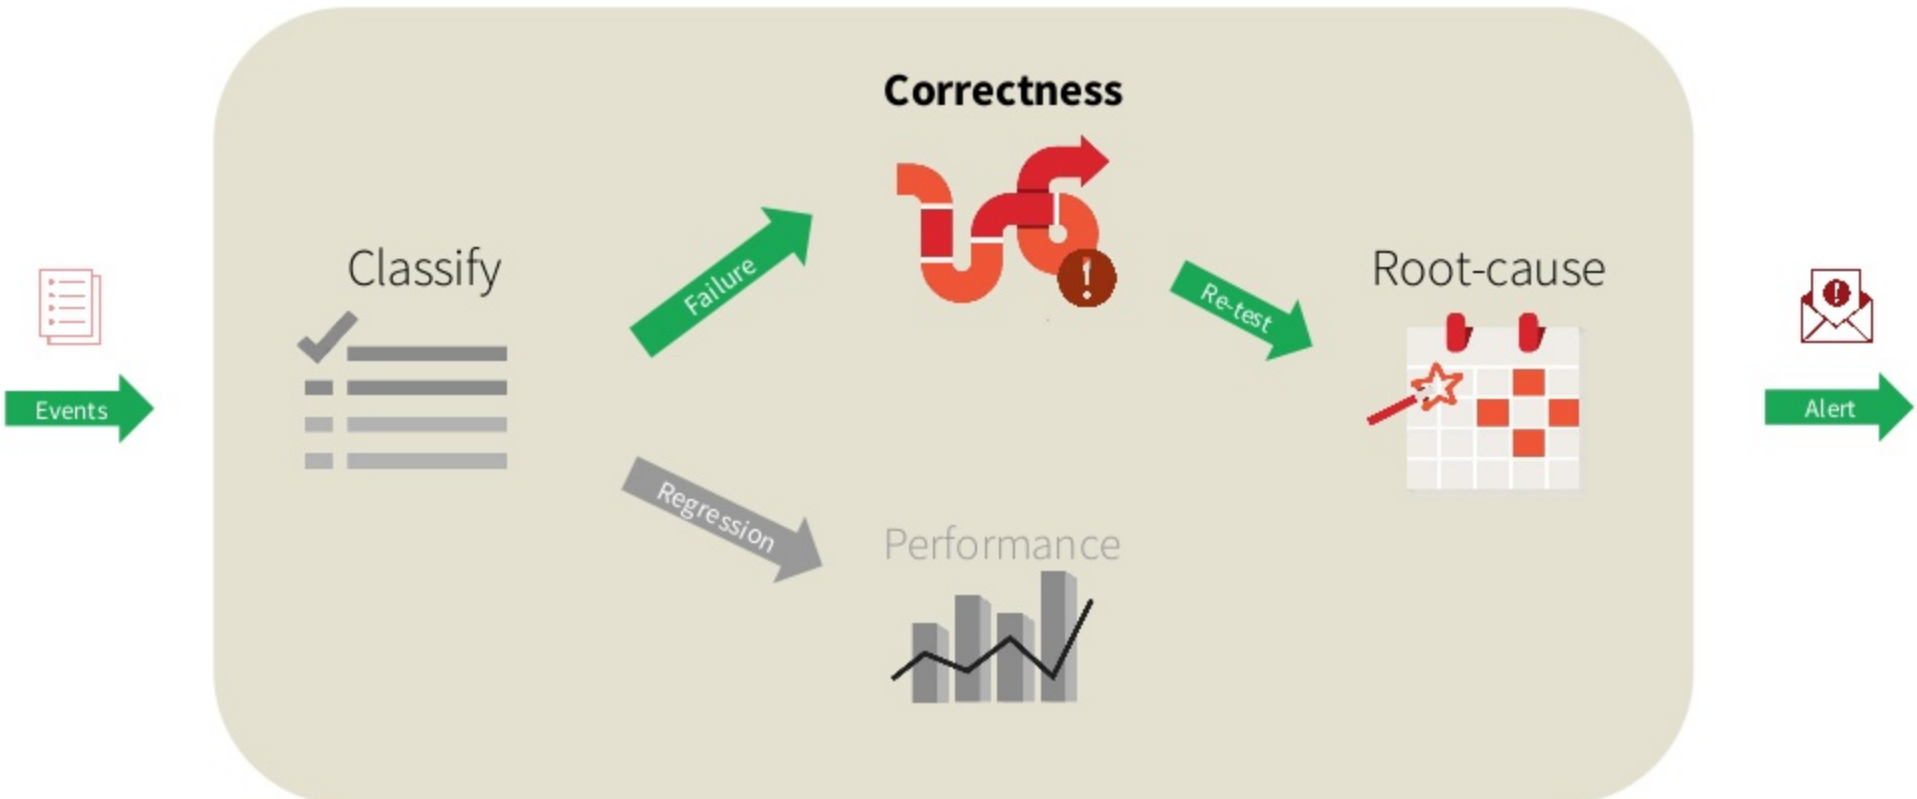

Random query generation

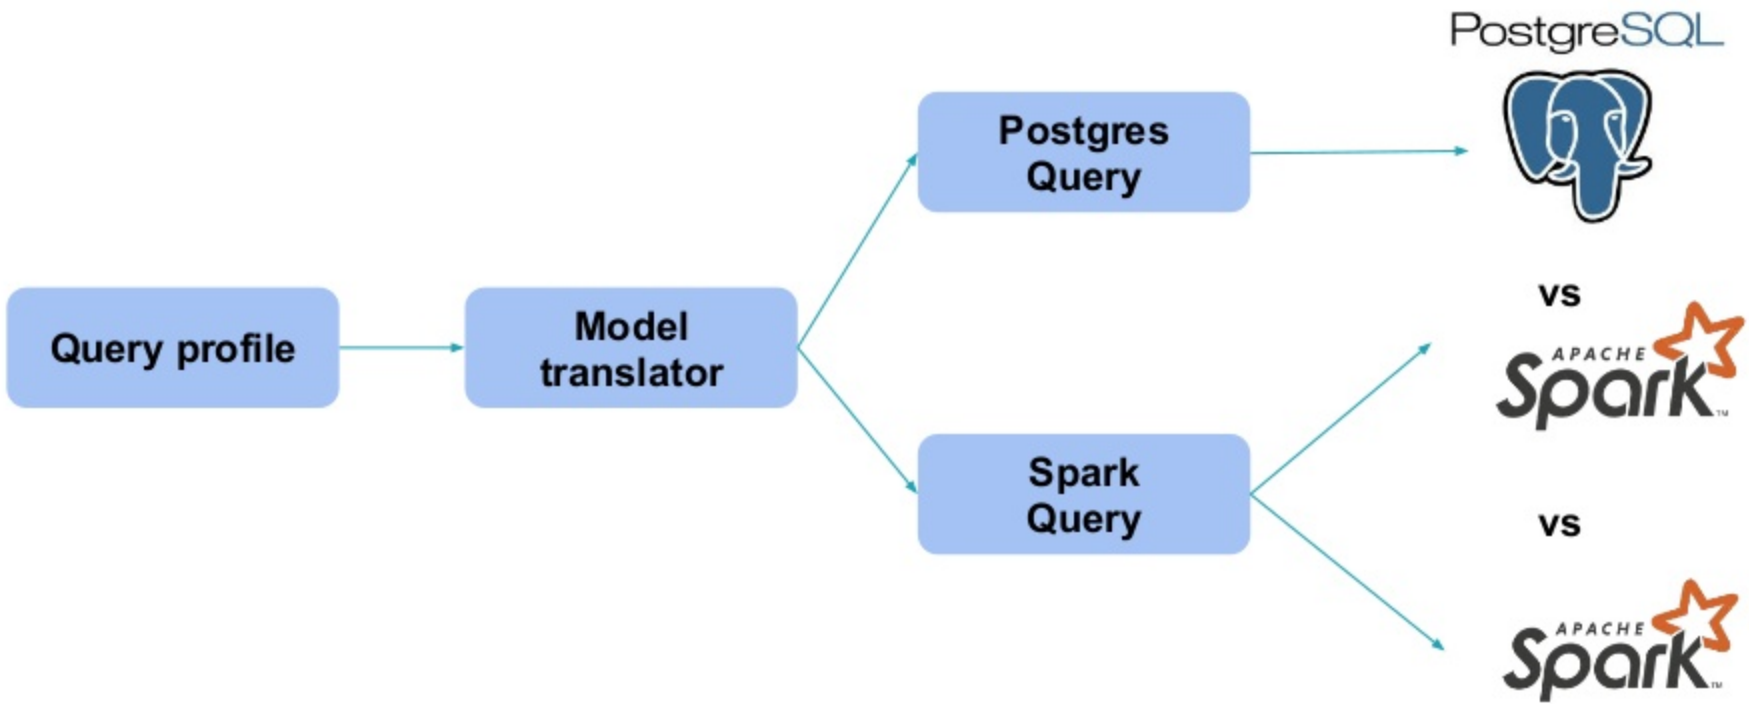

DDL and datagen

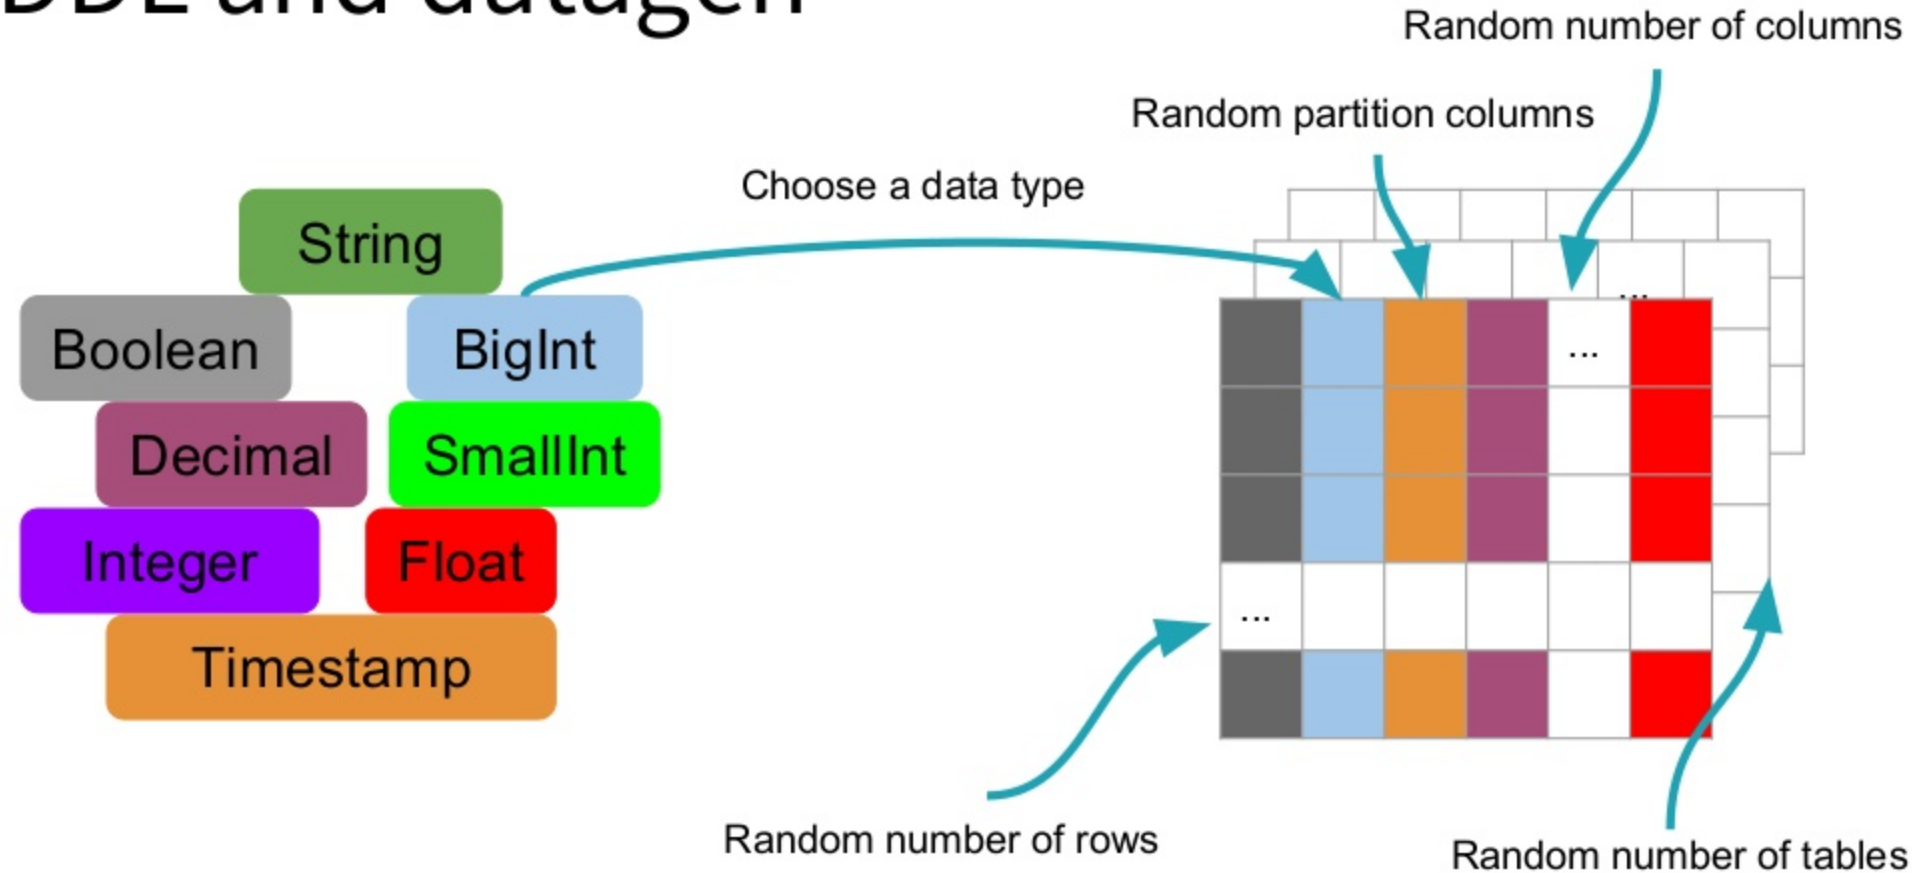

Recursive query model

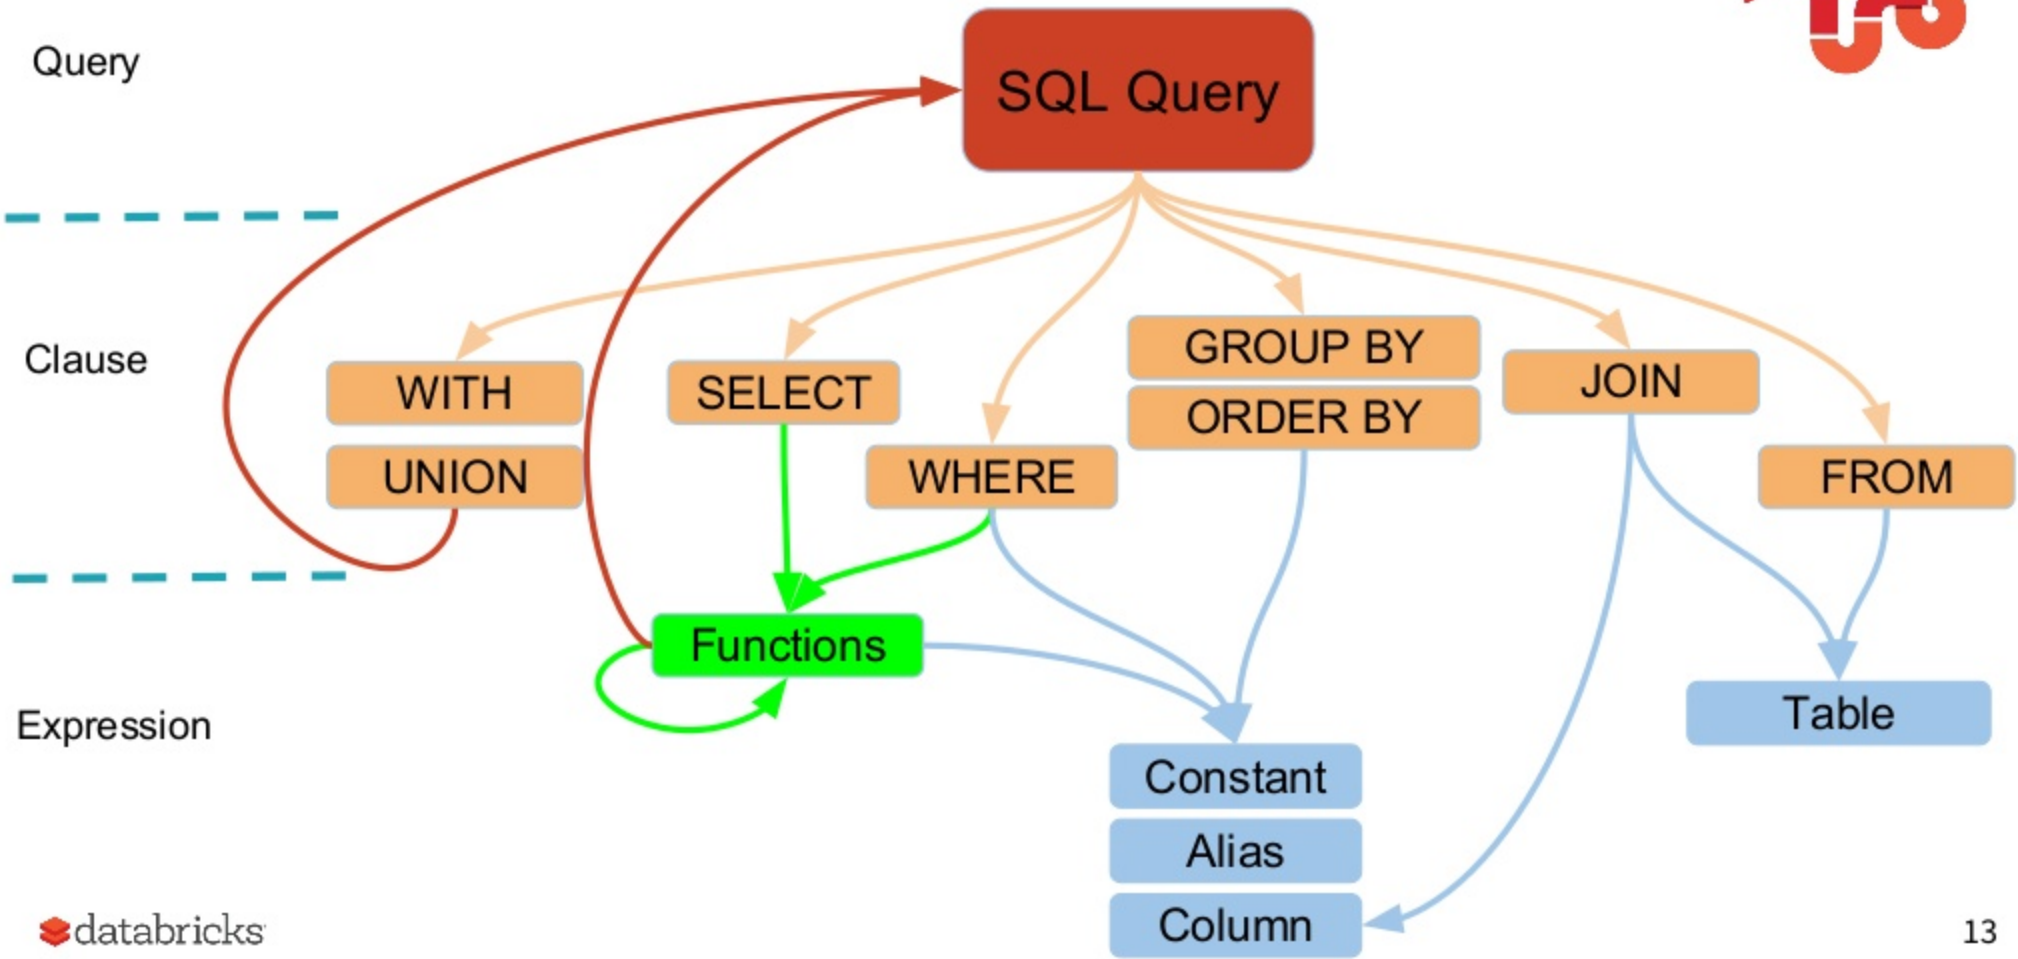

Probabilistic query profile

Independent weights

- Optional query clauses

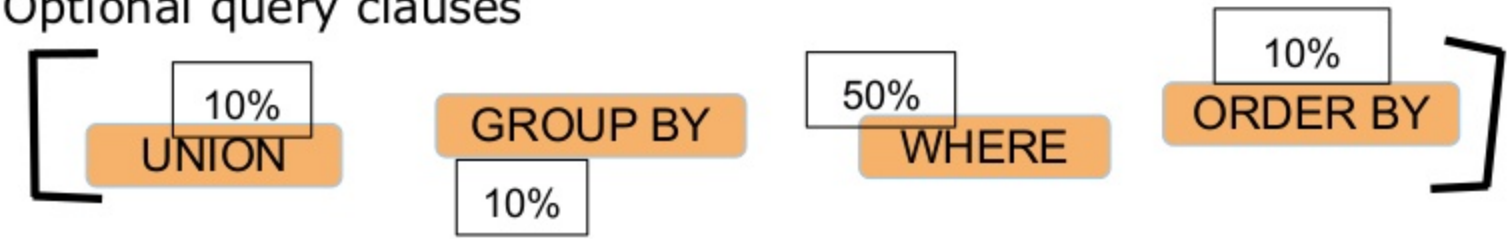

Inter-dependent weights

- Join types
- Select functions

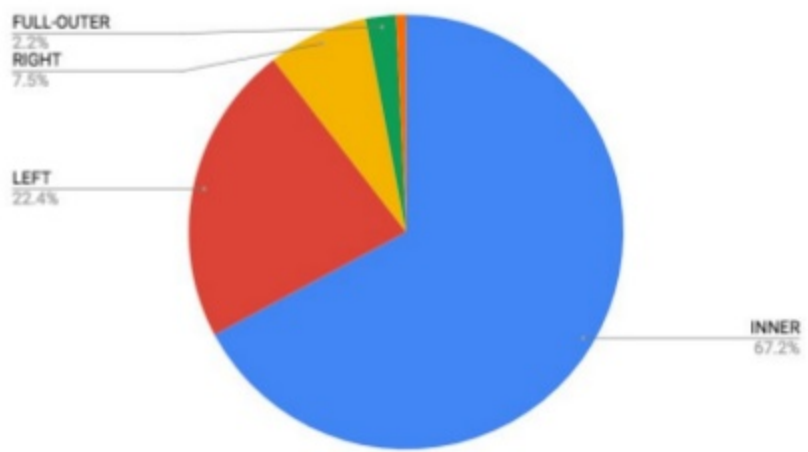

Coalesce flattening (1/4)

```
SELECT COALESCE(t2.smallint_col_3, t1.smallint_col_3, t2.smallint_col_3) AS int_col,
       IF(NULL, VARIANCE(COALESCE(t2.smallint_col_3, t1.smallint_col_3, t2.smallint_col_3)),
       COALESCE(t2.smallint_col_3, t1.smallint_col_3, t2.smallint_col_3)) AS int_col_1,
       STDDEV(t2.double_col_2) AS float_col,
       COALESCE(MIN((t1.smallint_col_3) - (COALESCE(t2.smallint_col_3, t1.smallint_col_3,
       t2.smallint_col_3))), COALESCE(t2.smallint_col_3, t1.smallint_col_3, t2.smallint_col_3),
       COALESCE(t2.smallint_col_3, t1.smallint_col_3, t2.smallint_col_3)) AS int_col_2
FROM table_4 t1
INNER JOIN table_4 t2 ON (t2.timestamp_col_7) = (t1.timestamp_col_7)
WHERE (t1.smallint_col_3) IN (CAST('0.04' AS DECIMAL(10,10)), t1.smallint_col_3)
GROUP BY COALESCE(t2.smallint_col_3, t1.smallint_col_3, t2.smallint_col_3)
```

Small dataset with 2 tables of 5x5 size
Within 10 randomly generated queries

Error: Operation is in ERROR_STATE

Coalesce flattening (2/3)

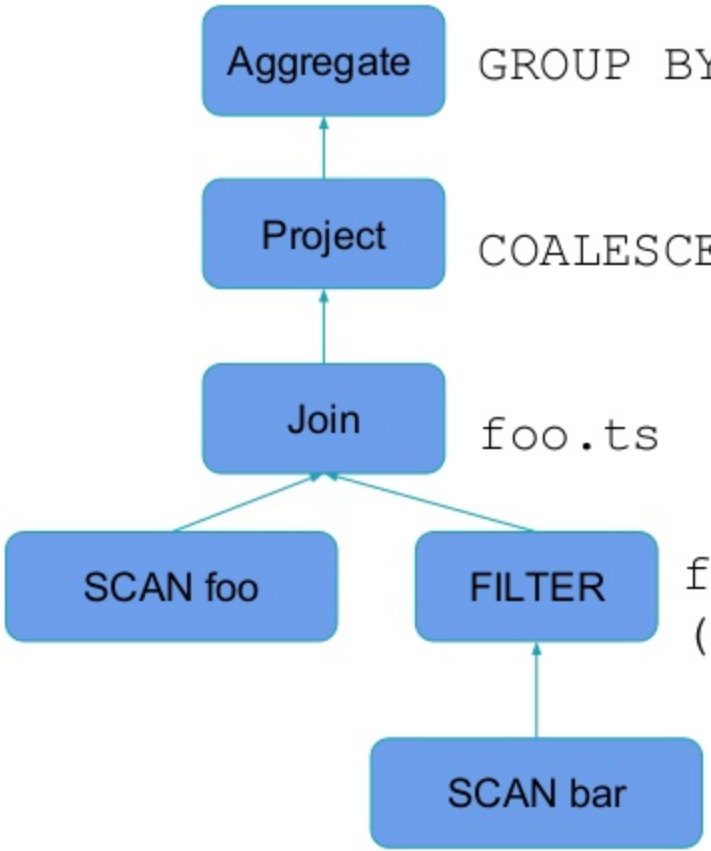


```
GROUP BY COALESCE(foo.id, foo.val)

COALESCE(COALESCE(foo.id, foo.val), 88)

foo.ts  = bar.ts

foo.id IN
(CAST('0.04' AS DECIMAL(10, 10)), foo.id)
```

Coalesce flattening (3/4)

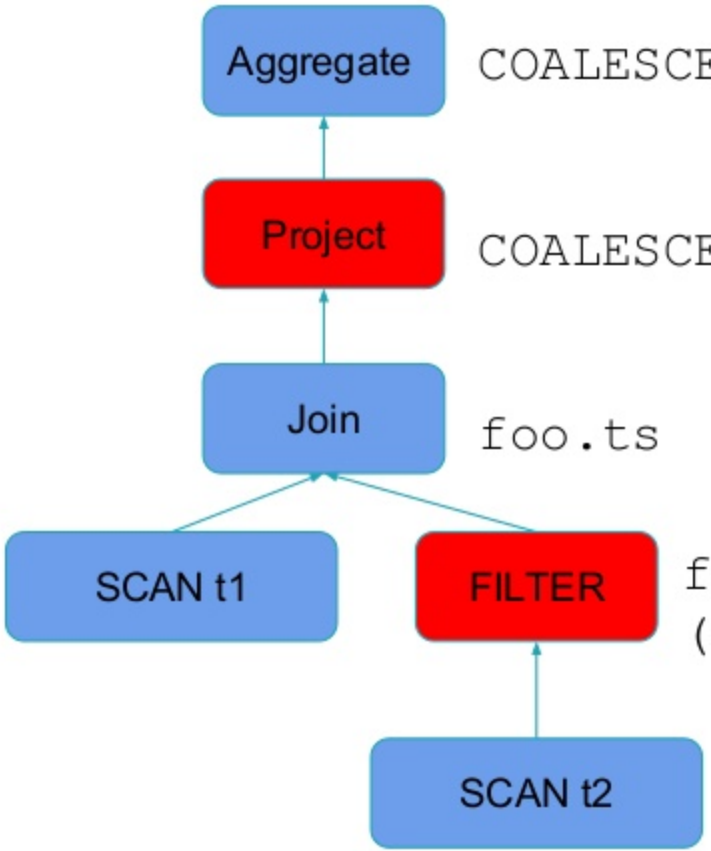


```
COALESCE(foo.id, foo.val)

COALESCE(COALESCE(foo.id, foo.val), 88)

foo.ts  = bar.ts

foo.id IN
(CAST('0.04' AS DECIMAL(10, 10)), foo.id)
```

Coalesce flattening (4/4)

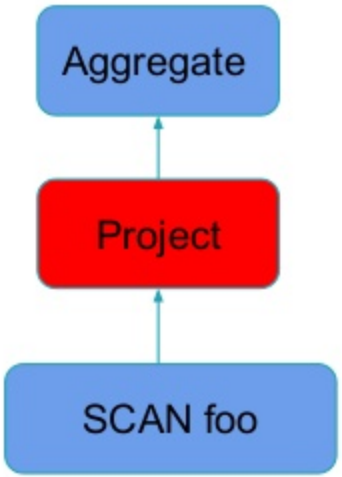

Minimized query:

```
SELECT
   COALESCE(COALESCE(foo.id, foo.val), 88)
FROM foo
GROUP BY
   COALESCE(foo.id, foo.val)
```

Analyzing the error

- The optimizer flattens the nested coalesce calls
- The SELECT clause doesn't contain the GROUP BY expression
- Possibly a problem with any GROUP BY expression that can be optimized

Lead function (1/3)

```
SELECT (t1.decimal0803_col_3) / (t1.decimal0803_col_3) AS decimal_col,
        CAST(696 AS STRING) AS char_col, t1.decimal0803_col_3,
       (COALESCE(CAST('0.02' AS DECIMAL(10,10)),
                 CAST('0.47' AS DECIMAL(10,10)),
                 CAST('-0.53' AS DECIMAL(10,10)))) +
       (LEAD(-65, 4) OVER (ORDER BY (t1.decimal0803_col_3) / (t1.decimal0803_col_3),
                 CAST(696 AS STRING))) AS decimal_col_1,
                 CAST(-349 AS STRING) AS char_col_1

FROM table_16 t1
WHERE (943) > (889)
```

Error: Column 4 in row 10 does not match:

[1.0, 696, -871.81, <<-64.98>>, -349] SPARK row

[1.0, 696, -871.81, <<None>>, -349] POSTGRESQL row

Lead function (2/3)

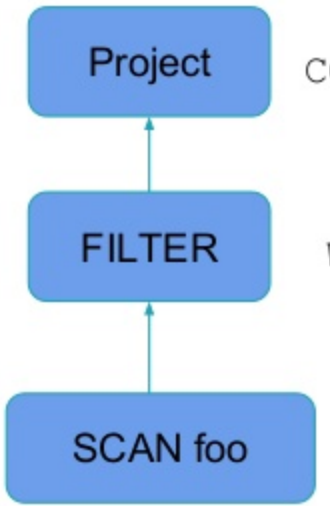


```
COALESCE(expr) + LEAD(-65, 4) OVER ORDER BY expr
```

```
WHERE expr
```

Lead function (3/3)

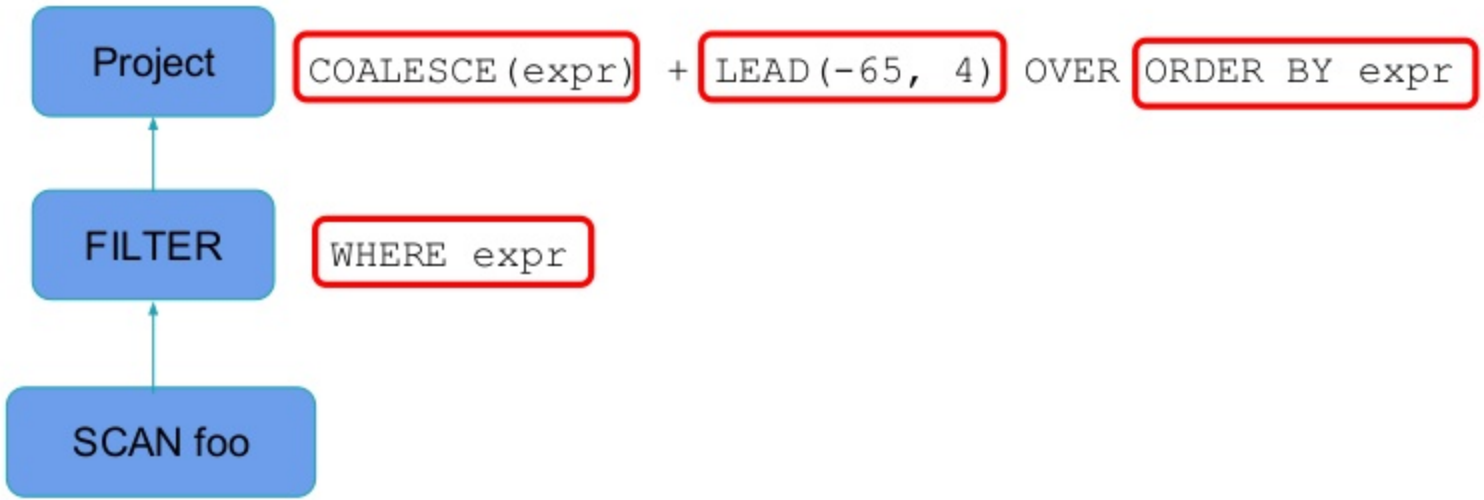

Analyzing the error

- Using constant input values breaks the behaviour of the LEAD function
- SC-16633: https://github.com/apache/spark/pull/14284

Performance

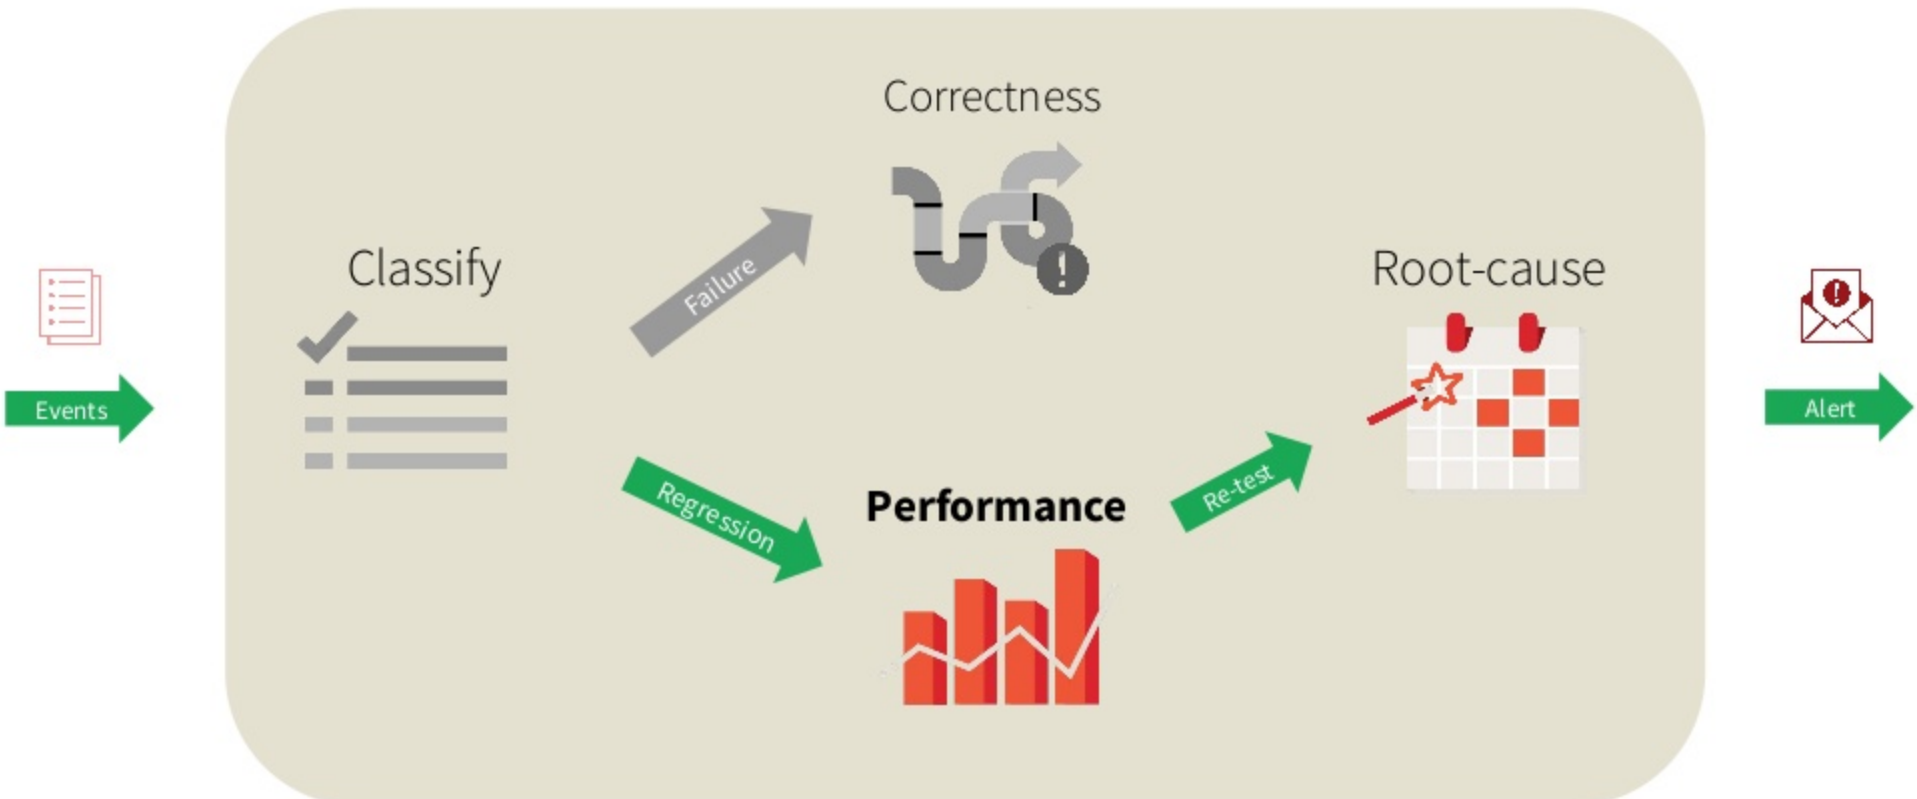

Benchmarking tools

- We use spark-sql-perf public library for TPC workloads
  - Provides datagen and import scripts
  - local, cluster, S3
  - Dashboards for analyzing results

- The Spark micro benchmarks
- And the async-profiler
  - to produce flamegraphs

README.md

Spark SQL Performance Tests

build passing

This is a performance testing framework for Spark SQL in Apache Spark 2.2+.

Note: This README is still under development. Please also check our source code for more information.

Quick Start

Running from command line.

https://github.com/databricks/spark-sql-perf

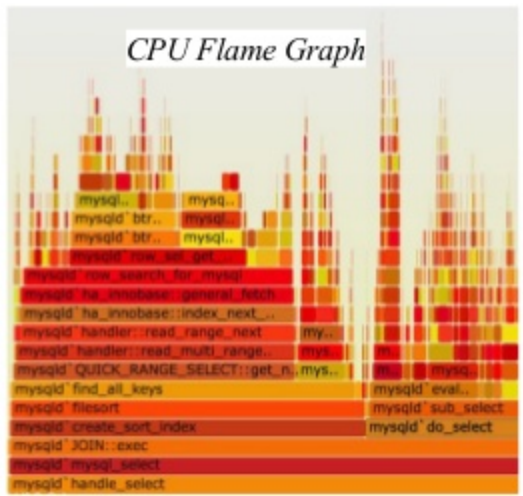

http://www.brendangregg.com/flamegraphs.html

DBR 5.0-beta (v2.4) performance ~~tracking~~ -- journey

daysBack : 30 daysTargetBack : 30 targetVersion : master-2.4

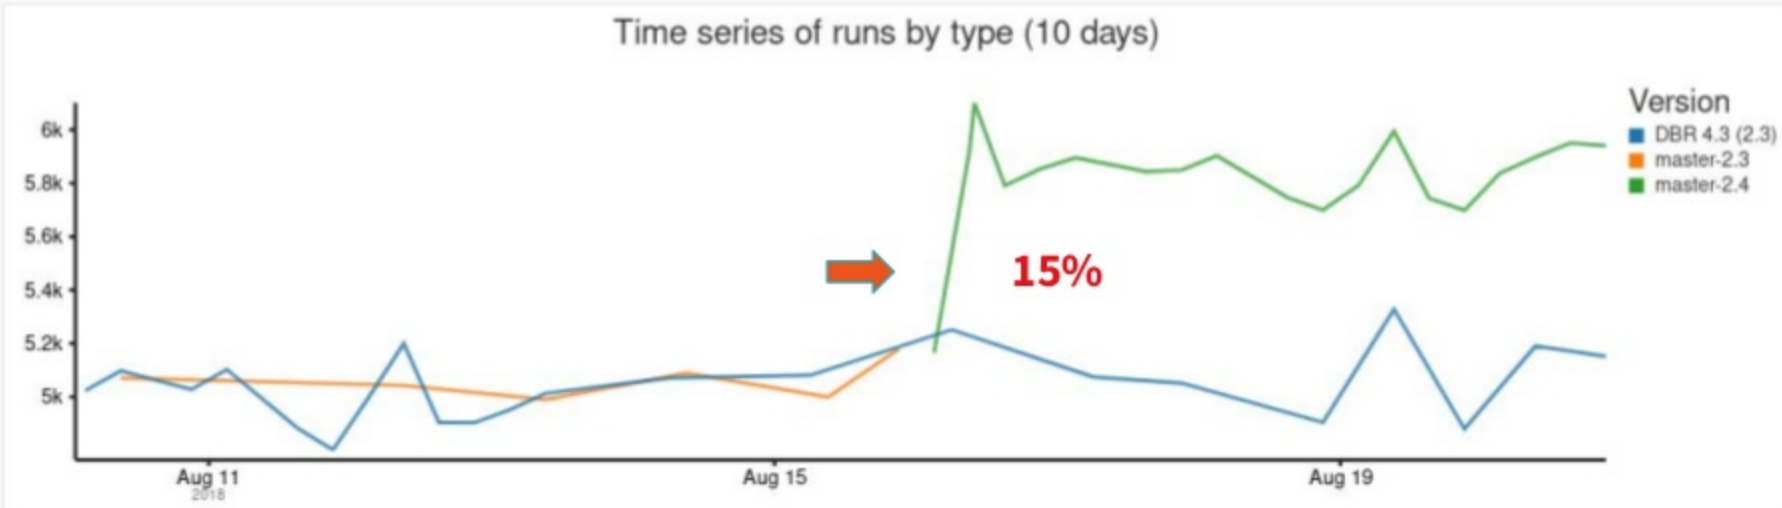

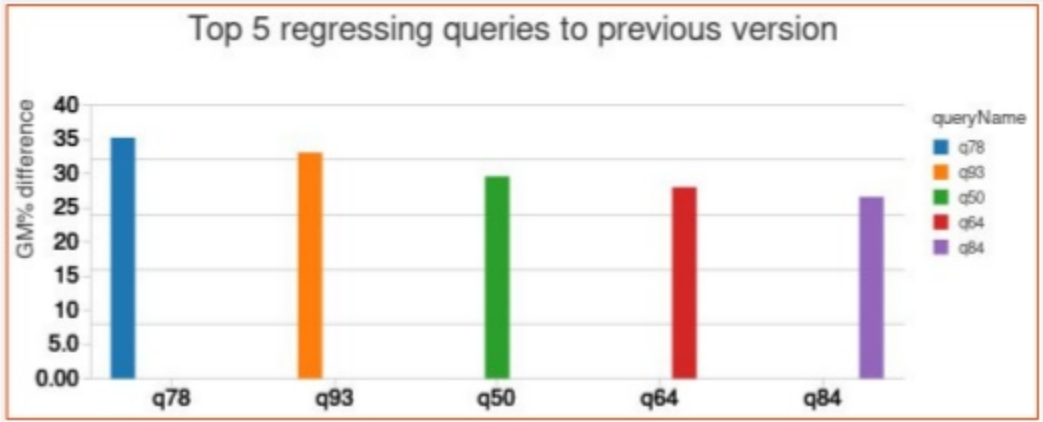

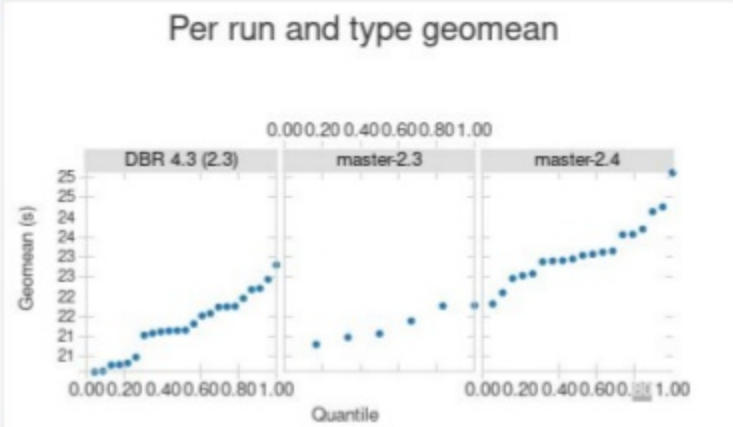

Per query drill-down: 67

First, **scope** and **validate**

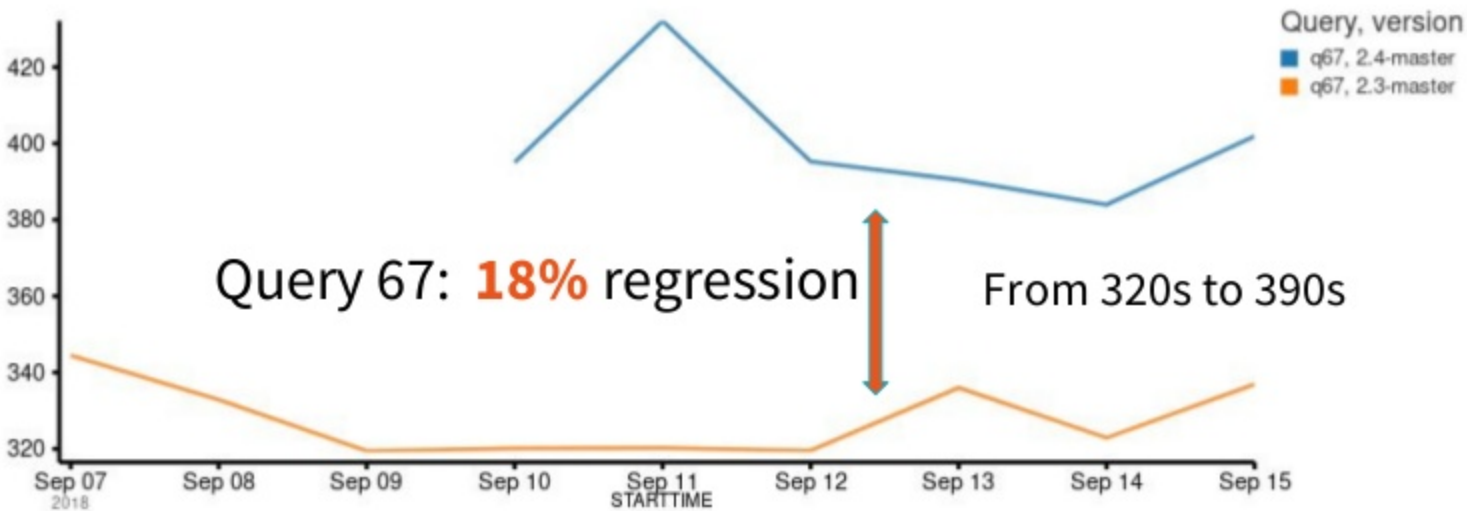

- in 2.4-master (dev) compared
- to 2.3 in DBR 4.3 (prod)

Q67 executor profile for Spark 2.4-master

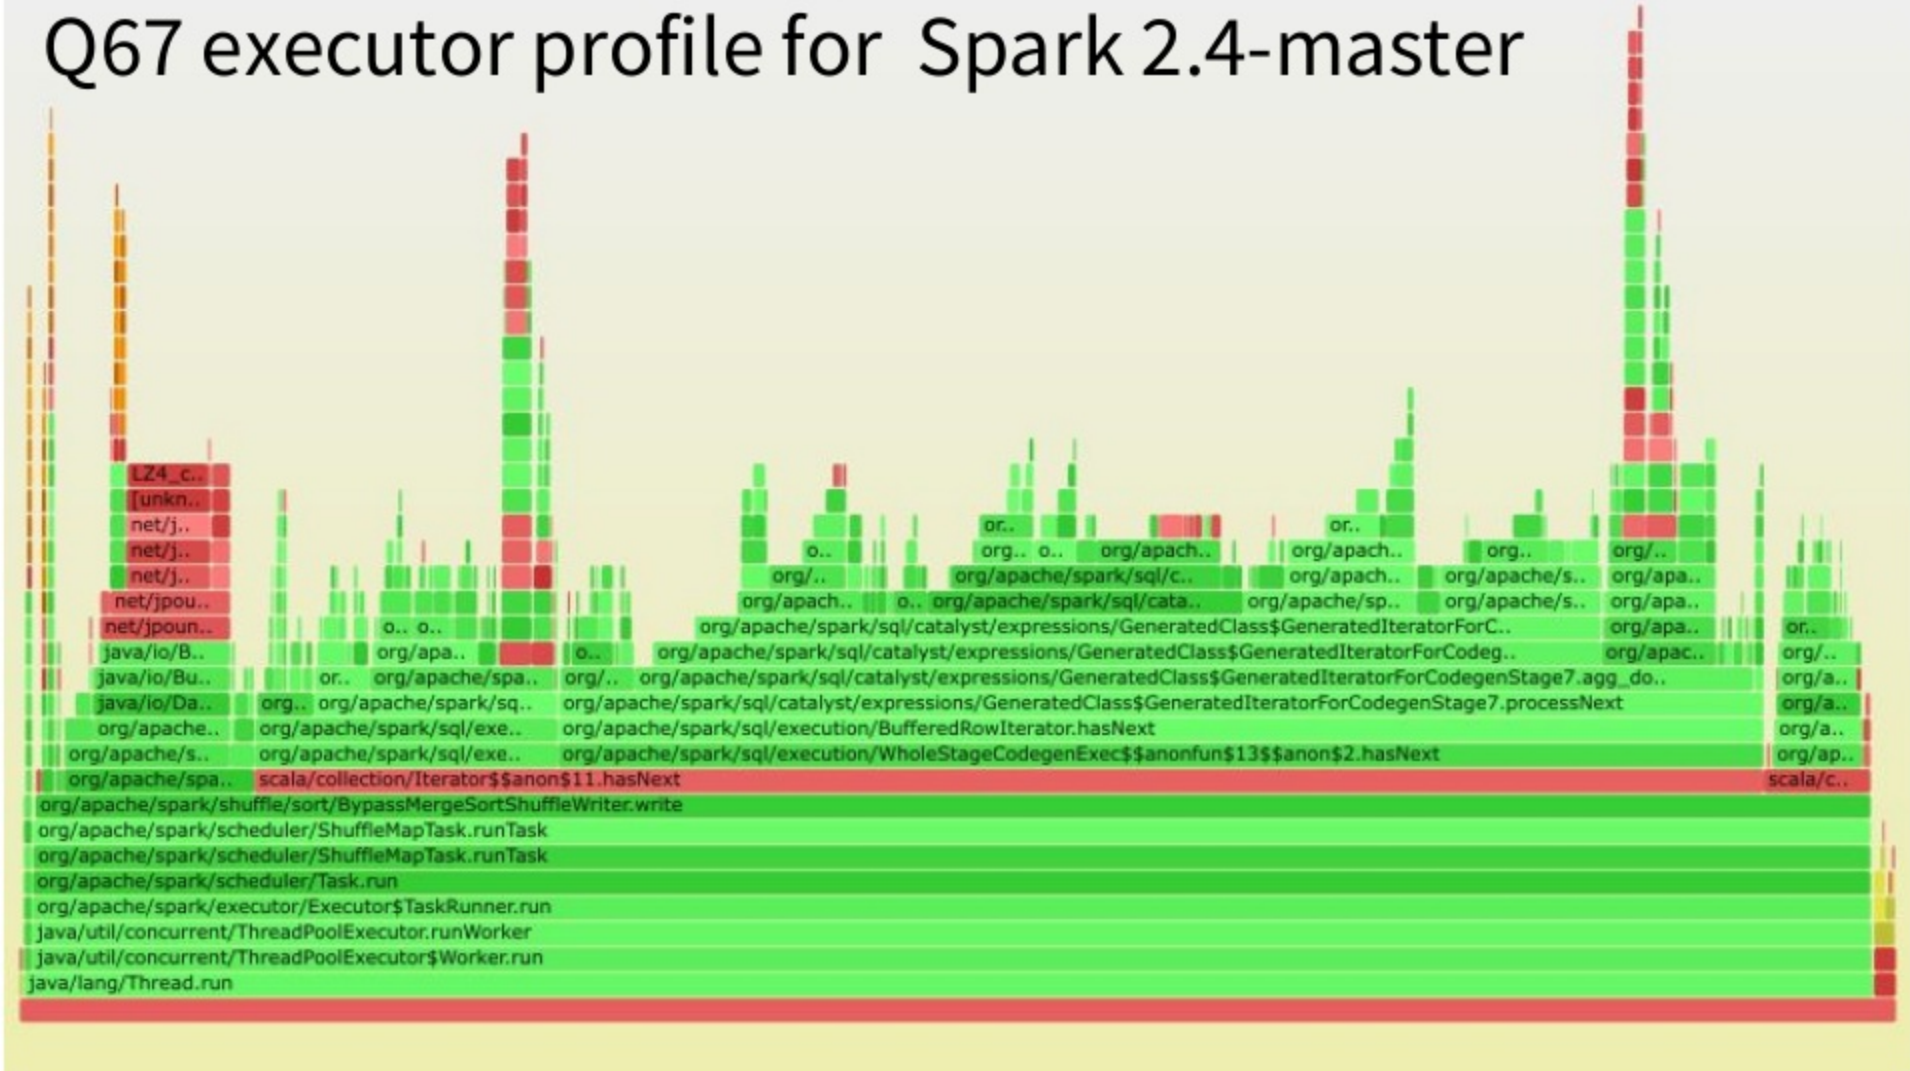

Side-by-side 2.3 vs 2.4: find the differences

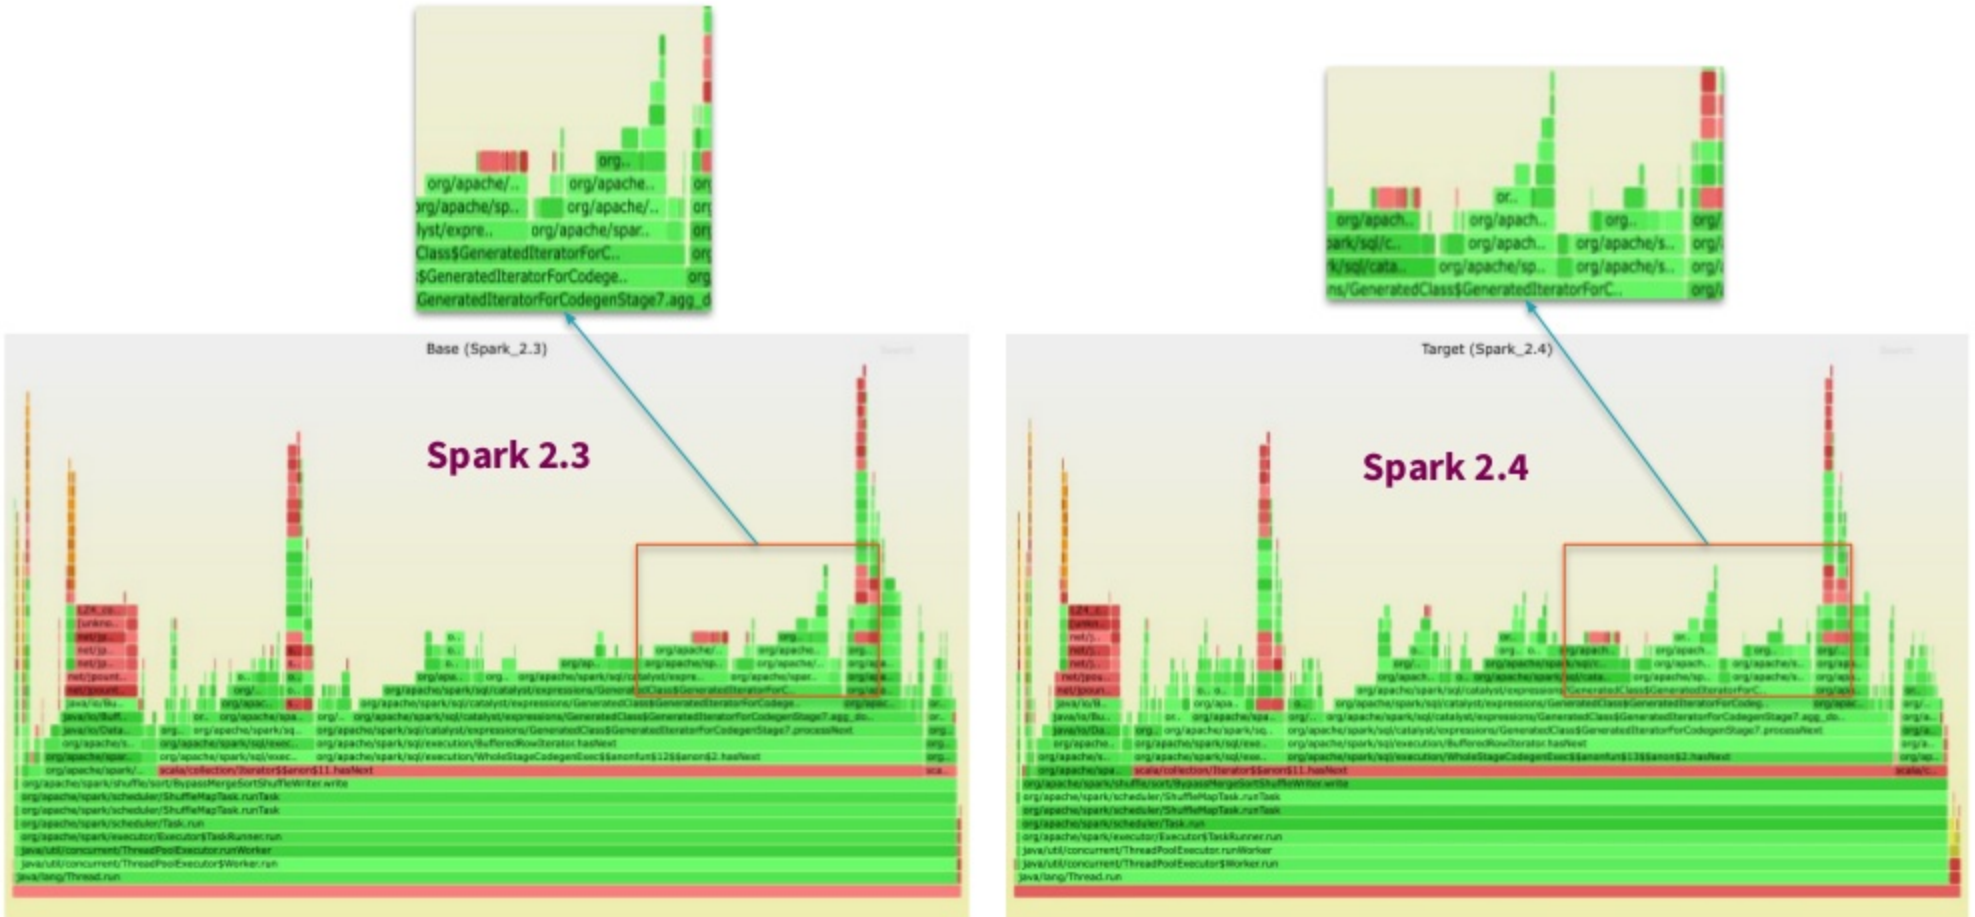

Framegraph diff zoom

Red slower **White** new

Look for hints:

- Mem mgmt
- Hashing
- unsafe

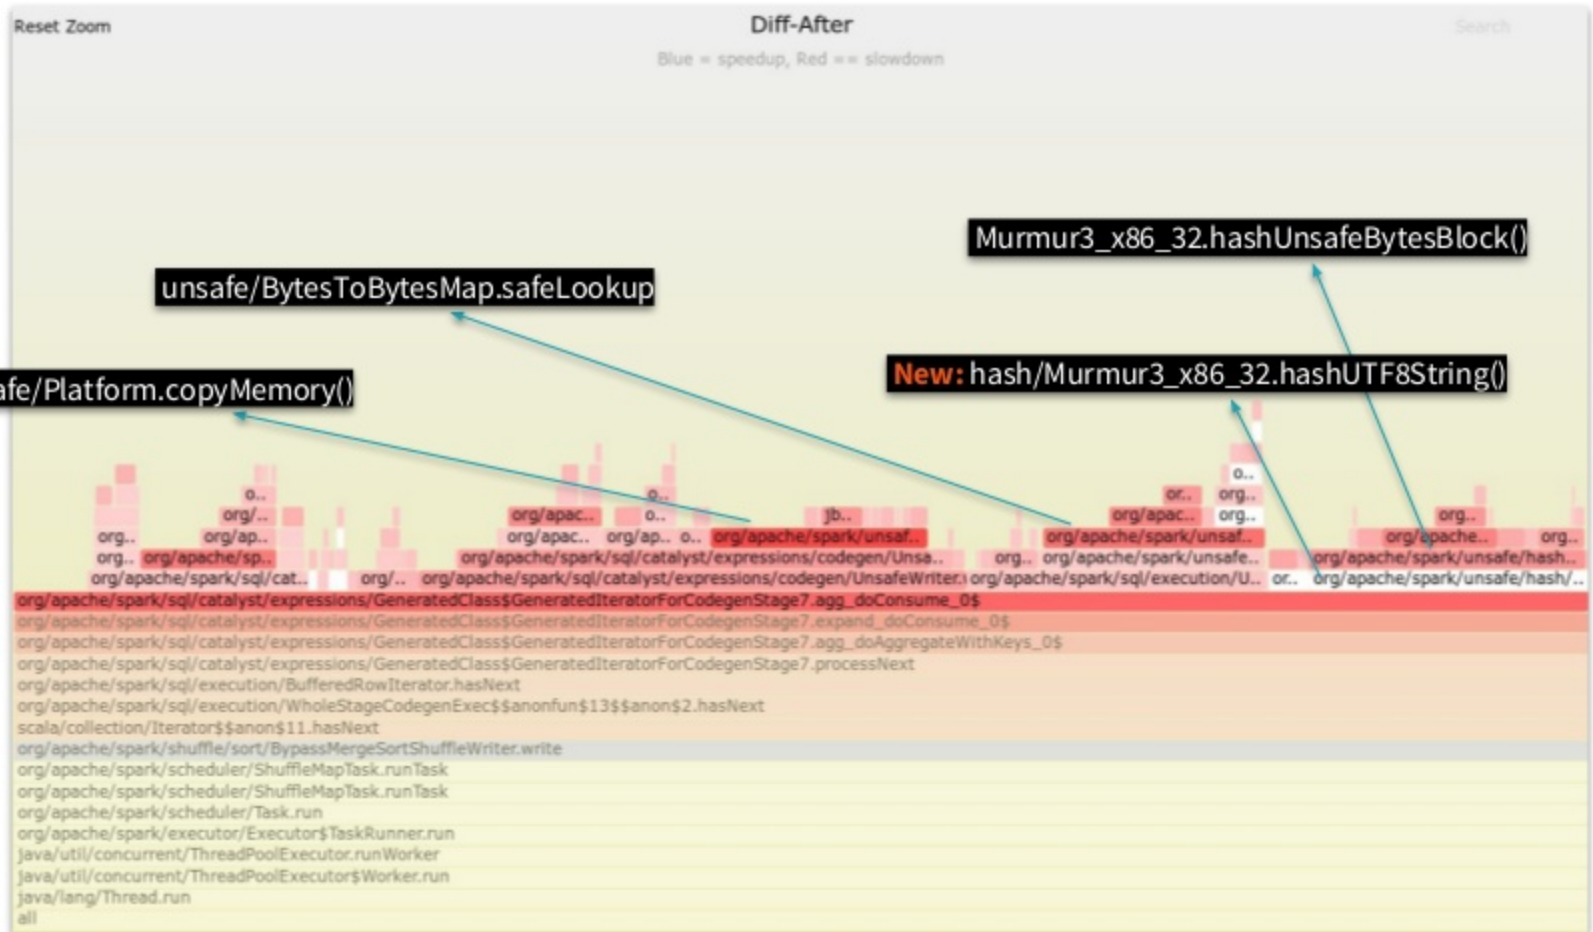

databricks

Root-causing

Microbenchmark for UTF8String

```
test("hashing") {
    import org.apache.spark.unsafe.hash.Murmur3_x86_32
    import org.apache.spark.unsafe.types.UTF8String
    val hasher = new Murmur3_x86_32(0)
    val str = UTF8String.fromString("b" * 10001)
    val numIter = 100000
    val start = System.nanoTime
    for(i <- 0 until numIter) {
      Murmur3_x86_32.hashUTF8String(str, 0)
```

Results:

- Spark 2.3: hashUnsafeBytes() -> **40µs**
- Spark 2.4 *hashUnsafeBytesBlock()* -> **140µs**

- also slower *UTF8String.getBytes()*

GIT BISECT

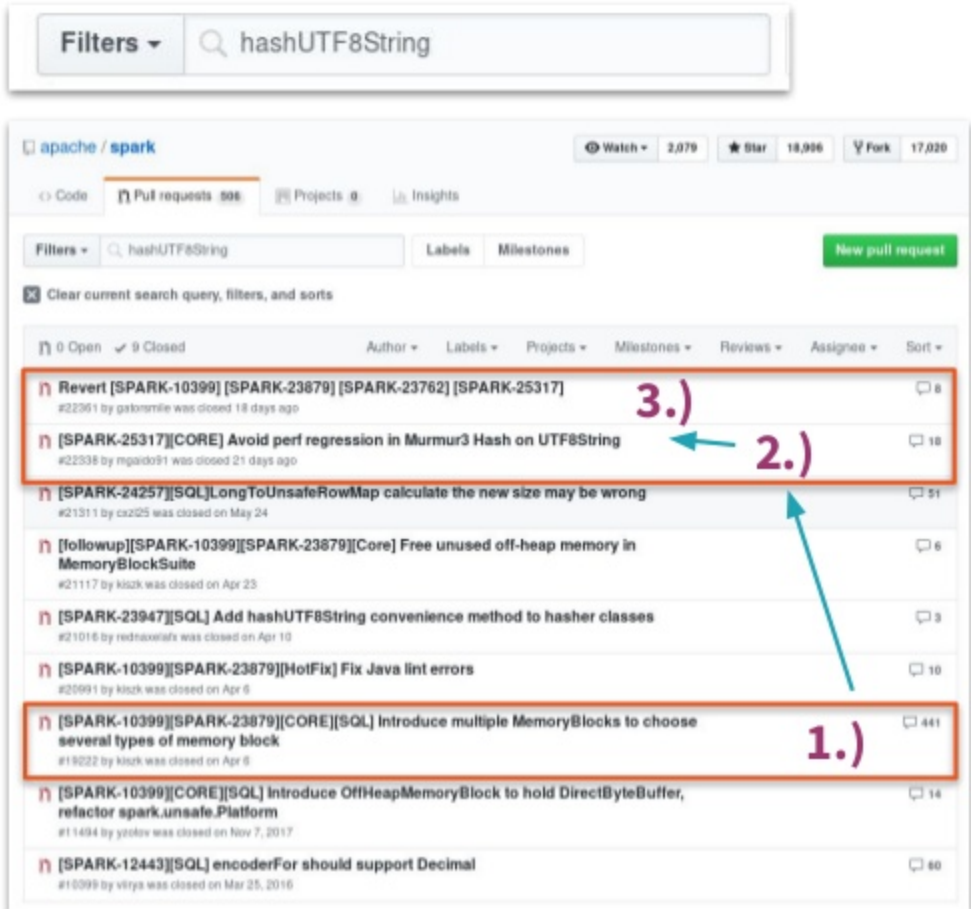

It is a journey to get a release out

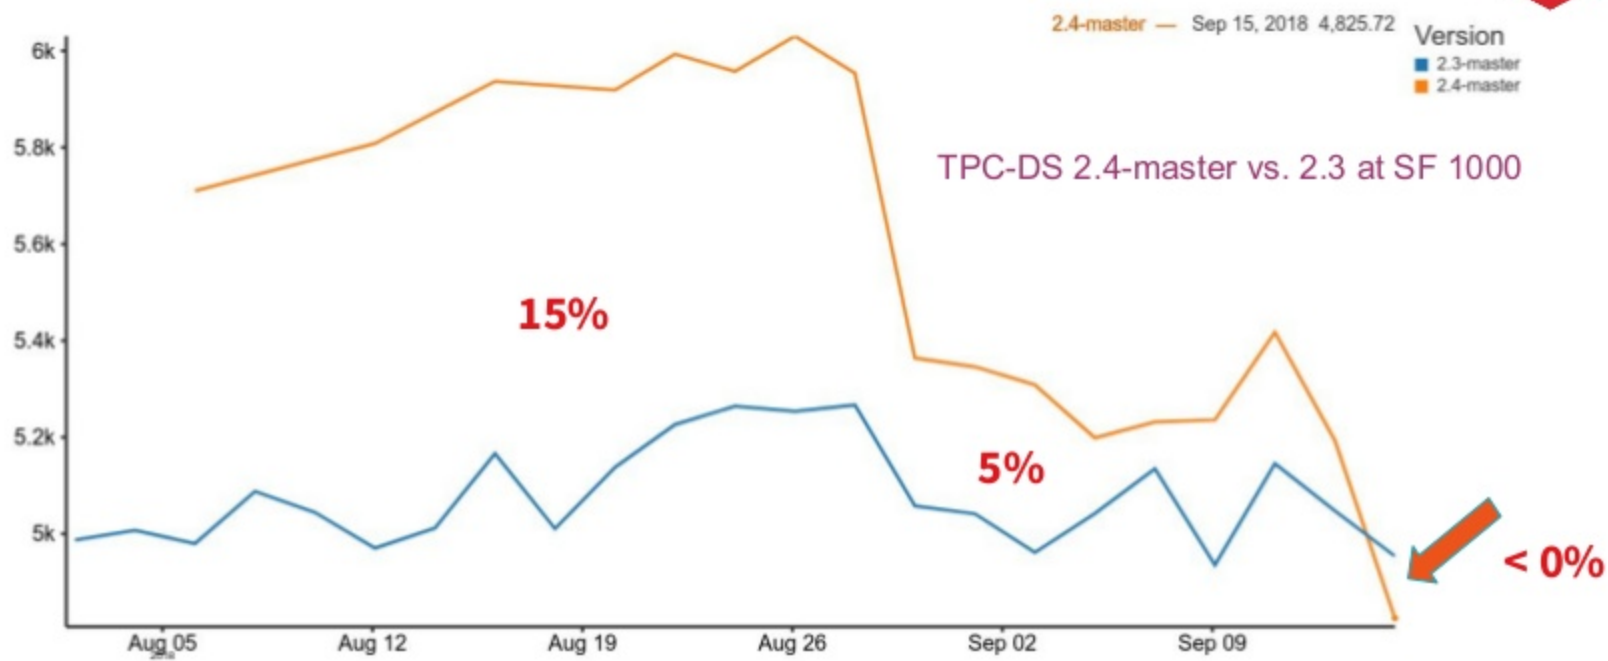

DBR and Spark testing and performance are a **continuous effort**

- **Over a month** effort to bring performance to improving

Conclusion

Spark in production is *not just the framework*

Unit and integration testing are not enough

We need Spark specific tools to automate the process

to ensure both correctness and performance

Thanks!

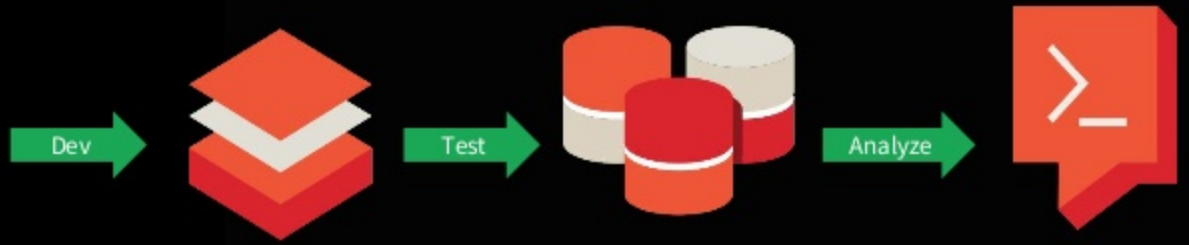

Correctness and Performance of Apache Spark SQL

October 2018

databricks