

Towards Writing Scalable Spark Applications

Philipp Brunenberg, Independent Big Data Consultant



philipp-brunenberg.de



[@p_brunenberg](https://twitter.com/p_brunenberg)



[Philipp Brunenberg](#)

#AmazingSpark

Data decomposition

*independently of
all others*

`map(k1, v1) -> list(k2, v2)`
`reduce(k2, iterator(v2))`

`map(k1, v1) -> list(k2, v2)`

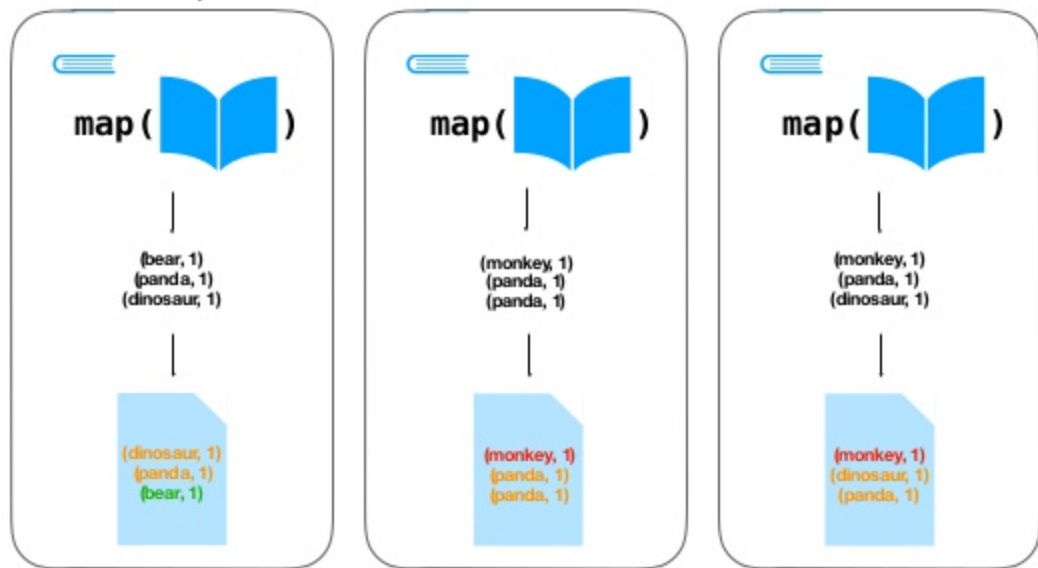
MASTER

Array[TaskState]
Array[MapOutputFileLocation]

M=6

R=3

$\text{hash}(\text{panda}) \% R = 0$
 $\text{hash}(\text{dinosaur}) \% R = 0$
 $\text{hash}(\text{bear}) \% R = 1$
 $\text{hash}(\text{monkey}) \% R = 2$



`reduce(k2, iterator(v2))`

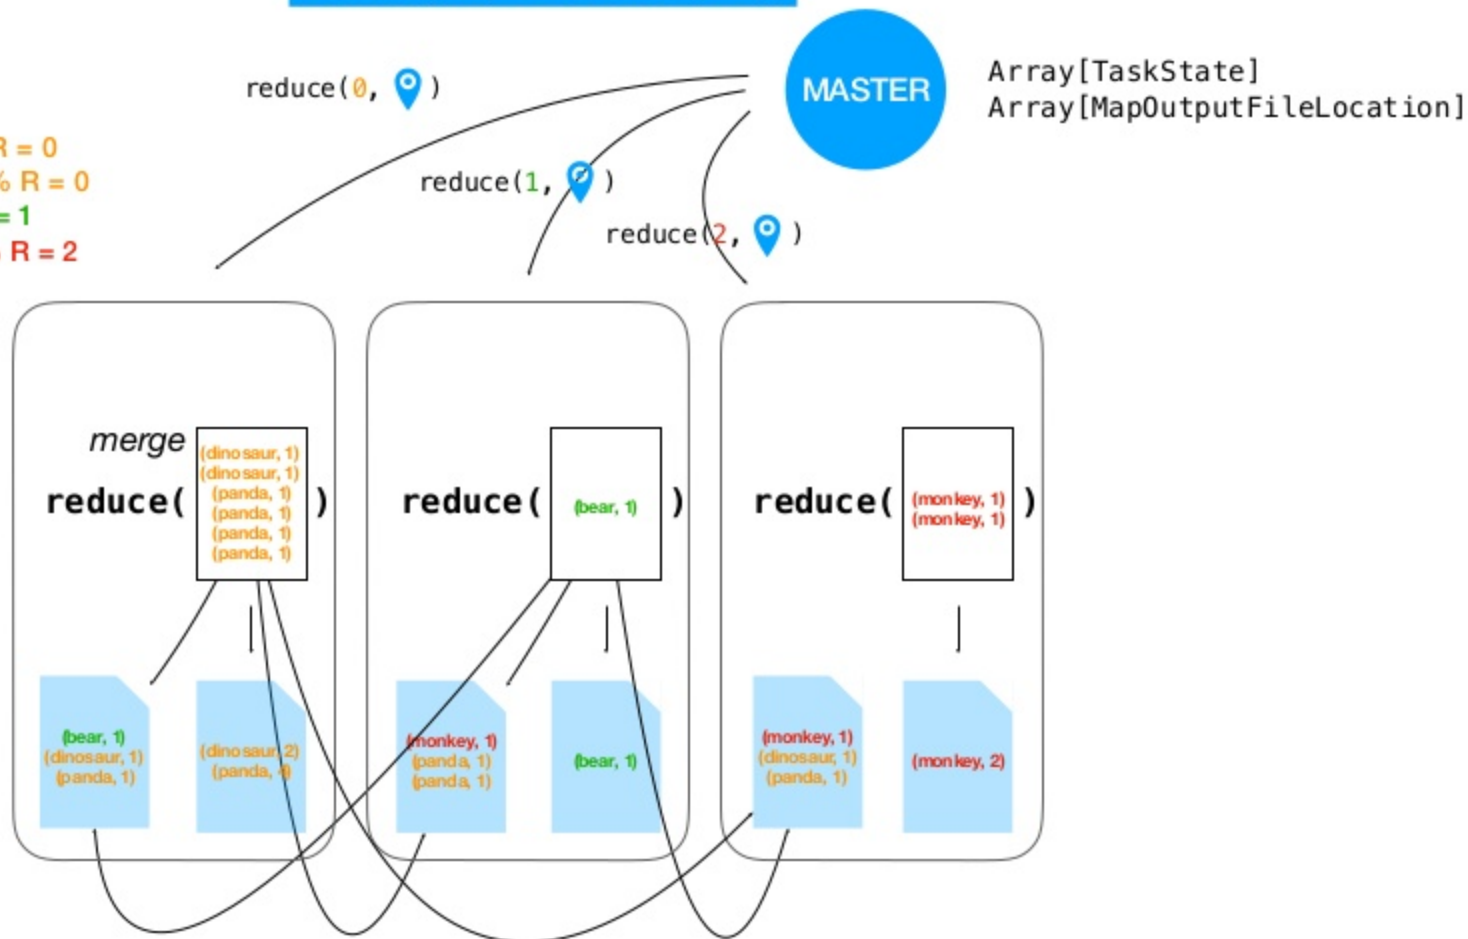
R=3

`hash(panda) % R = 0`

`hash(dinosaur) % R = 0`

`hash(bear) % R = 1`

`hash(monkey) % R = 2`



User Program

```
spark  
  .load()  
  .withColumn()  
  .groupBy()  
  .count()  
  .show()
```

```


rdd1 = sparkContext
    .parallelize(...)
    .repartition(3)

rdd2 = sparkContext
    .parallelize(...)
    .map(...)

rdd3 = sparkContext
    .parallelize(...)

rdd2.union(rdd2)
    .union(rdd)
    .join(rdd1)
    .count()

```



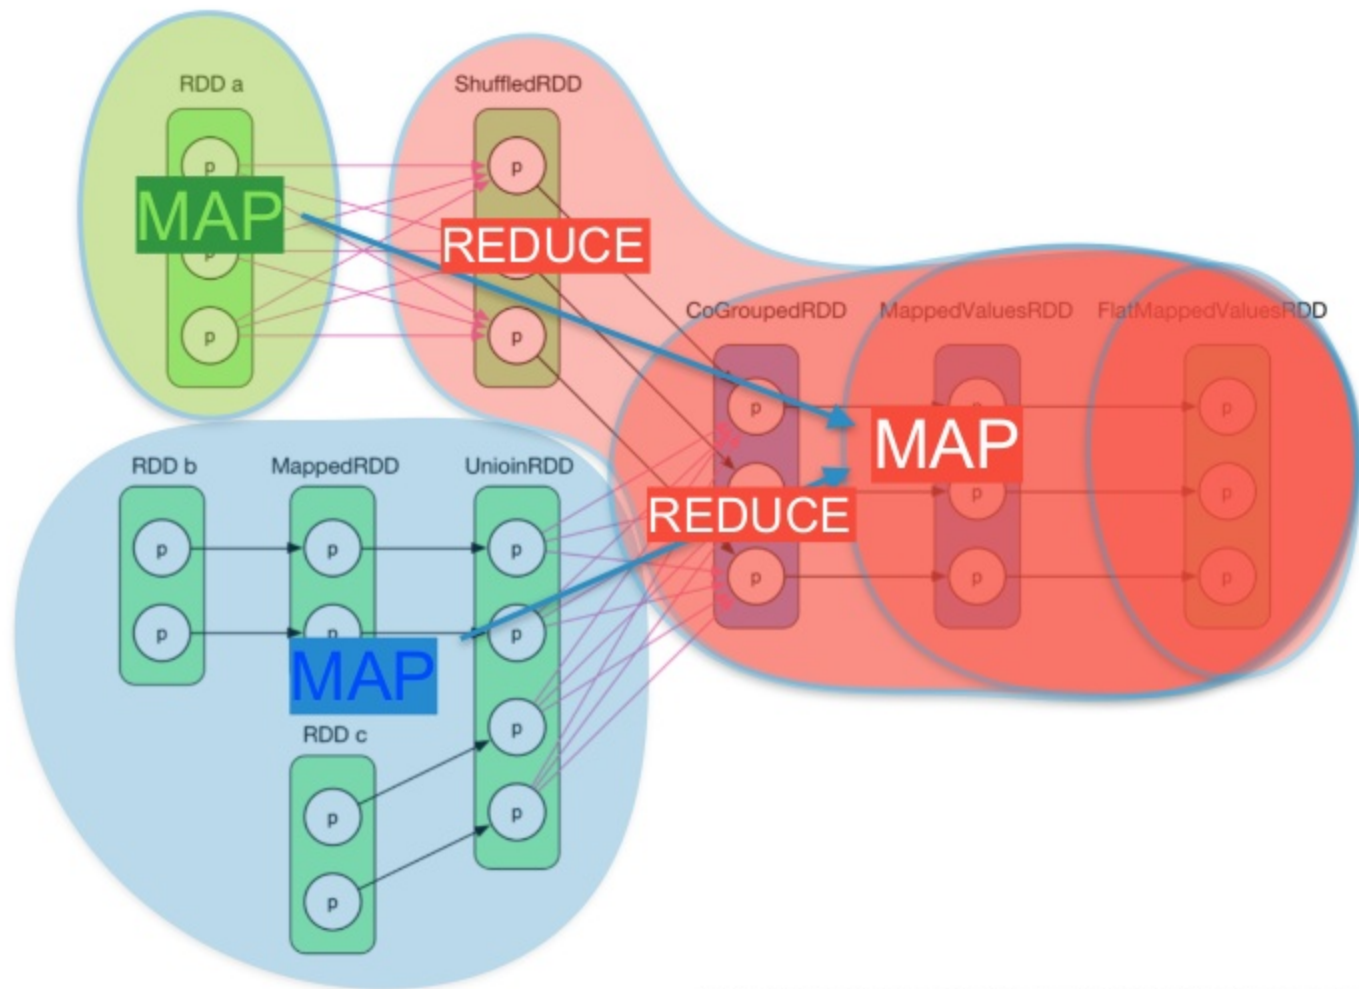

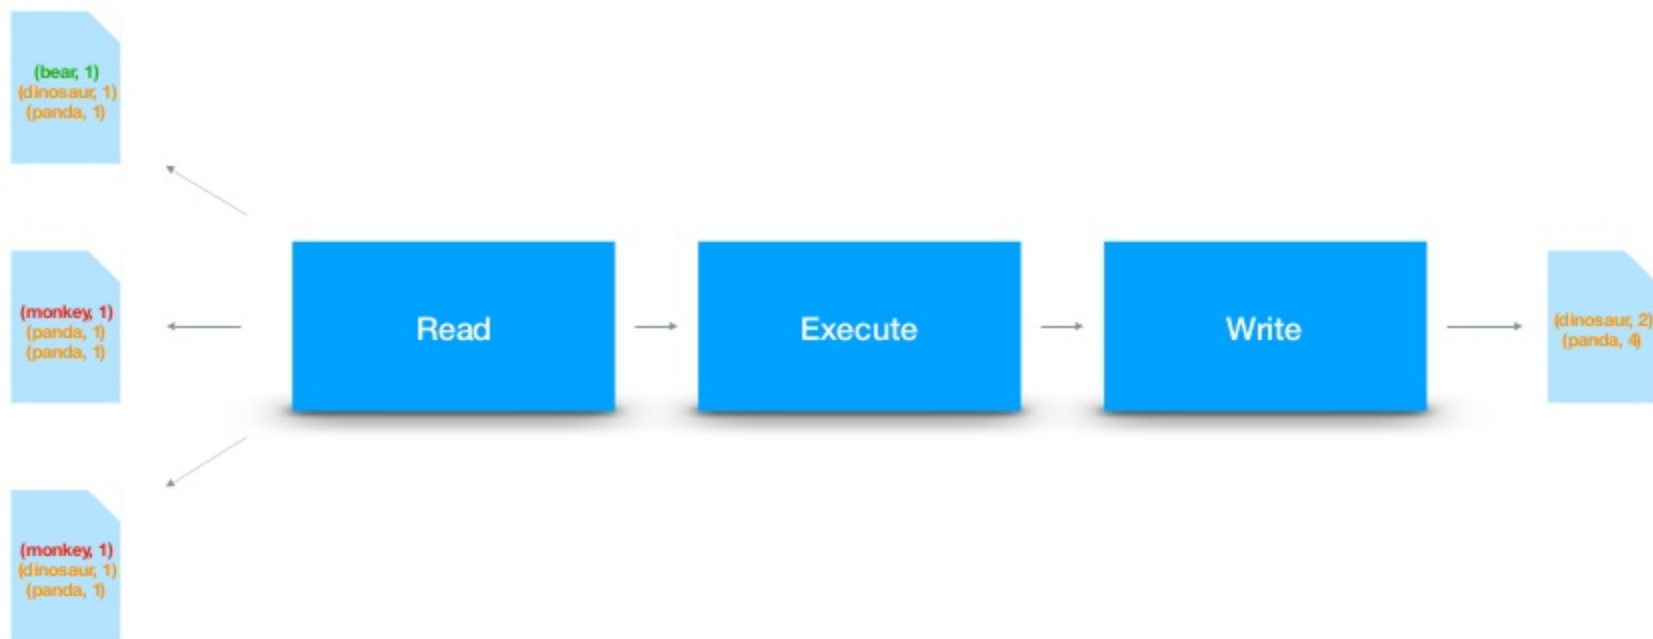


Image from: <https://github.com/JerryLead/SparkInternals/blob/master/PNGfigures/ComplexJob.png>


Performance Bottlenecks





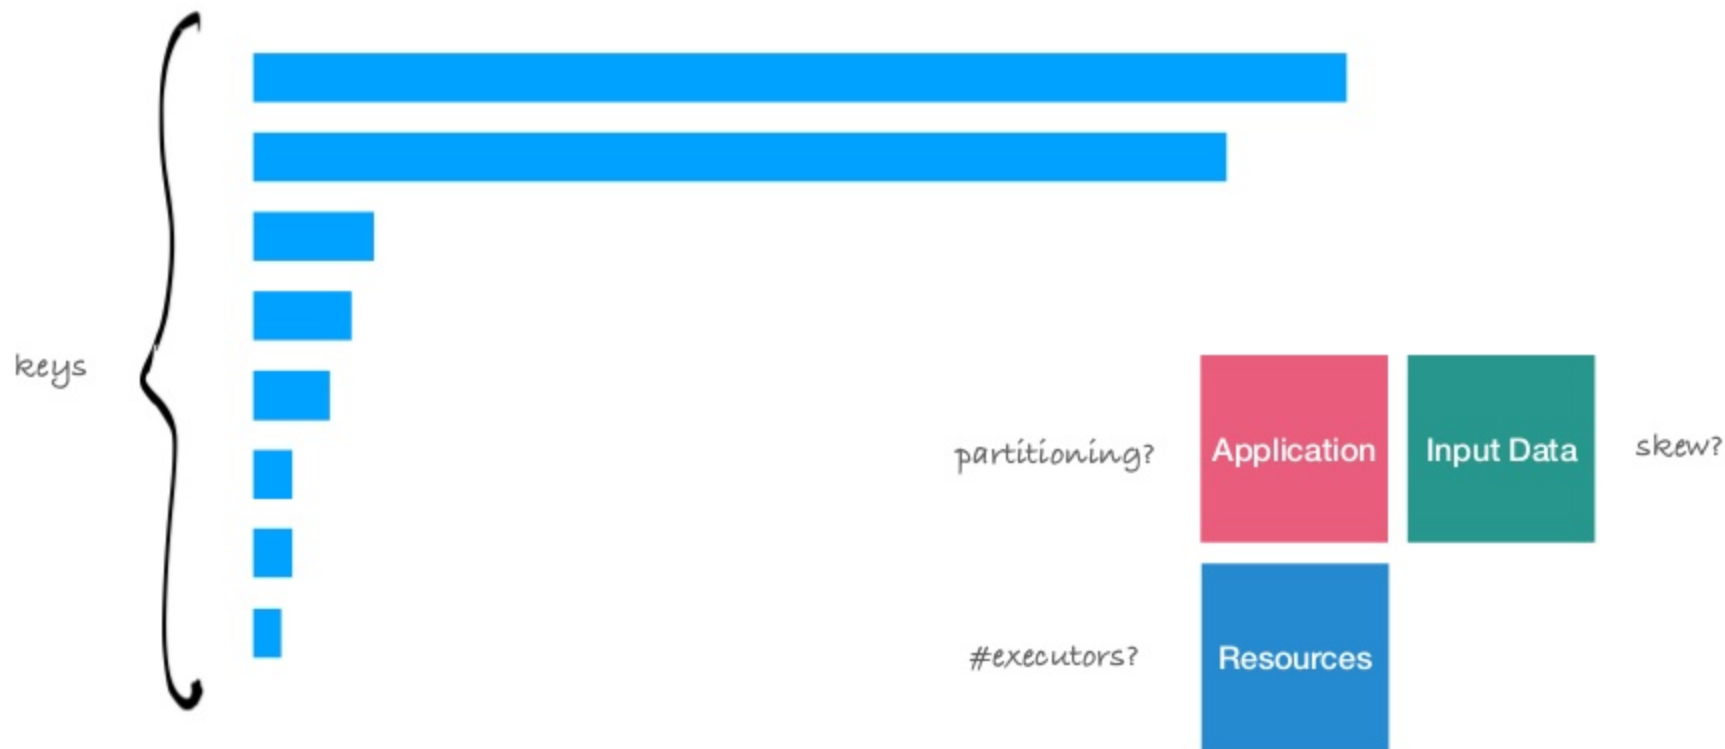
Stragglers

■ Scheduler Delay
■ Task Deserialization Time
■ Shuffle Read Time
■ Executor Computing Time
■ Shuffle Write Time
■ Result Serialization Time
■ Getting Result Time




Index	ID	Attempt	Status	Locality Level	Executor ID / Host	Launch Time	Duration	Scheduler Delay	Shuffle Read Size / Records	Shuffle Remote Reads	Write Time	Shuffle Write Size / Records
11	5236	0	SUCCESS	NODE_LOCAL	0 / 172.18.0.5 stdout stderr	2018/03/23 13:56:17	2.6 min	3 ms	44.5 MB / 535256	0.0 B	10 ms	2.5 MB / 258597
31	5256	0	SUCCESS	NODE_LOCAL	0 / 172.18.0.5 stdout stderr	2018/03/23 13:56:26	2.4 min	4 ms	52.6 MB / 530858	0.0 B	9 ms	2.5 MB / 257064
120	5345	0	SUCCESS	NODE_LOCAL	0 / 172.18.0.5 stdout stderr	2018/03/23 13:57:12	2.4 min	3 ms	39.9 MB / 535418	0.0 B	9 ms	2.5 MB / 259946
20	5245	0	SUCCESS	NODE_LOCAL	0 / 172.18.0.5 stdout stderr	2018/03/23 13:56:21	46 s	5 ms	58.8 MB / 121064	0.0 B	10 ms	213.1 KB / 4905
115	5340	0	SUCCESS	NODE_LOCAL	0 / 172.18.0.5 stdout stderr	2018/03/23 13:57:10	8 s	4 ms	160.1 MB / 26788	0.0 B	50 ms	234.9 KB / 5554
47	5272	0	SUCCESS	NODE_LOCAL	0 / 172.18.0.5 stdout	2018/03/23 13:56:34	7 s	4 ms	37.0 MB / 36130	0.0 B	10 ms	229.8 KB / 5339

Stragglers




Stragglers

Partitioning - worst case



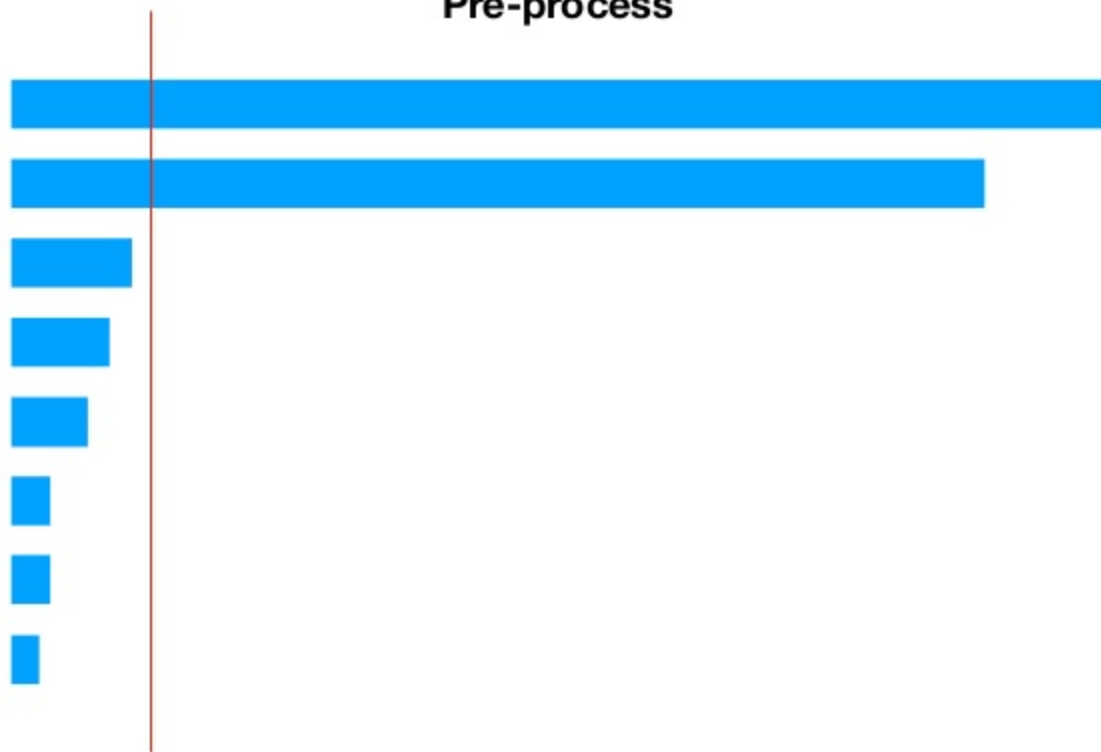
Stragglers

Partitioning - optimal case



Stragglers

Pre-process



Stragglers

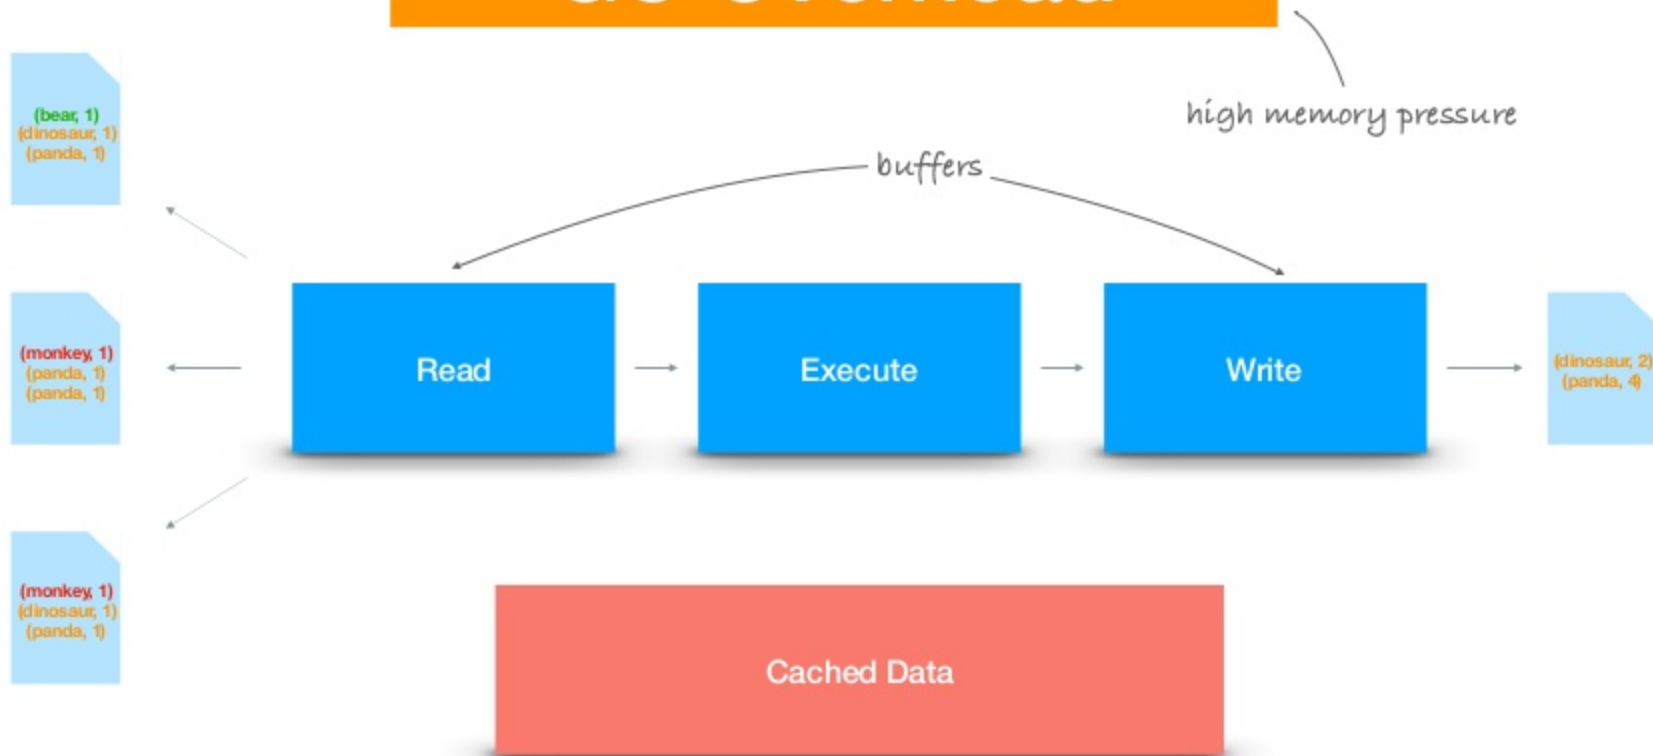
Salting



no idea how large
this will become

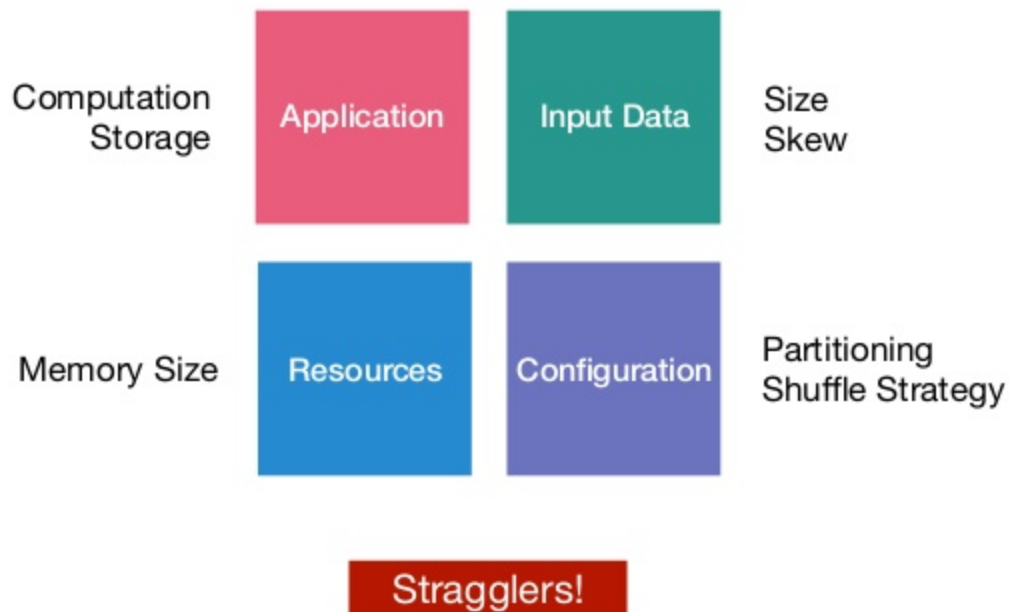
```
case class DuplicateClass(id: Long, docs: List[Document])
```


GC Overhead



GC Overhead

high memory pressure





<https://goo.gl/kRLy1t>

Questions?

Follow me for updates and more resources.



@p_brunenberg



Philipp Brunenberg



www.philipp-brunenberg.de