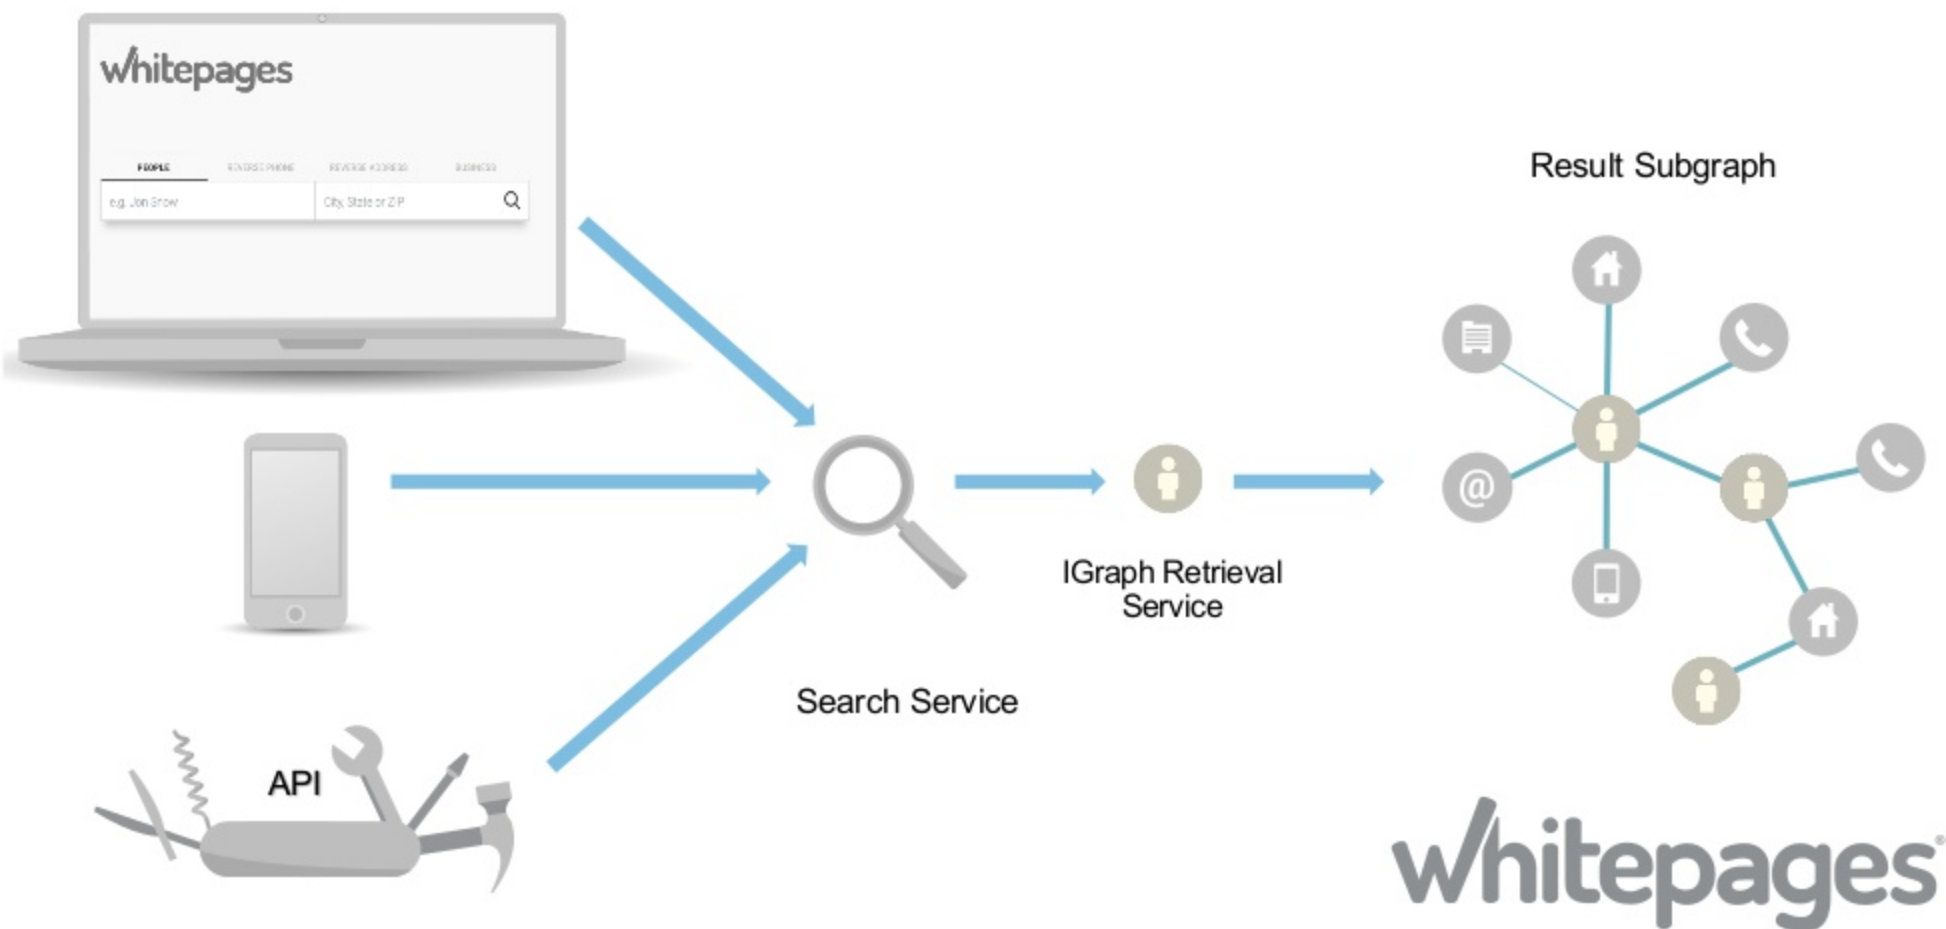


Spark Schema for Free

Dávid Szakállas, Whitepages
@szdavid92

whitepages®

#schema4free



Whitepages Identity Graph™



4B+ entities




6B+ links



whitepages®


Our story in a nutshell



> 30 000 SLOC

3rd PARTY
LIBRARIES

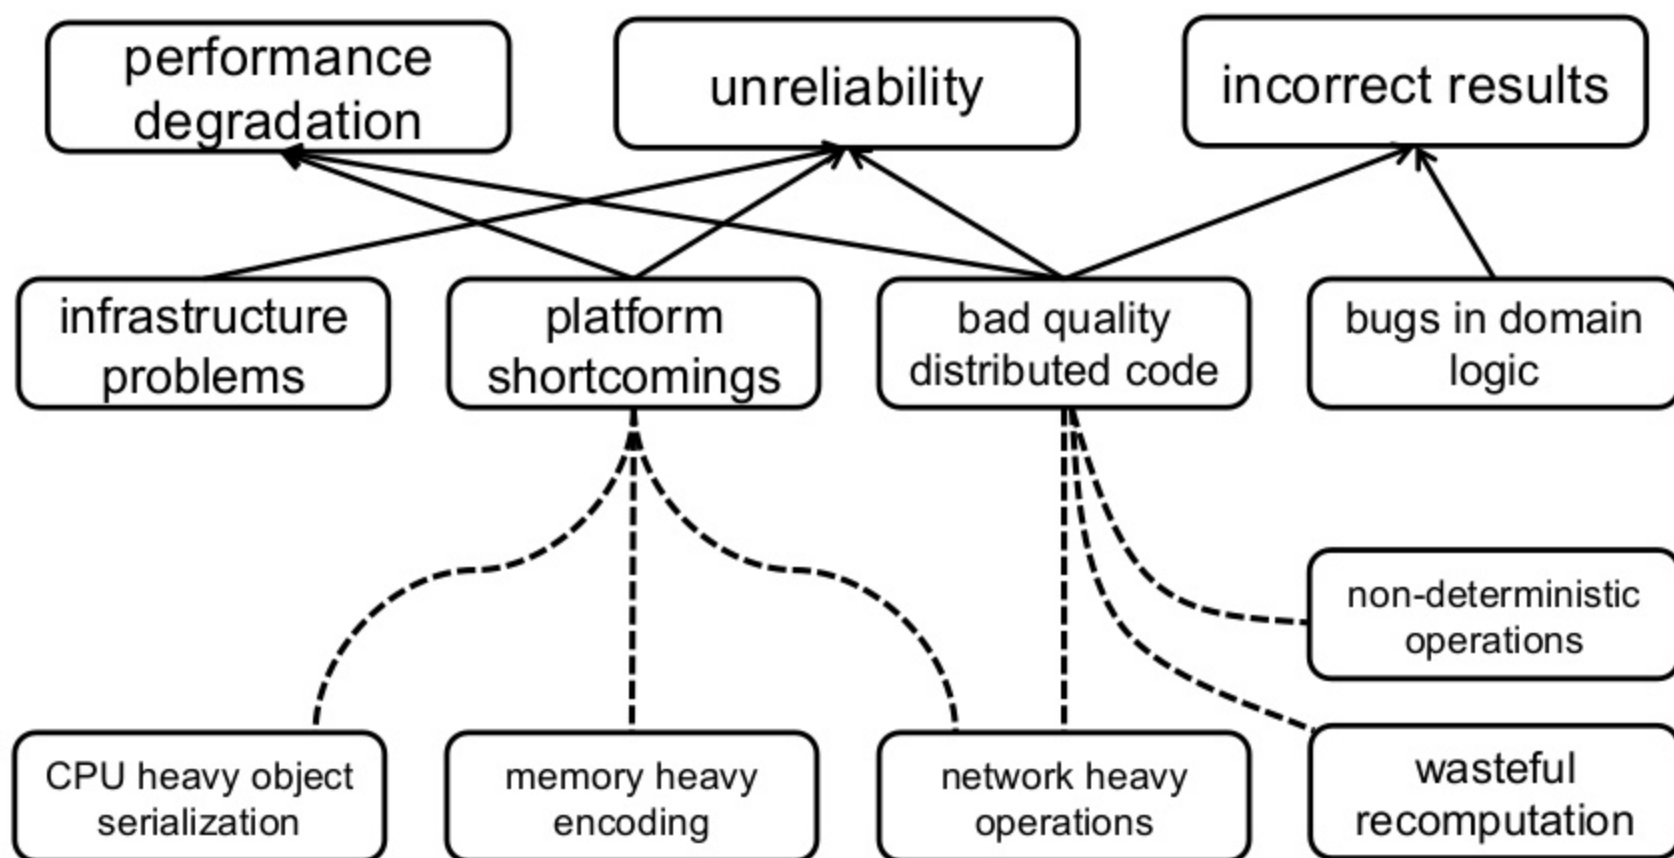
RICH DOMAIN
TYPES



Spark

RDD

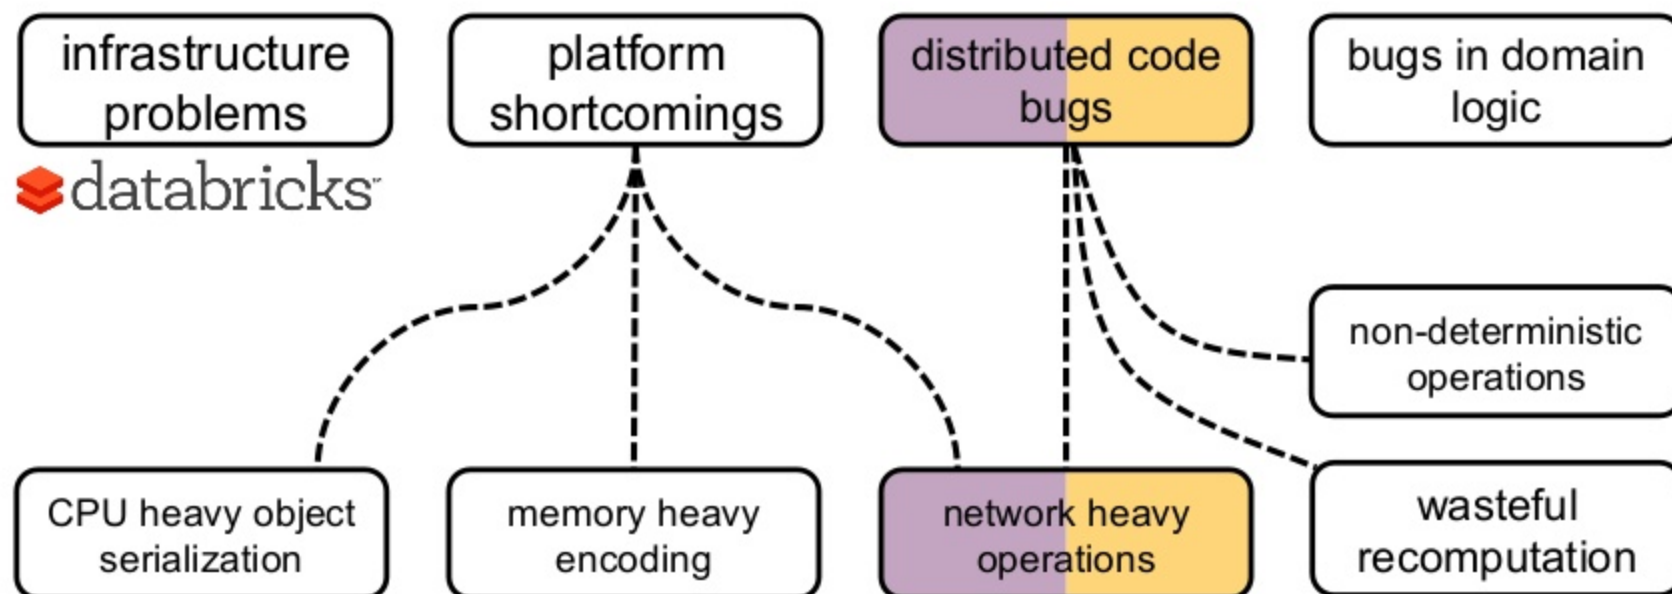
whitepages



Fixing logic and perf bugs

Moving to Databricks










Switching to Spark SQL



Requirements

- Use Scala
- Keep compatibility with existing output format
- Retain compile-time type safety
- Reuse existing domain types
 - ~ 30 core types in case classes
- Leverage the query optimizer
- Minimize memory footprint
 - spare object allocations where possible
- Reduce boilerplate

Two and a half APIs

	Dataset[T]	Dataset[Row]	SQL
	in logic	mistype the column name	syntax error
			
			

: Dataset[T]

S: Serialization cost
D: Deserialization cost
O: Optimization barrier
U: Untyped column
referencing

<code>filter(T => Boolean)</code>	DO
<code>map(T => U)</code>	SDO
<code>mapPartitions(T => U)</code>	SDO
<code>flatMap</code>	SDO
<code>groupByKey</code>	SDO
<code>reduce</code>	SDO
<code>joinWith</code>	U
<code>dropDuplicates</code>	U
<code>orderBy, ...</code>	U

- 1. operations with performance problems**
- 2. operations with type safety problems**
- 3. encoders are not extendable**

```
spark.emptyDataset[java.util.UUID]
```

COMPILE TIME ERROR

error: Unable to find encoder for type stored in a Dataset. Primitive types (Int, String, etc) and Product types (case classes) are supported by importing spark.implicits._ Support for serializing other types will be added in future releases. spark.emptyDataset[java.util.UUID]

```

val jvmRepr = ObjectType(classOf[UUID])

val serializer = CreateNamedStruct(Seq(
  Literal("msb"),
  Invoke(BoundReference(0, jvmRepr, false), "getMostSignificantBits", LongType),
  Literal("lsb"),
  Invoke(BoundReference(0, jvmRepr, false), "getLeastSignificantBits", LongType)
)).flatten

val deserializer = NewInstance(classOf[UUID],
  Seq(GetColumnByOrdinal(0, LongType), GetColumnByOrdinal(1, LongType)),
  ObjectType(classOf[UUID]))

implicit val uuidEncoder = new ExpressionEncoder[UUID](
  schema = StructType(Array(
    StructField("msb", LongType, nullable = false),
    StructField("lsb", LongType, nullable = false)
  )),
  flat = false,
  serializer = serializer,
  deserializer = deserializer,
  clsTag = classTag[UUID])

```

```
spark.emptyDataset[java.util.UUID]
```



```
spark.emptyDataset[(UUID, UUID)]
```

Message: No Encoder found for java.util.UUID

- field (class: "java.util.UUID", name: "_1")
- root class: "scala.Tuple2"

StackTrace: - field (class: "java.util.UUID",

- root class: "scala.Tuple2"

at org.apache.spark.sql.catalyst.ScalaReflection (...)

**RUNTIME
EXCEPTION!**



Creating extendable encoders

- Types are trees of products, sequences, maps with primitive serializable types as leaves
- Problem similar to JSON serialization
- Idea: use generic programming
 - Generate schema, serializers and deserializers
- Type-level programming with shapeless

```

trait ComposableEncoder[T] {
  // ???
}
object ComposableEncoder {
  // derive Spark Encoder
  implicit def getEncoder[T: ComposableEncoder]: Encoder[T] = ???
}

implicit val intEncoder: ComposableEncoder[Int] = ???
implicit val longEncoder: ComposableEncoder[Long] = ???
implicit val uuidEncoder: ComposableEncoder[UUID] = ???
// ...
// other primitive types

// compound types
implicit def productEncoder[G, Repr <: HList]: ComposableEncoder[T] =
  ???
implicit def arrayEncoder[T: ClassTag]: ComposableEncoder[Array[T]] =
  ???
// Option, Either, etc.

```

ComposableEncoder[(UUID, UUID)]



```
implicit val uuidEnc:  
ComposableEncoder[UUID]
```

```
implicit def productEnc[Repr <:  
HList]: ComposableEncoder[T]
```

```
implicit val intEnc:  
ComposableEncoder[Int]
```



ComposableEncoder[UUID]



```
implicit val uuidEnc:  
ComposableEncoder[UUID]
```

```
implicit def productEnc[Repr <:  
HList]: ComposableEncoder[T]
```

```
implicit val intEnc:  
ComposableEncoder[Int]
```

DONE

Expressive types for Spark.

scala

spark

typelevel

fp

functional-programming

📄 562 commits

🌿 5 branches

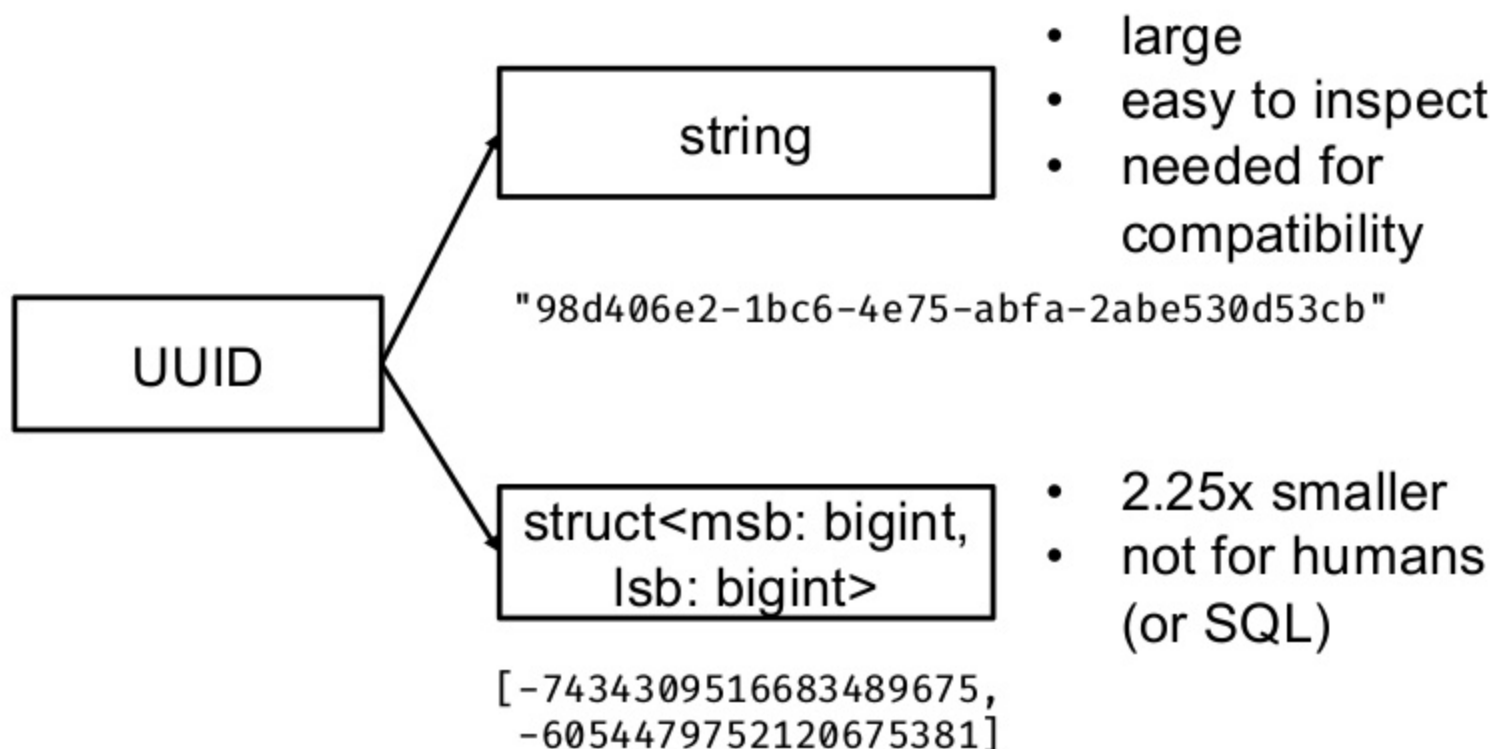
📦 9 releases

👤 32 contributors

📄 Apache-2.0

- Typesafe columns referencing
- Enhanced type signature for built-in functions
- Customizable, type safe encoders
 - Not fully compatible semantics 😞

Multiple encoders



- 1. operations with performance problems**
- 2. operations with type safety problems**
- ~~3. encoders are not extendable~~**

Expressive types for Spark.

scala

spark

typelevel

fp

functional-programming

📄 562 commits

🌿 5 branches

📦 9 releases

👤 32 contributors

📄 Apache-2.0

- Typesafe columns referencing
- Enhanced type signature for built-in functions
- Customizable, type safe encoders
 - Not fully compatible semantics 😞

```
case class Person(id: Long, name: String, gender: String, age: Int)
```

```
spark.read.parquet("data").as[Person]  
  .select($"name", $"age")  
  .filter($"age" >= 18)  
  .filter($"ag" <= 49)
```

RUNTIME ERROR

```
Name: org.apache.spark.sql.AnalysisException  
Message: cannot resolve '`ag`' given input  
columns: (...)
```

```
case class NameAge(name: String, age: Int)
```

```
val tds = TypedDataset.create(spark.read.parquet("data").as[Person])  
val pTds = tds.project[NameAge]  
val fTds = pTds.filter(pTds('age) >= 18)  
               .filter(pTds('ag) <= 49)
```

COMPILE TIME ERROR

```
error: No column Symbol with shapeless.tag.Tagged[String("ag")]  
of type A in NameAge  
      .filter(pTds('ag) <= 49)
```

```

val tds = TypedDataset.create(
  spark.read.parquet("data").as[Person])
val pTds = tds.project[NameAge]
val fTds = pTds.filter(pTds('age) >= 18)
  .filter(pTds('age) <= 49)

```

```

spark.read.parquet("data").as[Person]
  .select($"name", $"age")
  .filter($"age" >= 18)
  .filter($"age" <= 49)

```



```

== Physical Plan ==
*(1) Project
+- *(1) Filter
    +- *(1) FileScan parquet PushedFilters: [IsNotNull(age),
      GreaterThanOrEqual(age,18), LessThanOrEqual(age,49)], ReadSchema:
      struct<name:string,age:int>

```

filter

select

project

join

groupBy

orderBy

withColumn ...

- 1. operations with performance problems**
- ~~2. operations with type safety problems~~**
- ~~3. encoders are not extendable~~**

λs to Catalyst exprs

Compile simple closures to Catalyst expressions
[SPARK-14083]

<code>ds.map(_.name)</code>	<code>// ~ df.select(\$"name")</code>
<code>ds.groupByKey(_.gender)</code>	<code>// ~ df.groupBy(\$"gender")</code>
<code>ds.filter(_.age > 18)</code>	<code>// ~ df.where(\$"age" > 18)</code>
<code>ds.map(_.age + 1)</code>	<code>// ~ df.select(\$"age" + 1)</code>

λs to Catalyst exprs

```
spark.read.parquet("data").as[Person]  
  .map(x => NameAge(x.name, x.age))  
  .filter(_.age >= 18)  
  .filter(_.age <= 49)
```

CURRENTLY



```
== Physical Plan ==  
*(1) SerializeFromObject  
+- *(1) Filter  
    +- *(1) MapElements  
        +- *(1) DeserializeToObject  
            +- *(1) FileScan parquet PushedFilters: [], ReadSchema:  
struct<id:bigint,name:string,gender:string,age:int>
```

λs to Catalyst exprs

```
spark.read.parquet("data").as[Person]
  .map(x => Name(x.name, x.age))
  .filter(_.age >= 18)
  .filter(_.age <= 49)
```

```
spark.read.parquet("data")
  .select($"name", $"age")
  .filter($"age" >= 18)
  .filter($"age" <= 49)
```

[SPARK-14083]



CURRENTLY

```
== Physical Plan ==
*(1) Project
+- *(1) Filter
    +- *(1) FileScan parquet PushedFilters: [IsNotNull(age),
GreaterThanOrEqual(age,18), LessThanOrEqual(age,49)], ReadSchema:
struct<name:string,age:int>
```

Inlining deserialization into λ s

```
ds.filter( age: Int => age == 20 )
```

```
== Physical Plan ==
*(1) SerializeFromObject
+- *(1) Filter
   +- *(1) DeserializeToObject
      +- *(1) FileScan
```

```
root
|-- id: long (nullable = true)
|-- name: string (nullable = true)
|-- gender: string (nullable = true)
|-- age: integer (nullable = true)
```

Analyzed and rewritten on bytecode level



J. Wróblewski, K. Ishizaki, H. Inoue and M. Ohara, "Accelerating Spark DataFrames by Inlining Deserialization,"

DataFrame and Dataset

Enable Lambda Expression to Use Internal Data Format

- Our prototype modifies Java bytecode of **lambda expression** to access internal data format
- We improved performance by avoiding **Ser/De**

```
// ds: DataSet[Array[Int]] -> ds.map(a => Array(a[0]))
```

```
while (true) { ... }  
val array = (DataSet[Array[Int]])  
val array = mapPartitions(array) { (iter, partition) => {  
  iter.toArray().map { a => Array(a[0]) }  
}}
```

Diagram illustrating the transformation from a standard map operation to a lambda expression that directly accesses internal data format, bypassing serialization/deserialization.

SPARK SUMMIT

$\lambda \rightarrow \text{expr}$

- + type safe interface
- + catalyst analyzation & optimization
- + no refactor needed
- only *simple* closures

**NOT
IMPLEMENTED**

param inlining

- + type safe interface
- + no refactor needed
- + more complex closures
- only elides deser cost
- no optimization of body

**NOT
IMPLEMENTED**

Requirements

- Use Scala
- Keep compatibility with existing output format
- Retain compile-time type safety
- Reuse existing domain types
 - ~ 30 core types in case classes
- Leverage the query optimizer
- Minimize memory footprint
 - spare object allocations where possible
- Reduce boilerplate

Miscellaneous features

Strict parsing

- Spark infers input schema by default
- Specify the schema
 - validation
 - spare inference cost
- Schema could be generated
- Parsing extensions

```
import org.apache.spark.sql.types._  
  
import frameless._  
import org.apache.spark.sql.Encoder  
import org.apache.spark.sql.types._  
  
implicit def encoder[T: TypedEncoder] =  
  TypedExpressionEncoder[T]  
  
val persons = spark.read  
  .schema(encoder[Person].schema)  
  .csv("data.csv")  
  .as[Person]
```


UDF + tuples =

```
import org.apache.spark.sql.functions._  
  
val processPerson = udf((p: Person) => p.name)  
  
providerFile  
  .withColumn("results", processPerson($"person"))  
  .collect()
```

[SPARK-12823]

RUNTIME ERROR

```
Caused by: java.lang.ClassCastException:  
org.apache.spark.sql.catalyst.expressions.GenericRowWithSchema cannot be  
cast to Person at (...)  
... 16 more
```

UDF + tuples = 💖

```
import typedudf.TypedUdf
import typedudf.ParamEncoder._

val processPerson = TypedUdf((p : Person) => p.name)

spark.read.parquet("provider").as[ProviderFile]
  .withColumn("results", processPerson($"person"))
  .collect()
```

GitHub: [lesbroot/typedudf](https://github.com/lesbroot/typedudf)

Takeaways

- RDD -> Spark SQL: hard work
 - (with our requirements)
- needed to dig into Catalyst
- compile time overhead
 - Scala 2.12
- check out frameless, etc.

whitepages[®]

Questions

Thank you for listening!



#schema4free
@szdavid92

whitepages[®]
SEARCH. FIND. KNOW.

<https://www.whitepages.com/>