

# Spark-ITS: Indexing for Large-Scale Time Series Data on Spark

Liang Zhang (lzhang6@wpi.edu)  
Data Science Dept.,  
Worcester Polytechnic Institute

#SAISEco5

# Data Science Research Group @ Worcester Polytechnic Institute

Prof. Elke A. Rundensteiner

Prof. Mohamed Y. Eltabakh

Liang Zhang

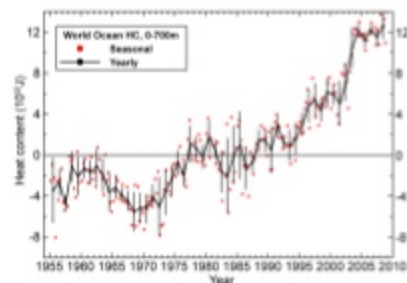
Noura Alghamdi

Liang Zhang, Noura Alghamdi, Mohamed Y. Eltabakh, Elke A. Rundensteiner. *TARDIS: Distributed Indexing Framework for Big Time Series Data*. Proceedings of 35th IEEE International Conference on Data Engineering **ICDE**, 2019

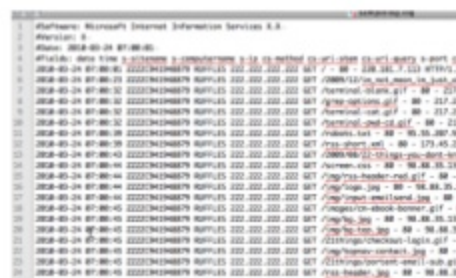
# Outline

- Motivation
- Background
- Spark-ITS Framework
  - Overview
  - Index Construction
  - Query Processing
- Performance Evaluation

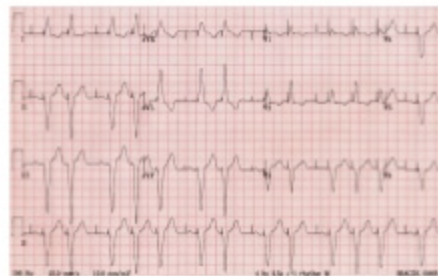
# Time Series are Continuously Produced Everywhere



Climate data



Web log



EEG

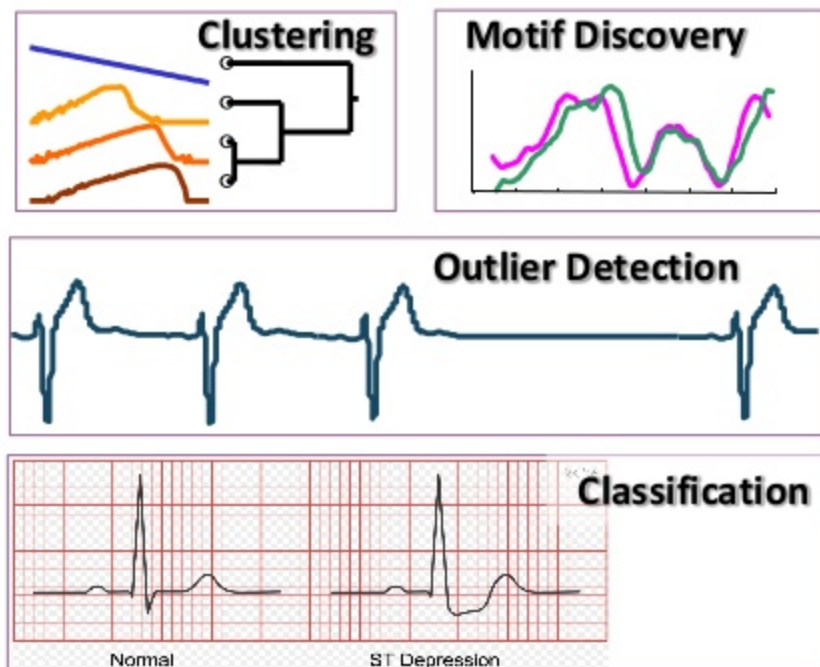


Stock price

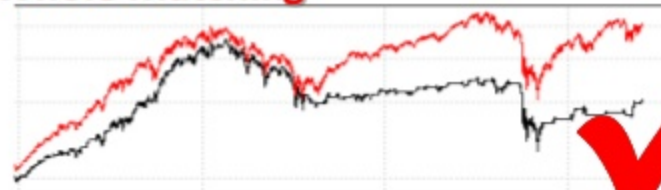
- How to deal with **billions** of time series?



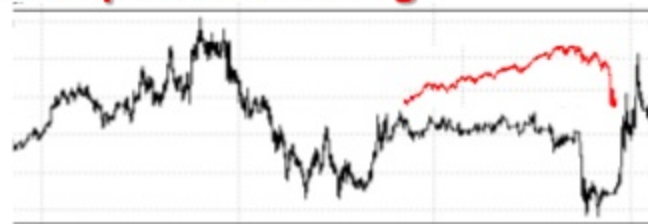
# Almost all Time Series Data Mining Tasks rely on **Similarity Query**



**Whole Matching**



**Subsequence Matching**



Esling, Philippe, and Carlos Agon. "Time-series data mining." *ACM (CSUR)* 45.1 (2012): 12.

# Spark-ITS

- A new **Index Tree** and an effective **Signature** to simplify the cardinality conversion and keep better similarity
- A **Distributed Index Framework** to support **large-scale** time series dataset
- Efficient algorithms for **Exact Match** and **kNN Approximate** queries process



# Spark-ITS Overview

1. Sampling
2. Node Statistic
3. Build Index Tree
4. Assign Partition ID

Global Index

Query

Local Index

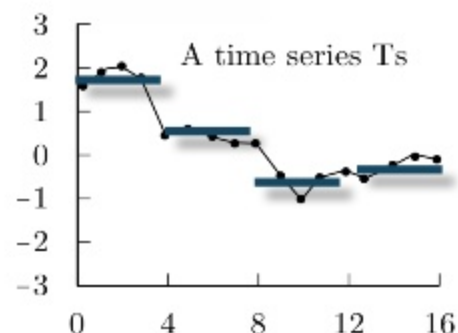
Partition

1. Read and convert data
2. Shuffle data

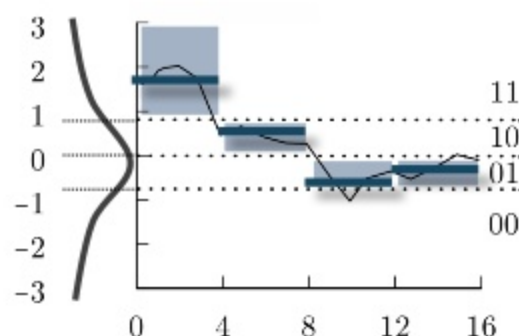
1. Construct Local Structure
2. Construct Bloom Filter

Indexed Data

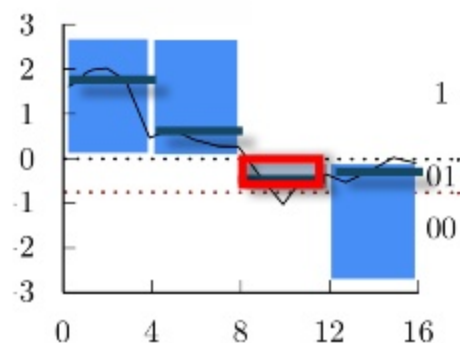
# Background: iSAX Representation



A time series of length 16  
PAA representation with 4 segments



SAX representation with  
4 segments and  
**cardinality** 4  
[11,10,01,00]



iSAX representation  
with 4 segments and  
**variable cardinality**  
[1<sub>2</sub>, 1<sub>2</sub>, **01**<sub>4</sub>, 0<sub>2</sub>]

PAA: Piecewise Aggregate Approximation

iSAX: indexable Symbolic Aggregate approximation

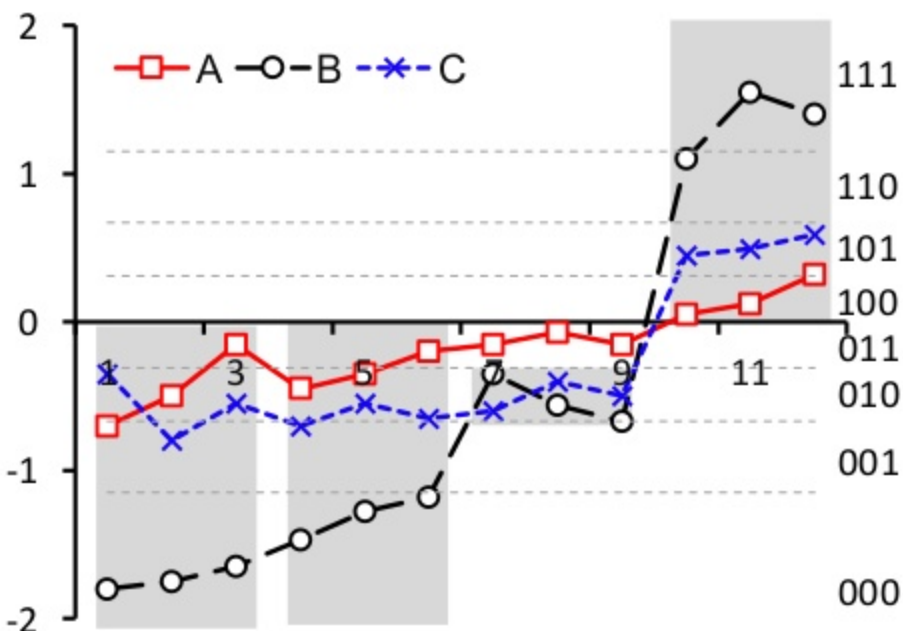
Shieh, Jin, and Eamonn Keogh. "iSAX: indexing and mining terabyte sized time series." *SIGKDD ACM*, 2008.

Camerra, A., Palpanas, T., Shieh, J., & Keogh, E. "iSAX 2.0: Indexing and mining one billion time series." *ICDM*, 2010



# Word-level Similarity

State-of-the-art: Character-level Similarity



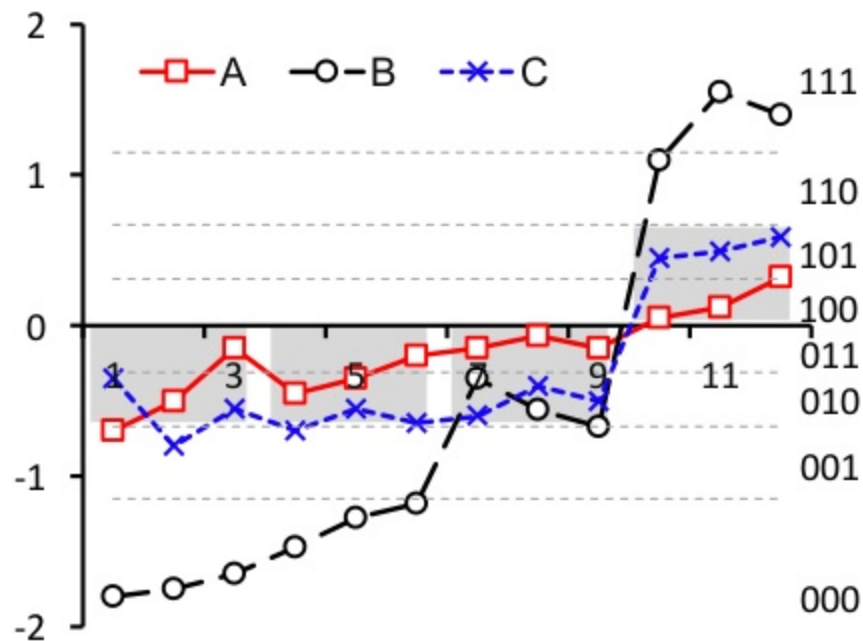
**A:**  $[0_1, 0_1, 011_3, 1_1]$

**B:**  $[0_1, 0_1, 010_3, 1_1]$

**C:**  $[0_1, 0_1, 010_3, 1_1]$

**B** and **C**  
are similar

Proposed: Word-level Similarity



**A:**  $[01_2, 01_2, 01_2, 10_2]$

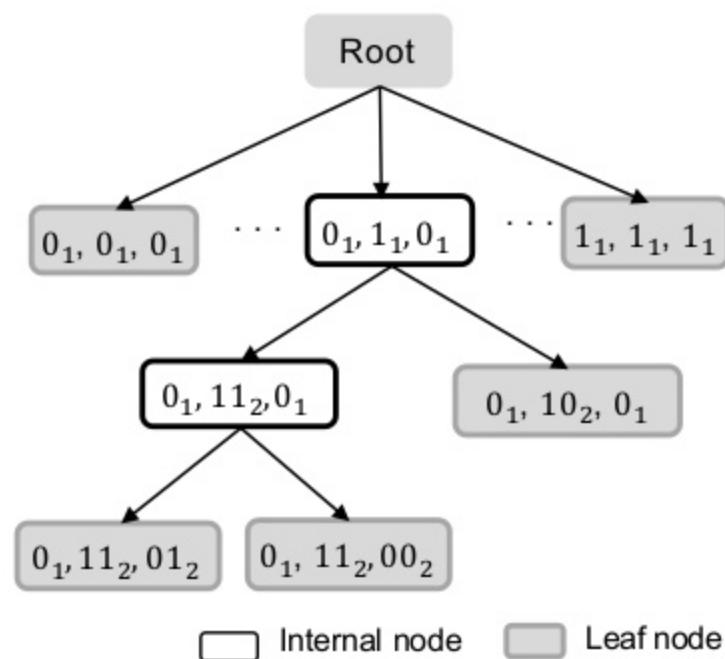
**B:**  $[00_2, 00_2, 01_2, 11_2]$

**C:**  $[01_2, 01_2, 01_2, 10_2]$

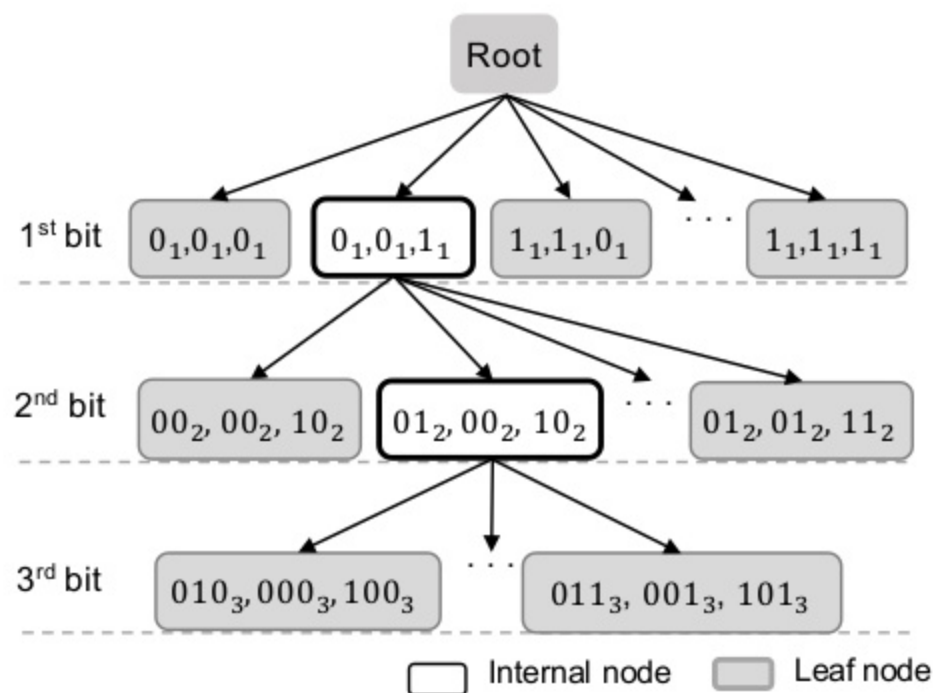
**A** and **C**  
are similar

# New Index Tree Supports Word-level Similarity

**State-of-the-art:** iSAX Binary Tree



**Proposed:** iSAX-T K-ary Tree



# iSAX-T(Transpose) Signature

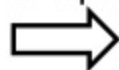
Time series:

[1100, 1101, 0110, 0001]



$$\begin{bmatrix} 1 & 1 & 0 & 0 \\ 1 & 1 & 0 & 1 \\ 0 & 1 & 1 & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix}$$

Transpose



$$\begin{bmatrix} 1 & 1 & 0 & 0 \\ 1 & 1 & 1 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 1 & 0 & 1 \end{bmatrix}$$

Hex



$$\begin{bmatrix} C \\ E \\ 2 \\ 5 \end{bmatrix}$$

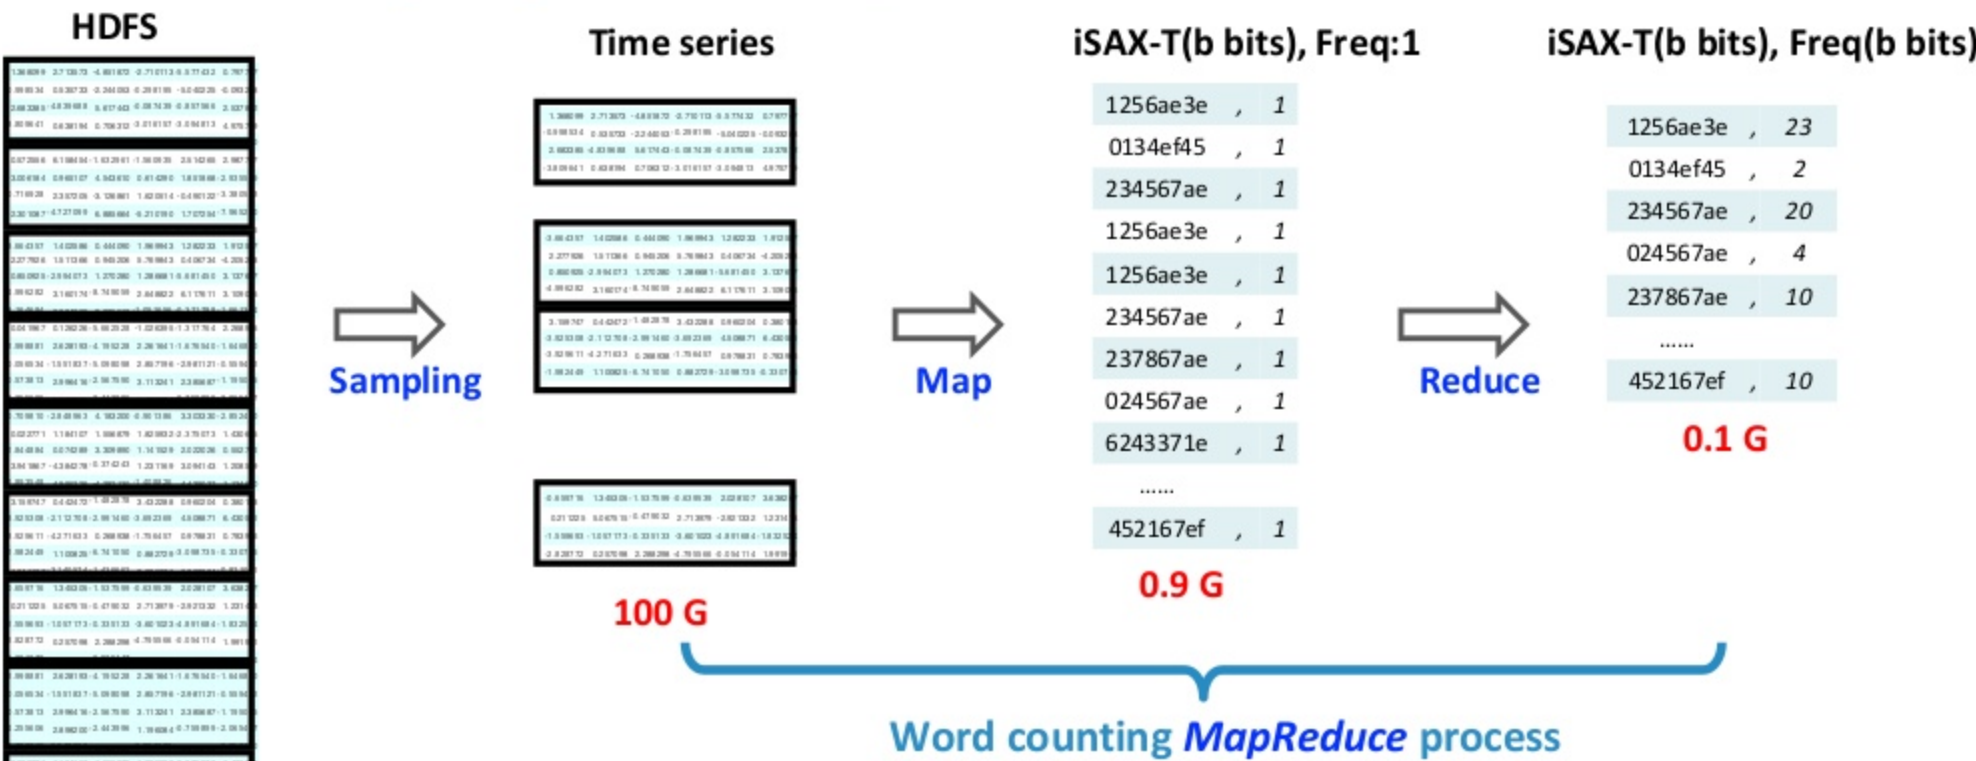


		iSAX-T
SAX(T,4,16) =	{1100, 1101, 0110, 0001}	= CE25
SAX(T,4,8) =	{110, 110, 011, 000}	= CE2
SAX(T,4,4) =	{11, 11, 01, 00}	= CE
SAX(T,4,2) =	{1, 1, 0, 0}	= C

# Outline

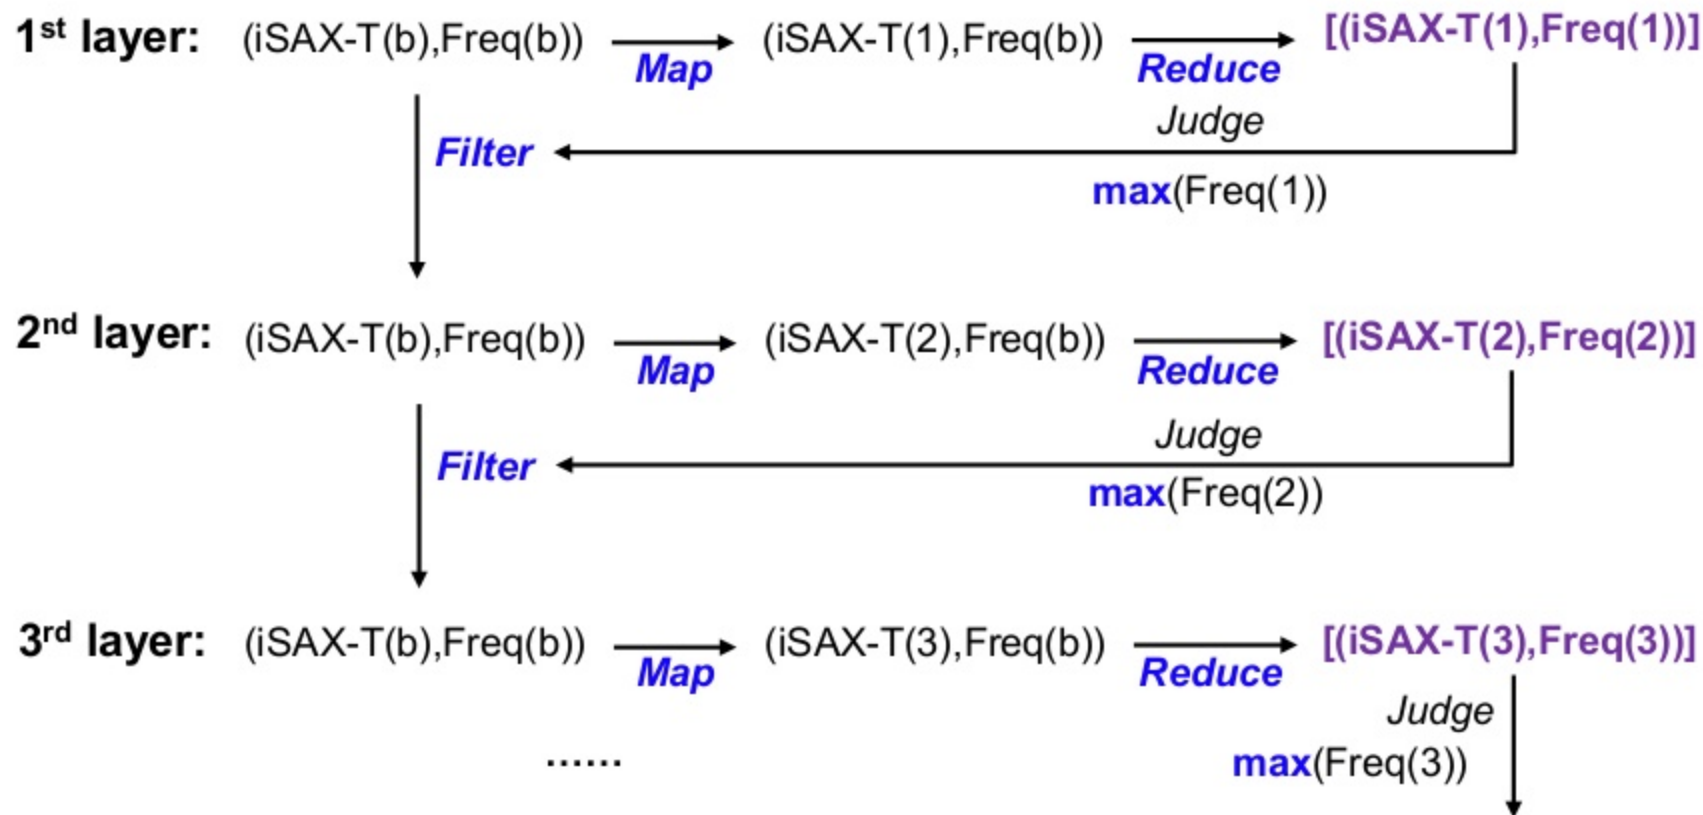
- Motivation
- Background
- Spark-ITS Framework
  - Overview
  - Index Construction
  - Query Processing
- Performance Evaluation

# Global Index[1/4]: Sampling



Segment Number: 8, so use 2 letters to represent 1 bit  
Initial cardinality: **b bit level**  
The data size is based on 1 billion time series with 256 length

## Global Index[2/4]: Node Statistic





# Global Index[3/4]: Build Tree

## 1<sup>st</sup> layer (iSAX-T, Freq)

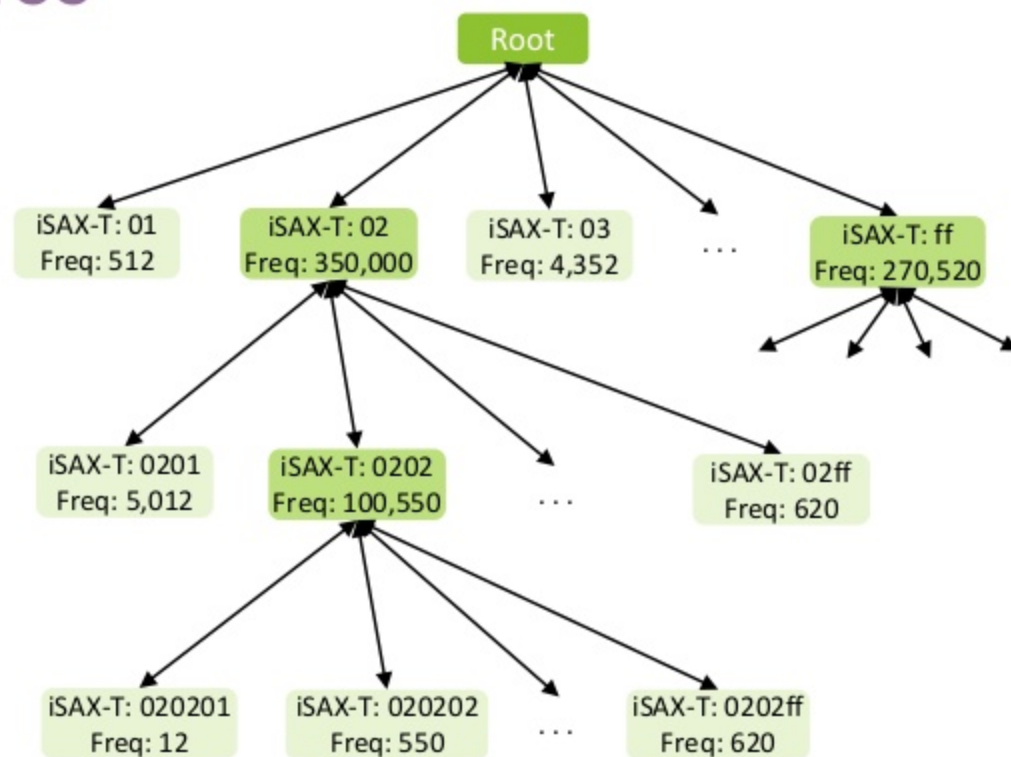
- ("01", 512)
- ("02", 355,000)
- ....
- ("ff", 270,520)

## 2<sup>nd</sup> layer (iSAX-T, Freq)

- ("0201", 5,012)
- ("0202", 100,550)
- ....
- ("ffff", 10,520)

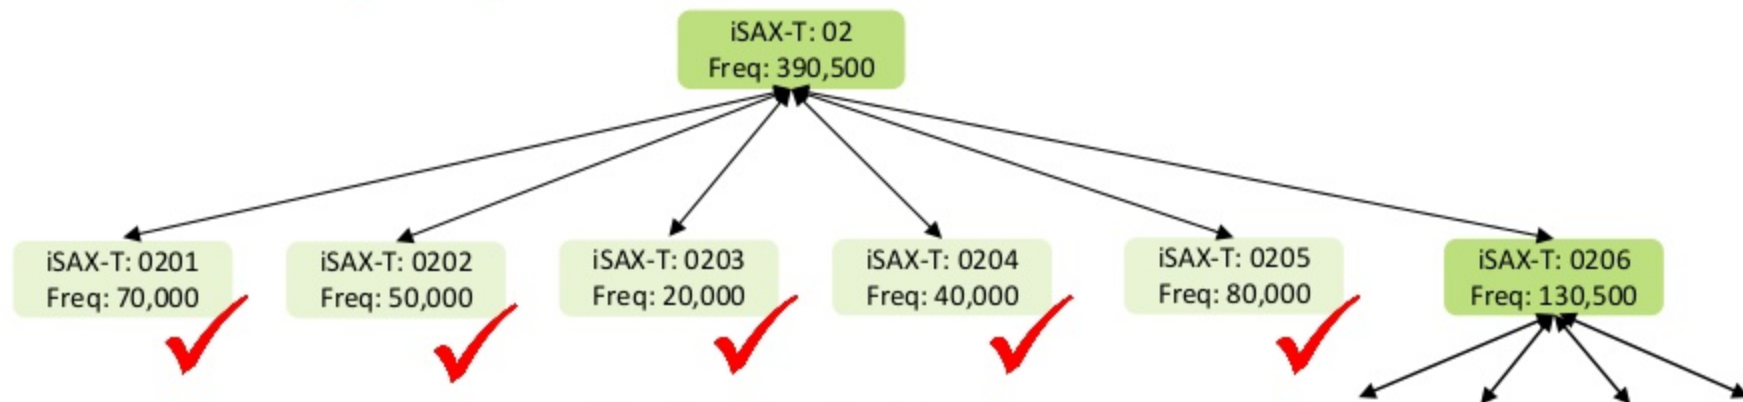
## 3<sup>rd</sup> layer (iSAX-T, Freq)

- ("020201", 12)
- ("020202", 550)
- ....
- ("0202ff", 620)



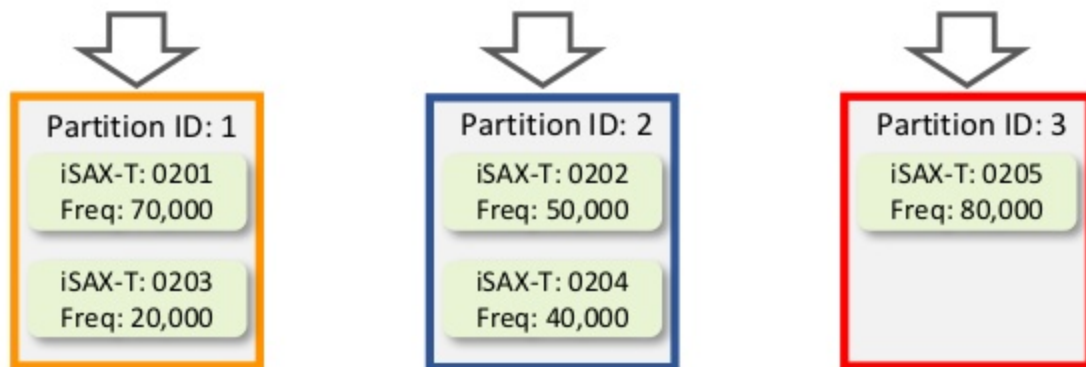
Segment number: 8  
Partition Capacity: 100,000

## Global Index[4/4]: Assign Partition Id to Leaf Nodes



### Bin Packing Problem:

How to fit a set of nodes in the smallest numbers of partitions?



Partition capacity: 100,000

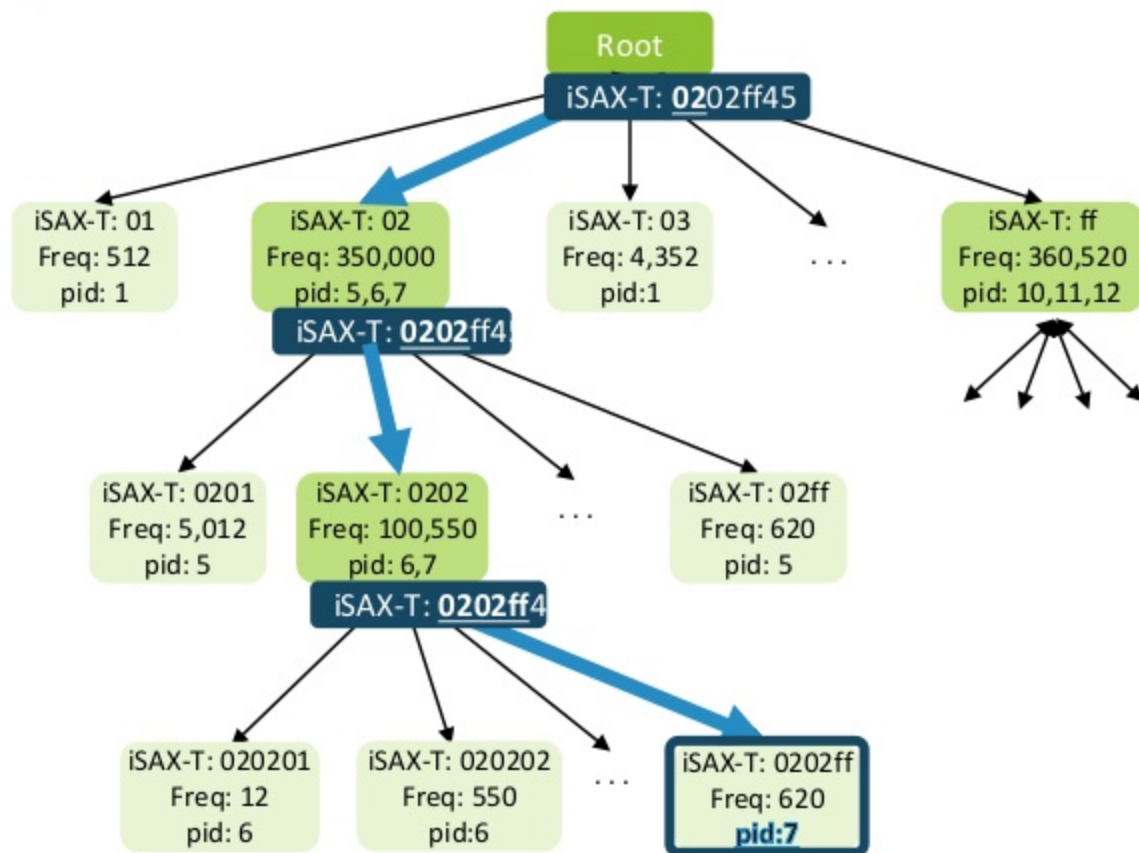
# Repartition: Wrap Global Index as the Partitioner

## A Time Series

iSAX-T: 0202ff45  
TS: [0.34, 0.31, 1.14...]



pid:7



# Local Index: Construction Within Each Partition

## Time series in one partition

1.368 008 2.713073 -0.803875 -0.710113 0.077422 0.787787

0.000000 0.000000 0.000000 0.000000 0.000000 0.000000 0.000000

3.801801 0.000000 0.000000 0.000000 0.000000 0.000000 0.000000

0.000000 1.000000 0.000000 0.000000 0.000000 0.000000 0.000000

0.000000 0.000000 0.000000 0.000000 0.000000 0.000000 0.000000

0.000000 0.000000 0.000000 0.000000 0.000000 0.000000 0.000000

0.000000 0.000000 0.000000 0.000000 0.000000 0.000000 0.000000

0.000000 0.000000 0.000000 0.000000 0.000000 0.000000 0.000000

0.000000 0.000000 0.000000 0.000000 0.000000 0.000000 0.000000

0.000000 0.000000 0.000000 0.000000 0.000000 0.000000 0.000000

0.000000 0.000000 0.000000 0.000000 0.000000 0.000000 0.000000

0.000000 0.000000 0.000000 0.000000 0.000000 0.000000 0.000000

0.000000 0.000000 0.000000 0.000000 0.000000 0.000000 0.000000

0.000000 0.000000 0.000000 0.000000 0.000000 0.000000 0.000000

0.000000 0.000000 0.000000 0.000000 0.000000 0.000000 0.000000

0.000000 0.000000 0.000000 0.000000 0.000000 0.000000 0.000000

0.000000 0.000000 0.000000 0.000000 0.000000 0.000000 0.000000

0.000000 0.000000 0.000000 0.000000 0.000000 0.000000 0.000000

0.000000 0.000000 0.000000 0.000000 0.000000 0.000000 0.000000

0.000000 0.000000 0.000000 0.000000 0.000000 0.000000 0.000000

0.000000 0.000000 0.000000 0.000000 0.000000 0.000000 0.000000

0.000000 0.000000 0.000000 0.000000 0.000000 0.000000 0.000000

0.000000 0.000000 0.000000 0.000000 0.000000 0.000000 0.000000

0.000000 0.000000 0.000000 0.000000 0.000000 0.000000 0.000000

0.000000 0.000000 0.000000 0.000000 0.000000 0.000000 0.000000

0.000000 0.000000 0.000000 0.000000 0.000000 0.000000 0.000000

0.000000 0.000000 0.000000 0.000000 0.000000 0.000000 0.000000

0.000000 0.000000 0.000000 0.000000 0.000000 0.000000 0.000000

0.000000 0.000000 0.000000 0.000000 0.000000 0.000000 0.000000

0.000000 0.000000 0.000000 0.000000 0.000000 0.000000 0.000000

0.000000 0.000000 0.000000 0.000000 0.000000 0.000000 0.000000

0.000000 0.000000 0.000000 0.000000 0.000000 0.000000 0.000000

0.000000 0.000000 0.000000 0.000000 0.000000 0.000000 0.000000

0.000000 0.000000 0.000000 0.000000 0.000000 0.000000 0.000000

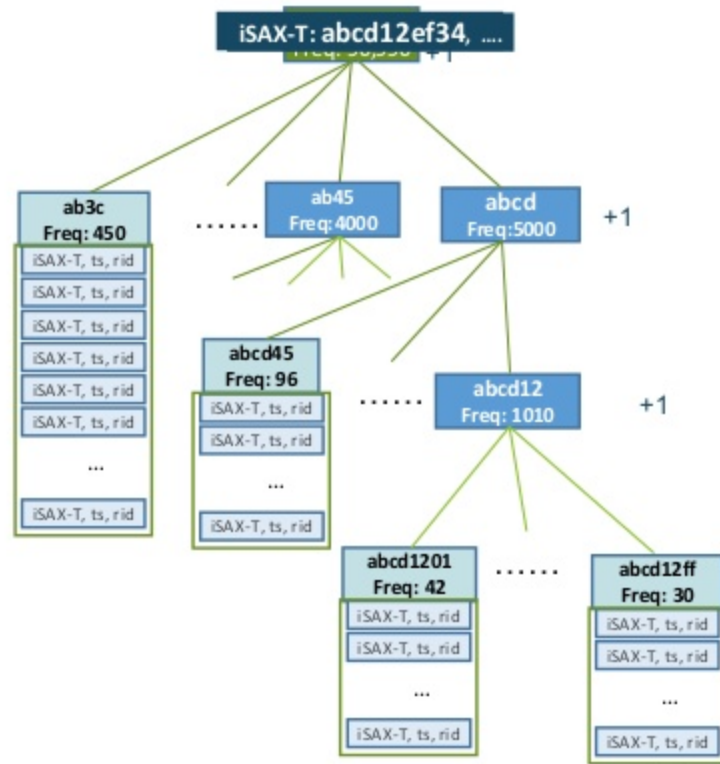
0.000000 0.000000 0.000000 0.000000 0.000000 0.000000 0.000000

0.000000 0.000000 0.000000 0.000000 0.000000 0.000000 0.000000

0.000000 0.000000 0.000000 0.000000 0.000000 0.000000 0.000000

0.000000 0.000000 0.000000 0.000000 0.000000 0.000000 0.000000

## Local Index



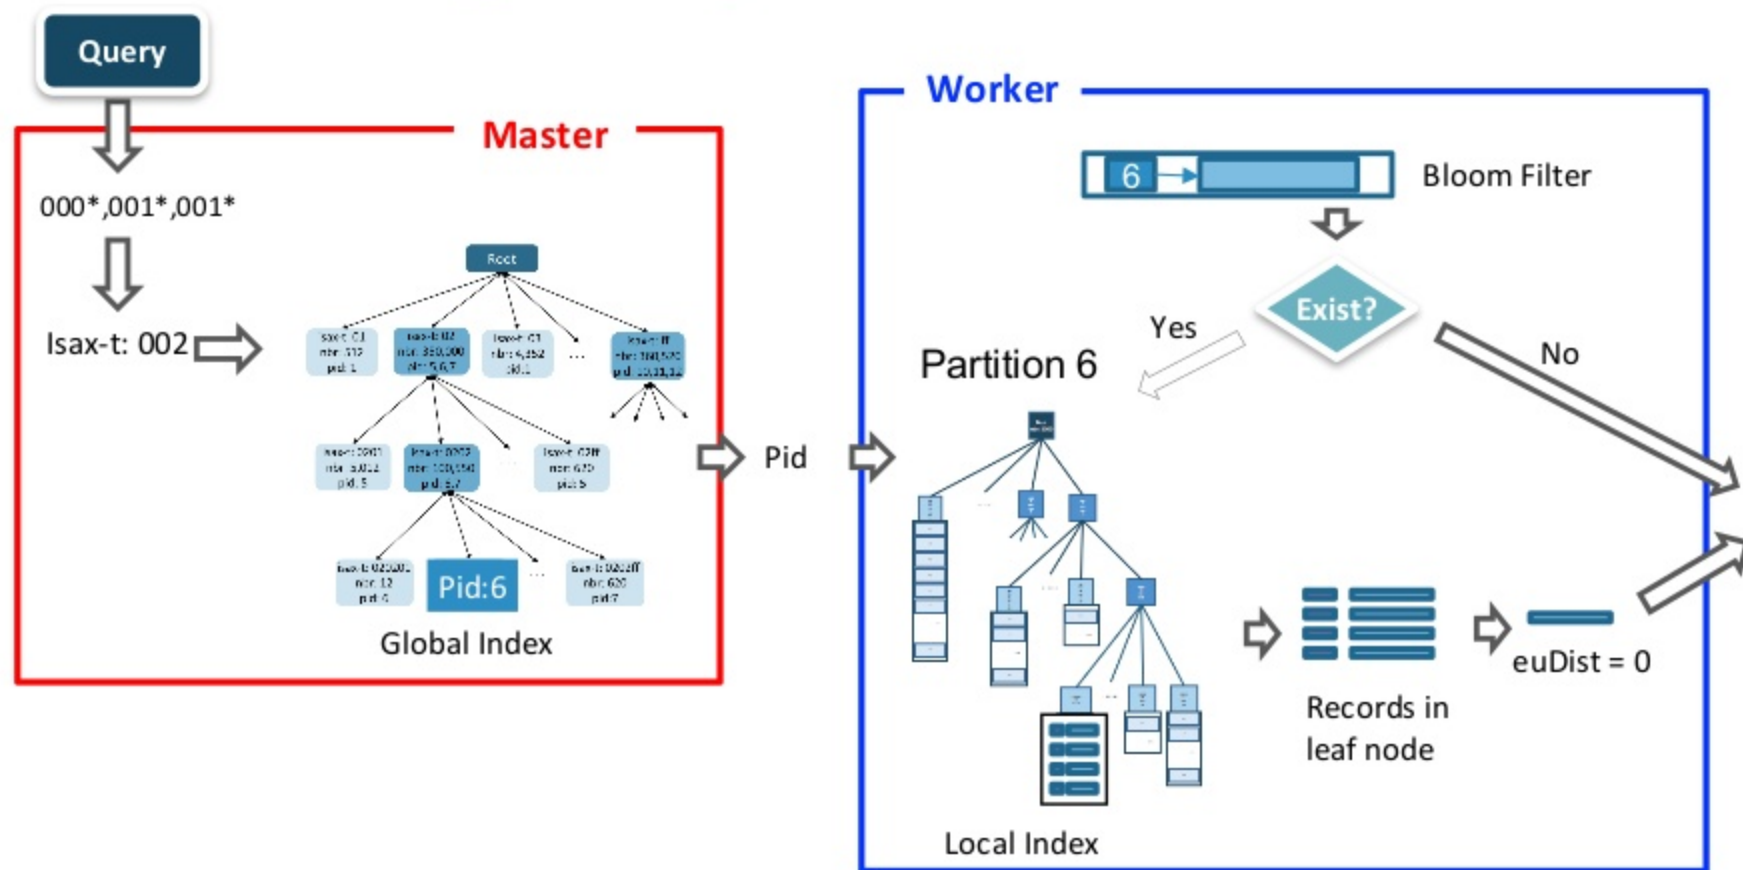
## Bloom Filter

Partition capacity: 100,000  
Node split threshold: 1000  
Segment Number: 8

# Outline

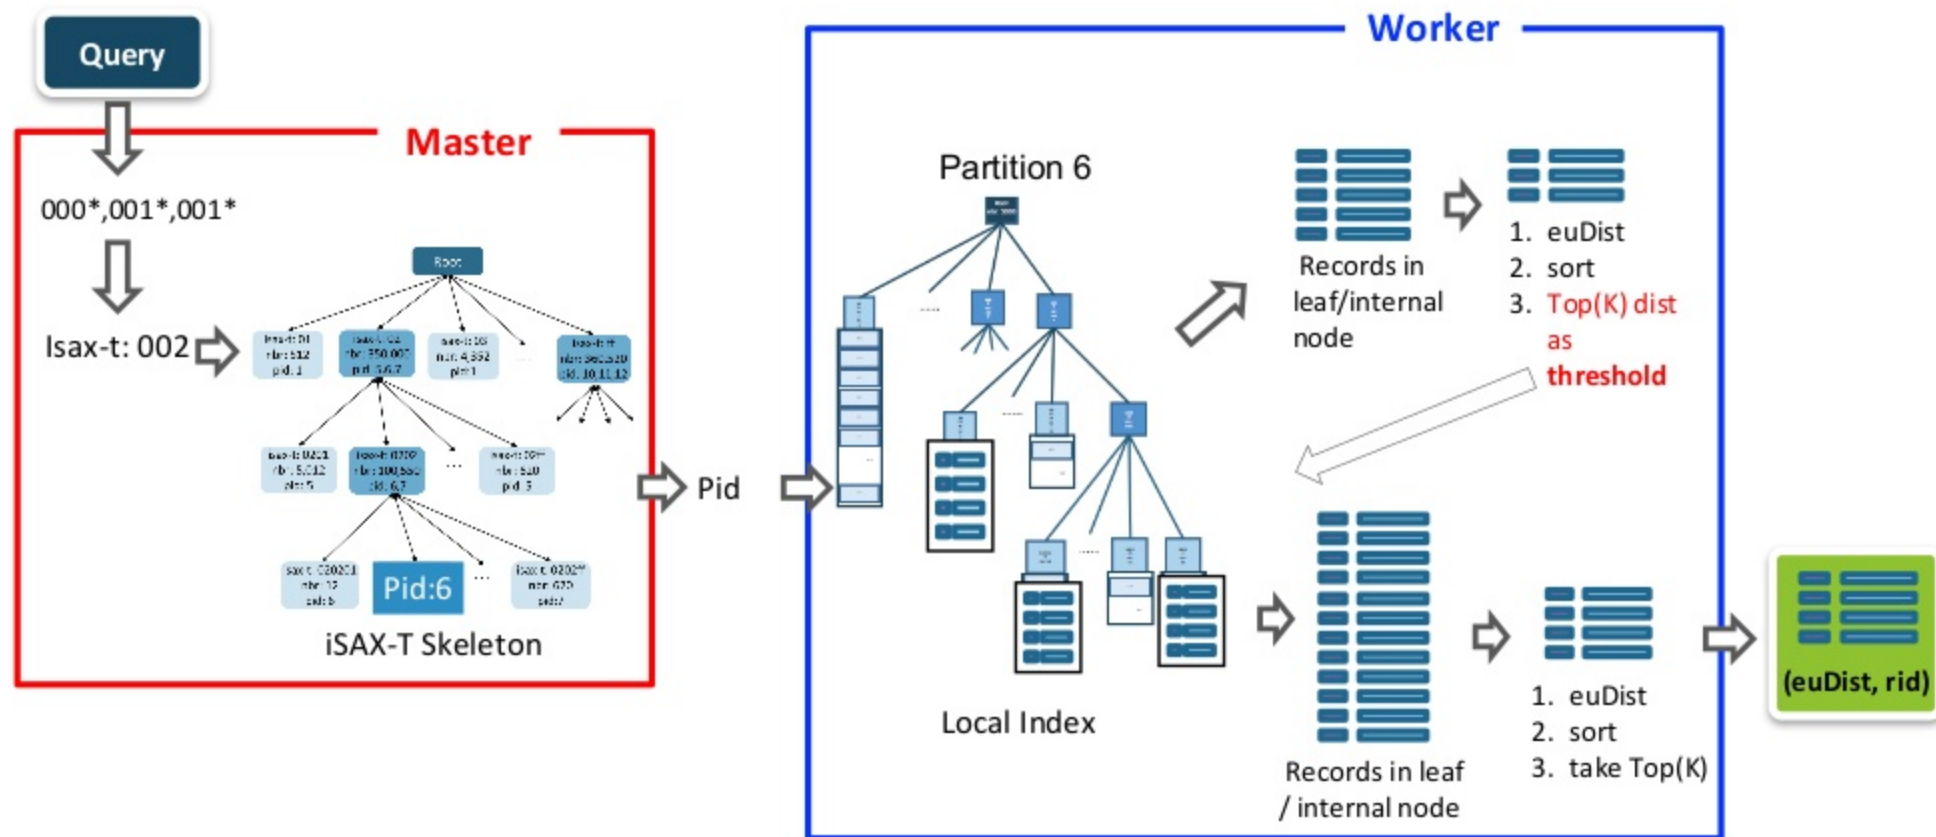
- Motivation
- Background
- Spark-ITS Framework
  - Overview
  - Index Construction
  - Query Processing
- Performance Evaluation

# Exact Matching Query

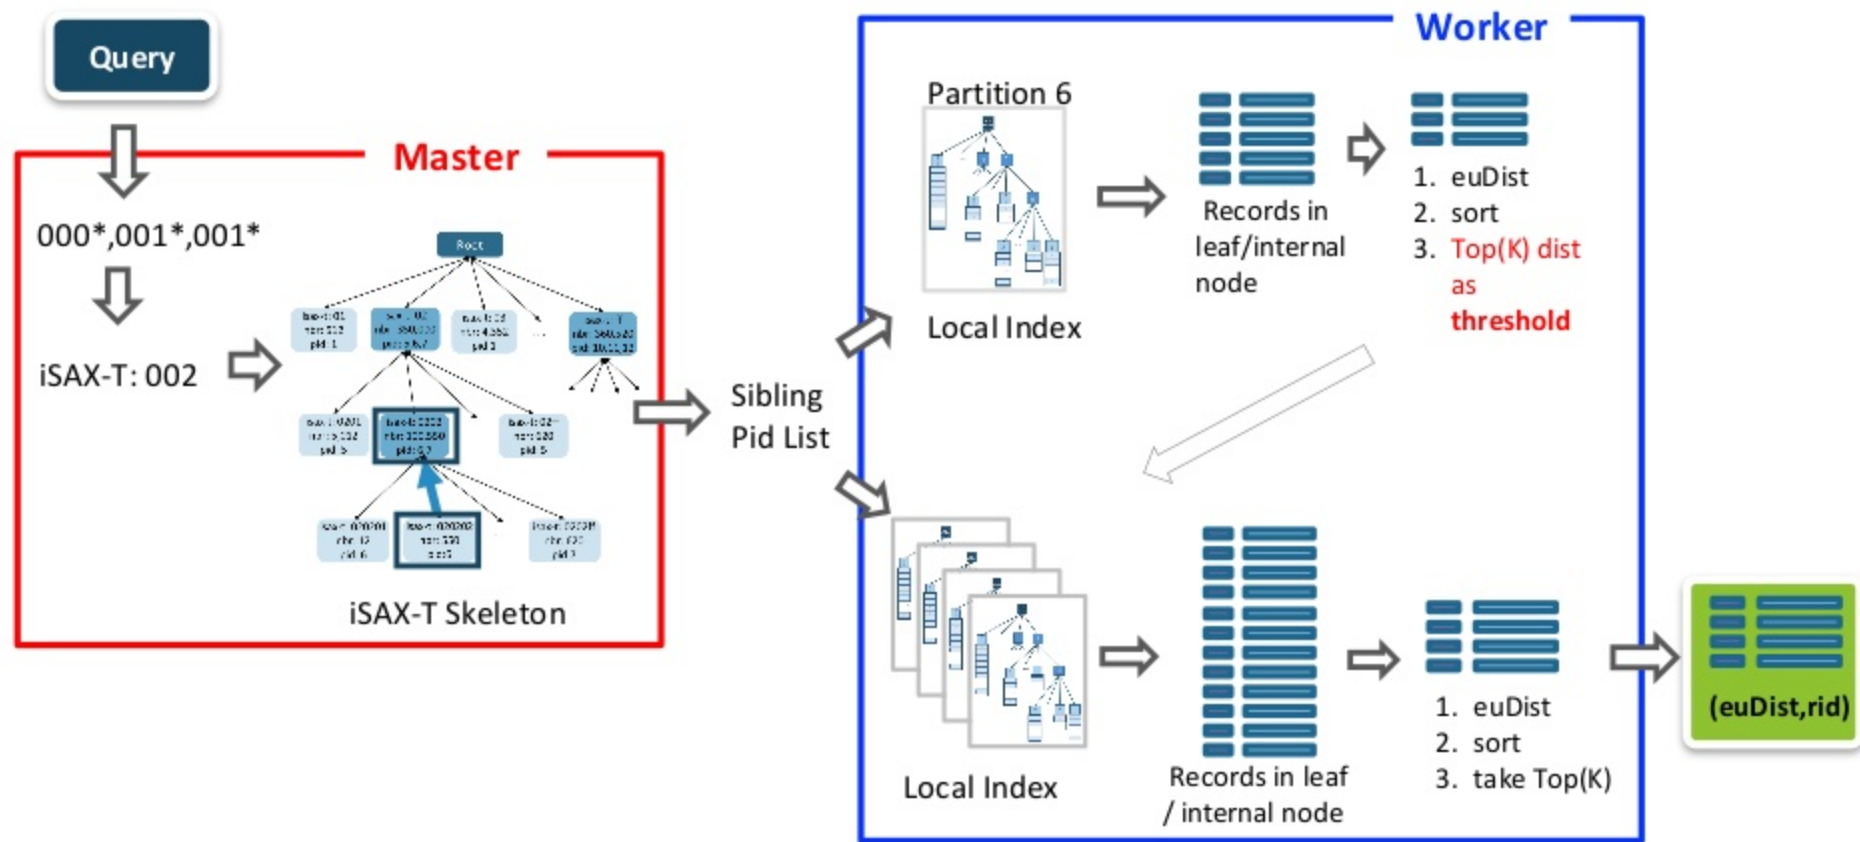




# KNN Approximate Query: One Partition Access



# KNN Approximate Query: Multi-Partitions Access



# Outline

- Motivation
- Background
- Spark-ITS Framework
  - Overview
  - Index Construction
  - Query Processing
- Performance Evaluation

# Experimental Setup

HW&SW	Configuration
Spark	2.0.2, Standalone mode
Hadoop	2.7.3
Platform	Ubuntu 16.04. LTS
HW	2 nodes, each node consist of 56 Xeon E5 processors, 500G RAM, 7TB SATA hard drive

Dataset	Size	Length
Random Walk	1 billion	256
Texmex <sup>1</sup>	1 billion	128
DNA <sup>2</sup>	200 million	192
Noaa Climate <sup>3</sup>	200 million	64

	Baseline	Spark-ITS
Initial cardinality	512	64
Word length	8	8
Sampling percent	10%	10%
Leaf node split threshold of Local index	1000	1000

**State-of-the-Art:** Yagoubi, Djamel-Edine, et al. "DPiSAX: Massively Distributed Partitioned iSAX." *ICDM 2017*

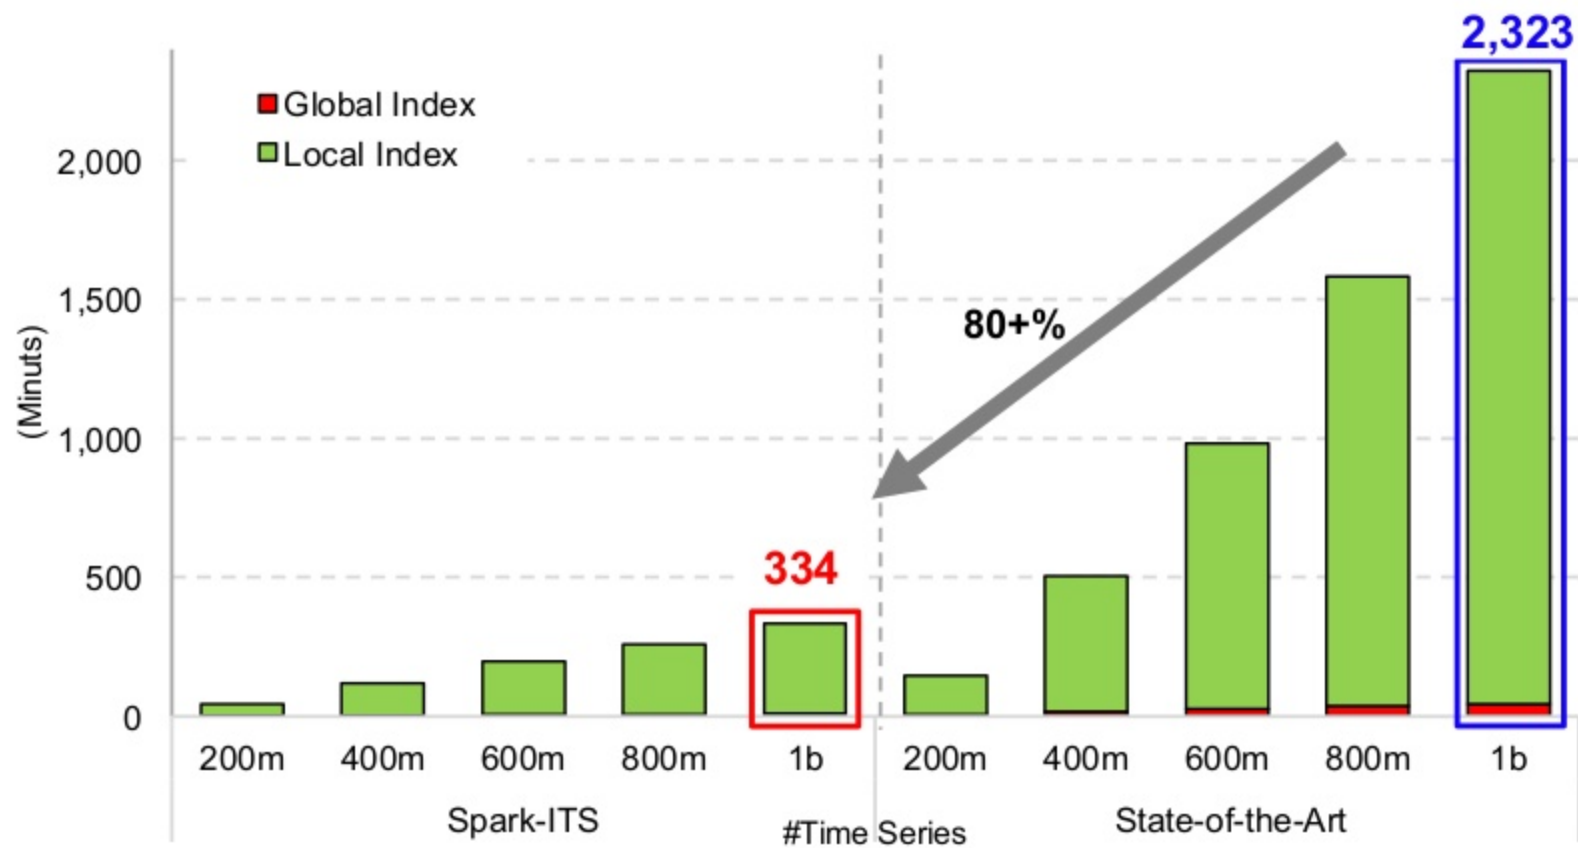
The initial cardinality of the baseline system is the default value and it needs a large initial value to guarantee enough bit level for binary split.

The dataset is normalized  
Each point is saved as float format

Source:

1. <http://corpus-texmex.irisa.fr/>
2. <https://genmone.ucsc.edu>
3. <https://www.ncdc.gov/>

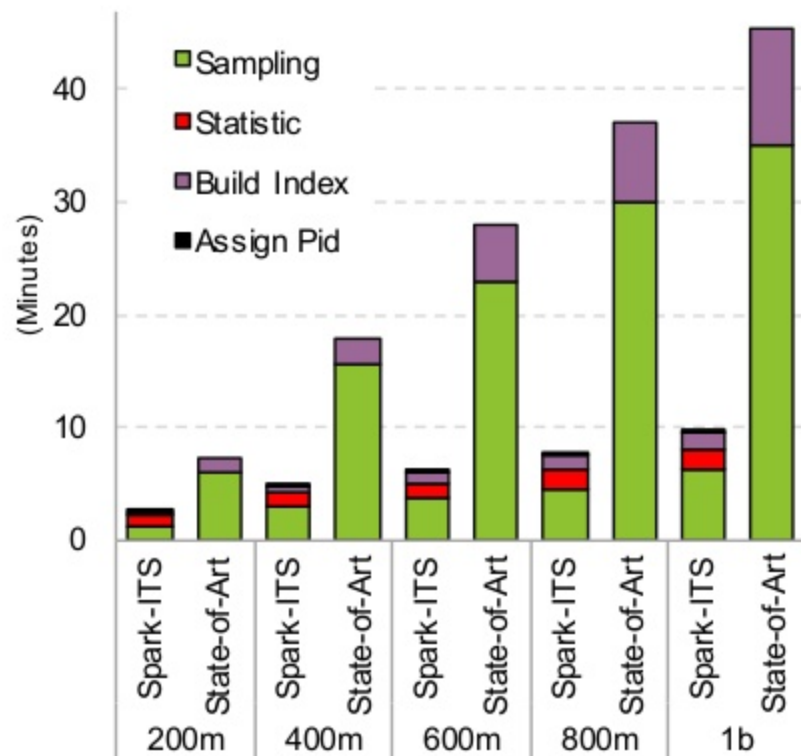
# Index Construction Time



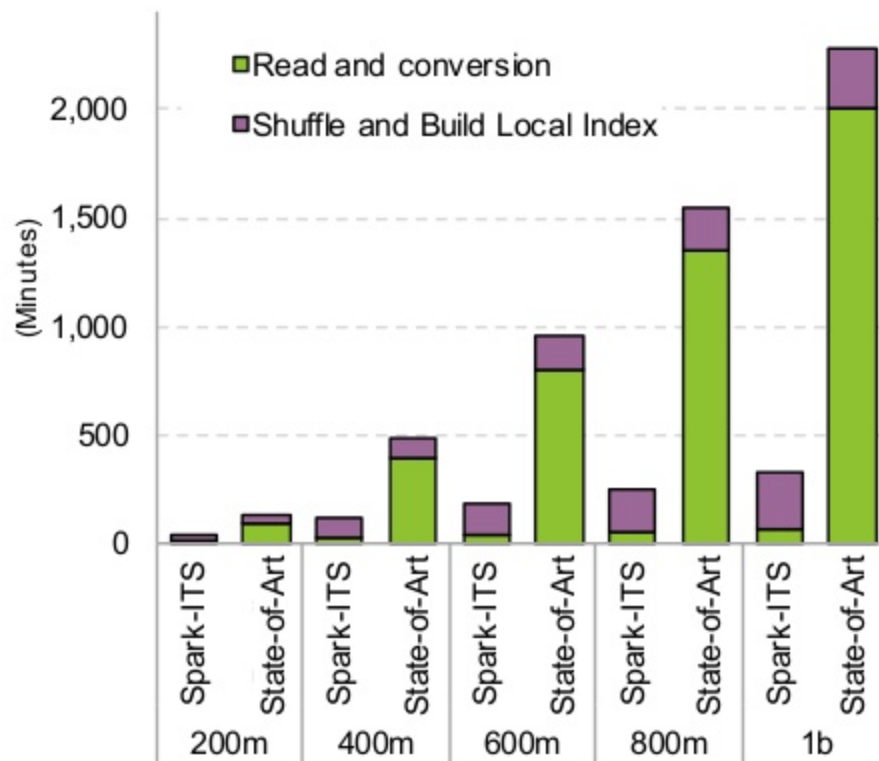
Dataset: Random Walk Benchmark

# Index Construction Time: Breakdown

## Global Index Time Breakdown



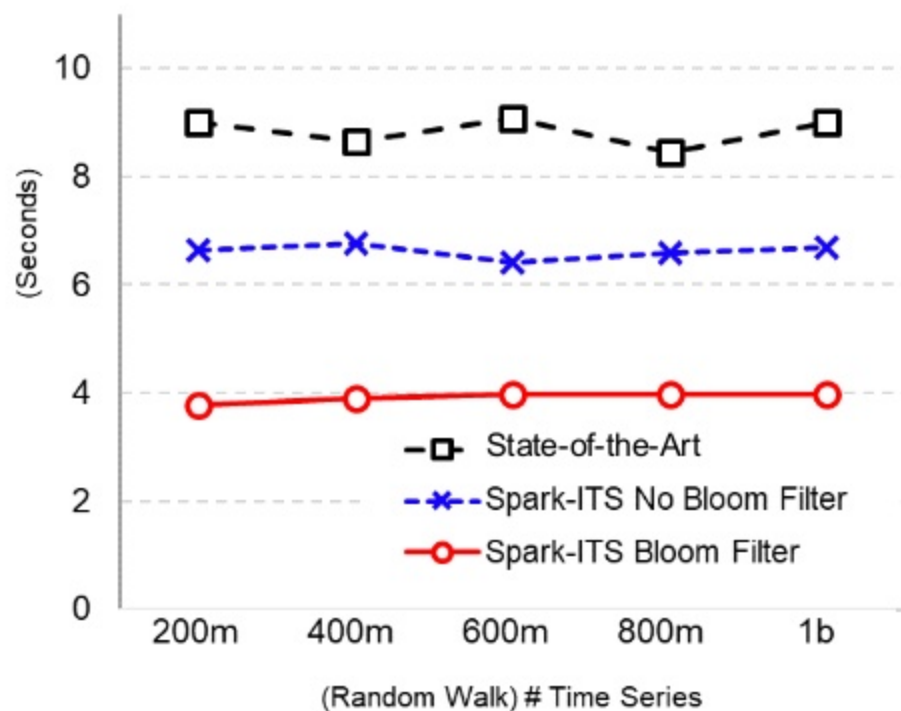
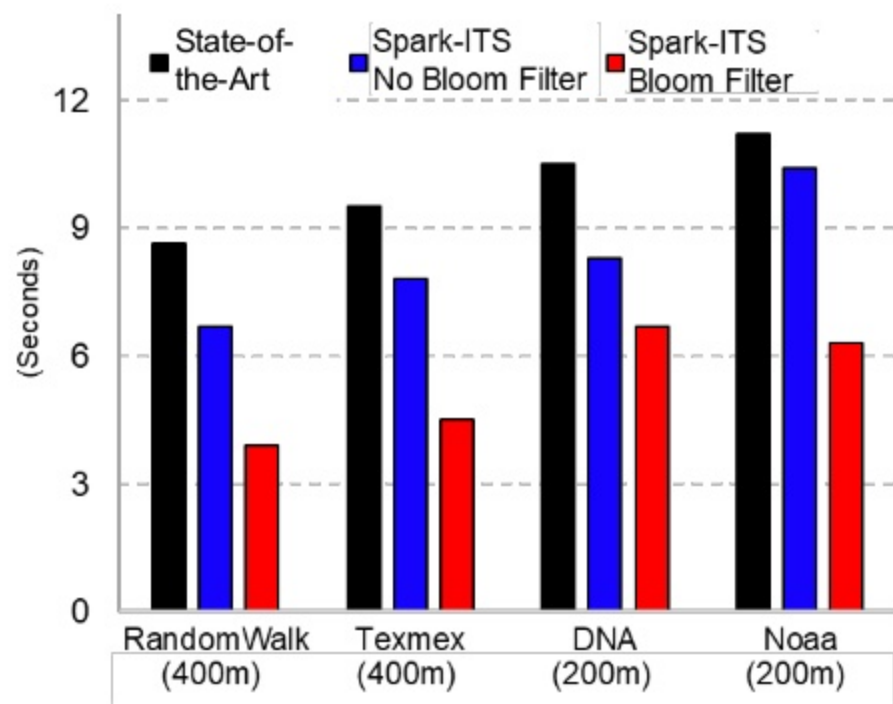
## Repartition and Local Index Time Breakdown



Dataset: Random Walk Benchmark

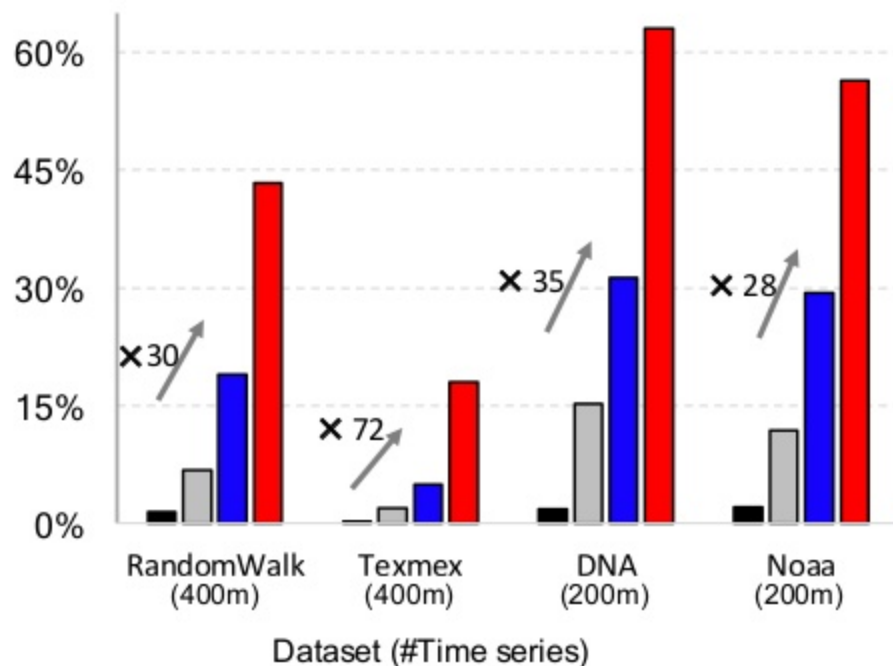


# Exact Matching Query

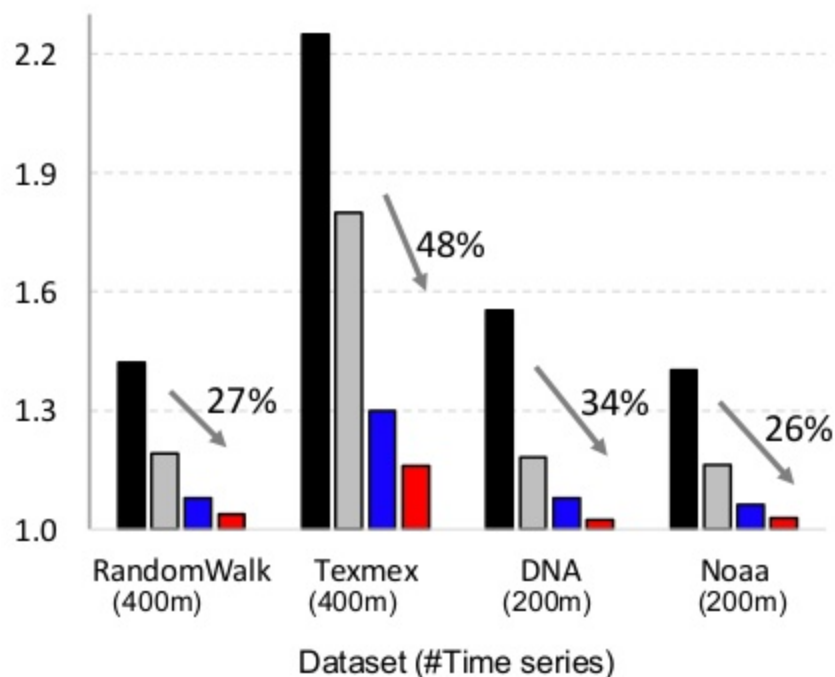


# kNN-Approximate Query Performance

## Recall



## Error Ratio



State-of-the-Art Target Node Access One Partition Access Multi-Partitions Access

# Conclusion

- Index Tree
  - Large fan-out decreases the depth of leaf nodes
  - Keeps better similarity at Word-level
  - The signature simplifies the conversion of cardinality
- Spark-ITS: Index Construction
  - Block-sampling and node statistic collection to **fast build global index**
  - **Synchronously** build local indices within a partition
  - Constructs Index faster **80+%**.
- Spark-ITS: Query
  - Exact Matching: the time decreases **by 50%**.
  - kNN approximate: the accuracy increases more than **10 fold**.

## Acknowledge Funding from...

Xianjin Tech Co., Ltd.

Saudi Arabian Cultural Mission

WPI Computer Science Dept.,

NSF CNS: 305258 II-EN

NSF CRI: 0551584

