

Unsupervised Object Detection using the Azure Cognitive Services on Spark



Mark Hamilton, Microsoft
Anand Raman, Microsoft

#SAISExp4



Snow
Leopard
Trust



NEWS

Home | Video | World | US & Canada | UK | Business | Tech | Science | Stories | Enter

Asia | China | India

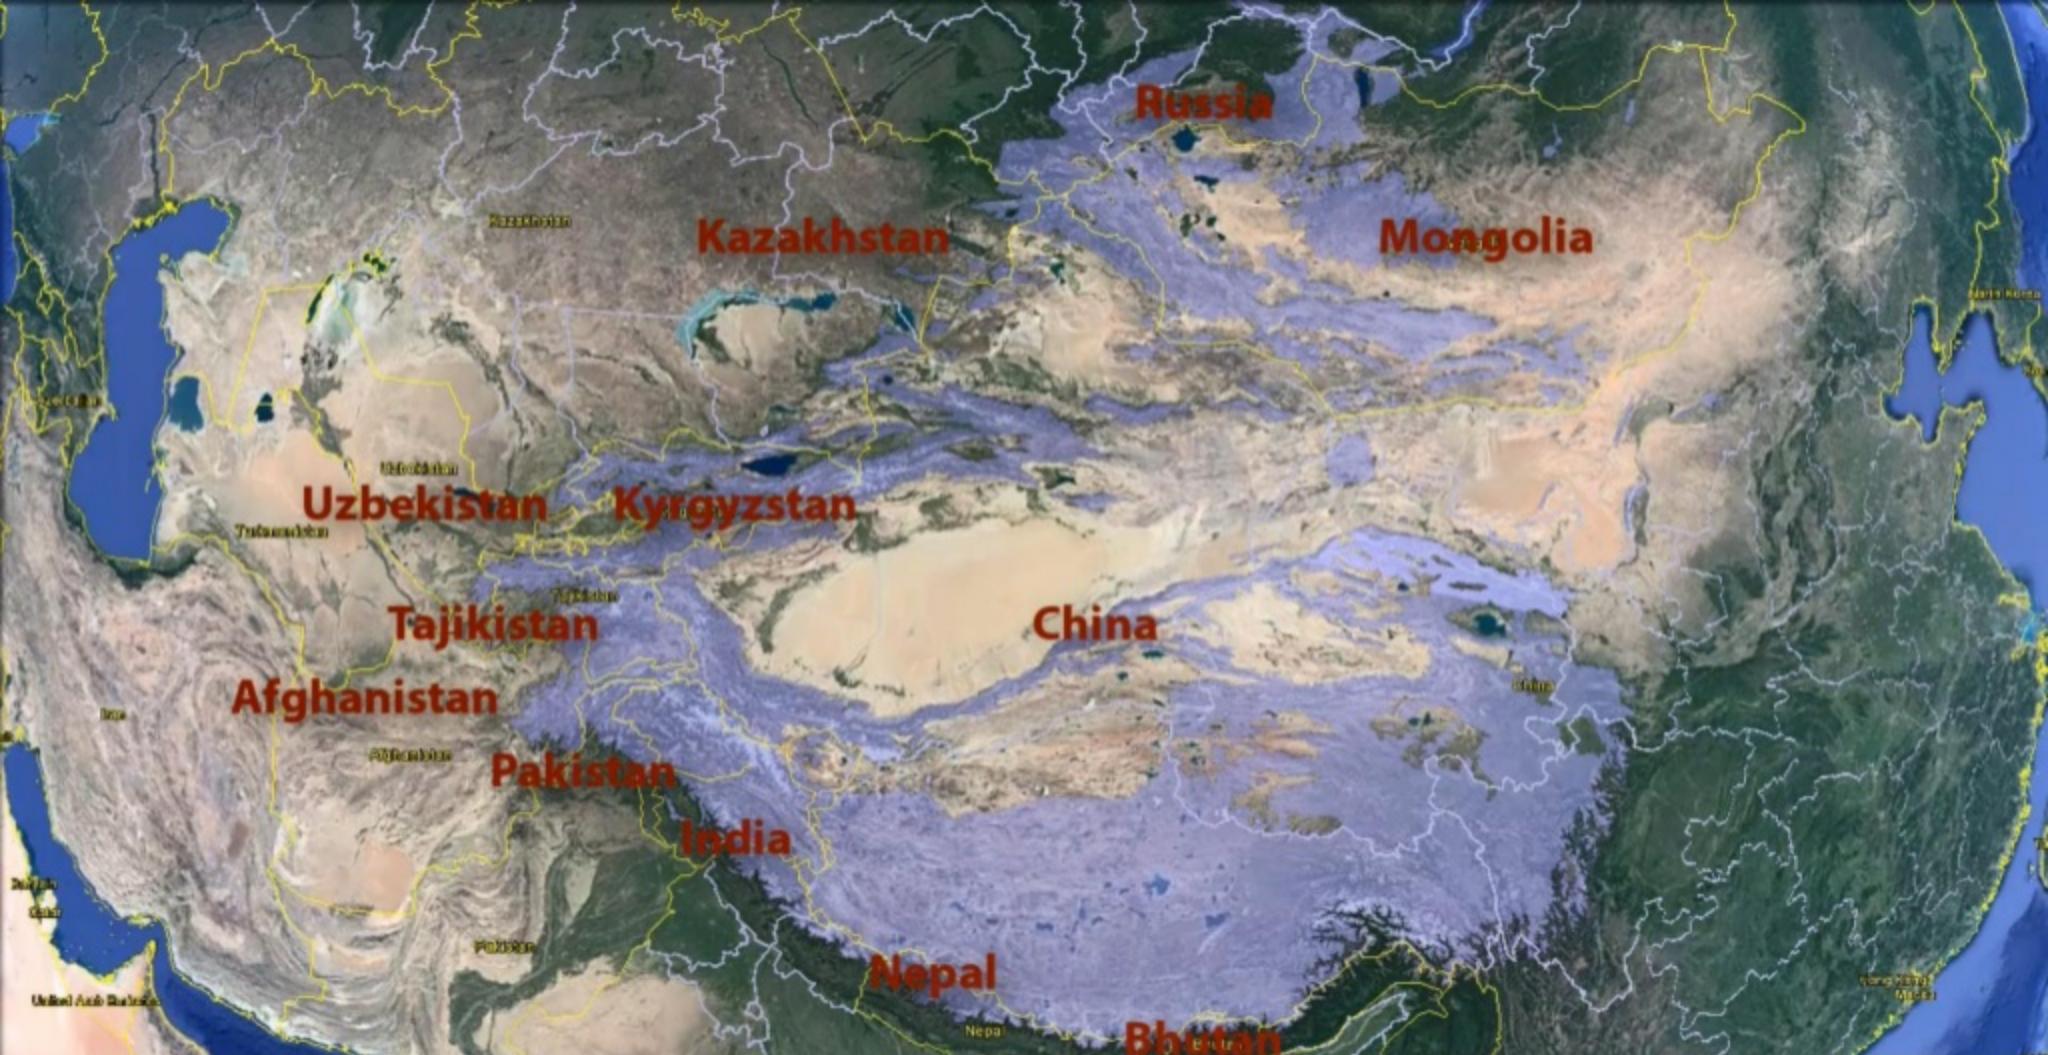
Snow leopard no longer 'endangered'

© 14 September 2017

[f](#) [t](#) [m](#) [e](#) [Share](#)

Statement on IUCN Red List Status Change of the Snow Leopard

The Snow Leopard Trust, one the leading conservation organizations working to protect this cat, opposes the IUCN's decision to change the snow leopard's Red List status from 'Endangered' to 'Vulnerable'.





Snow
Leopard
Trust

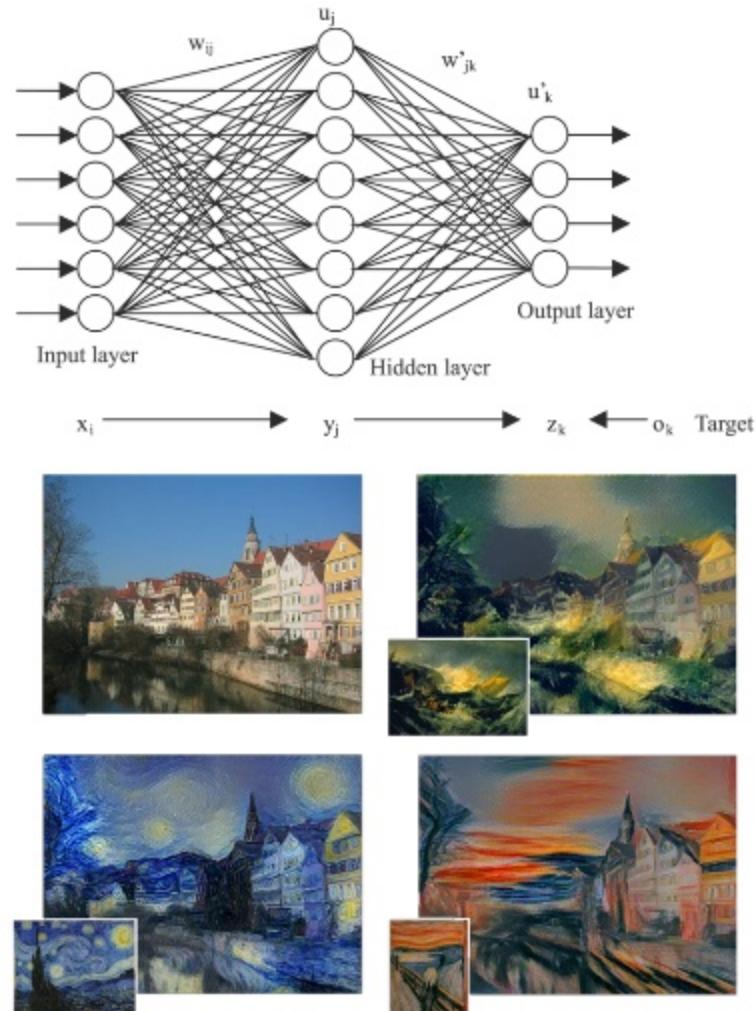
Classifying all 1.3 million images will
take ***20k hours***

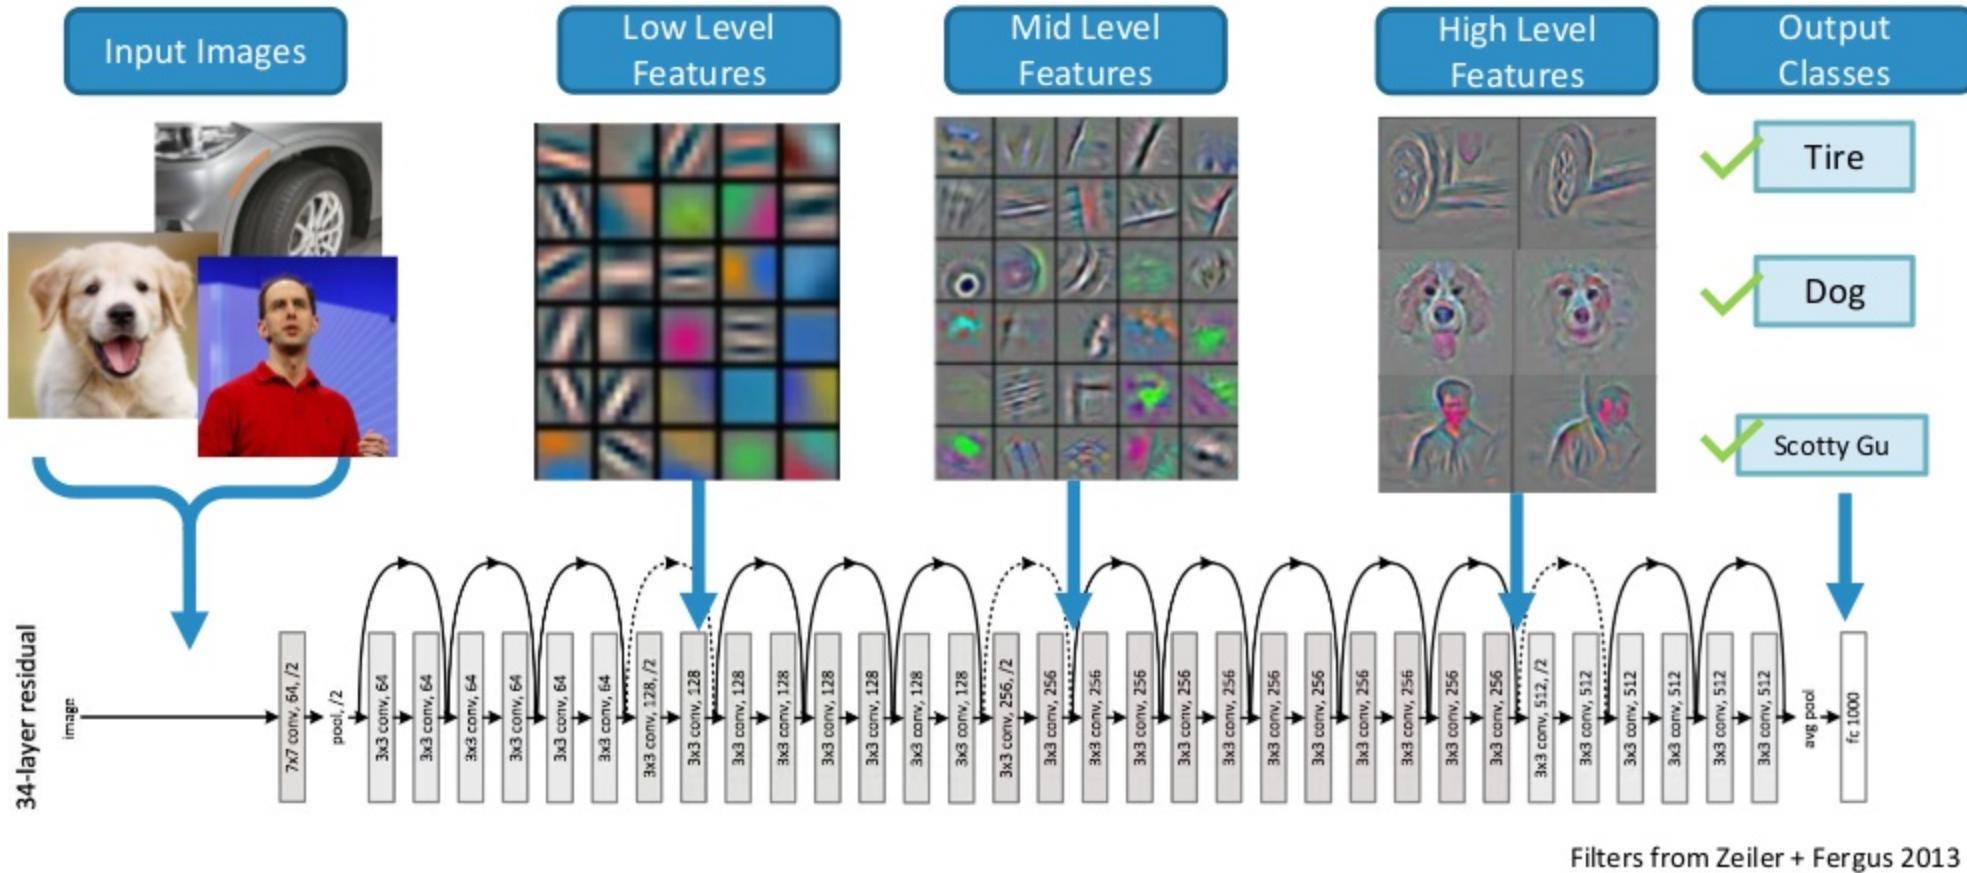


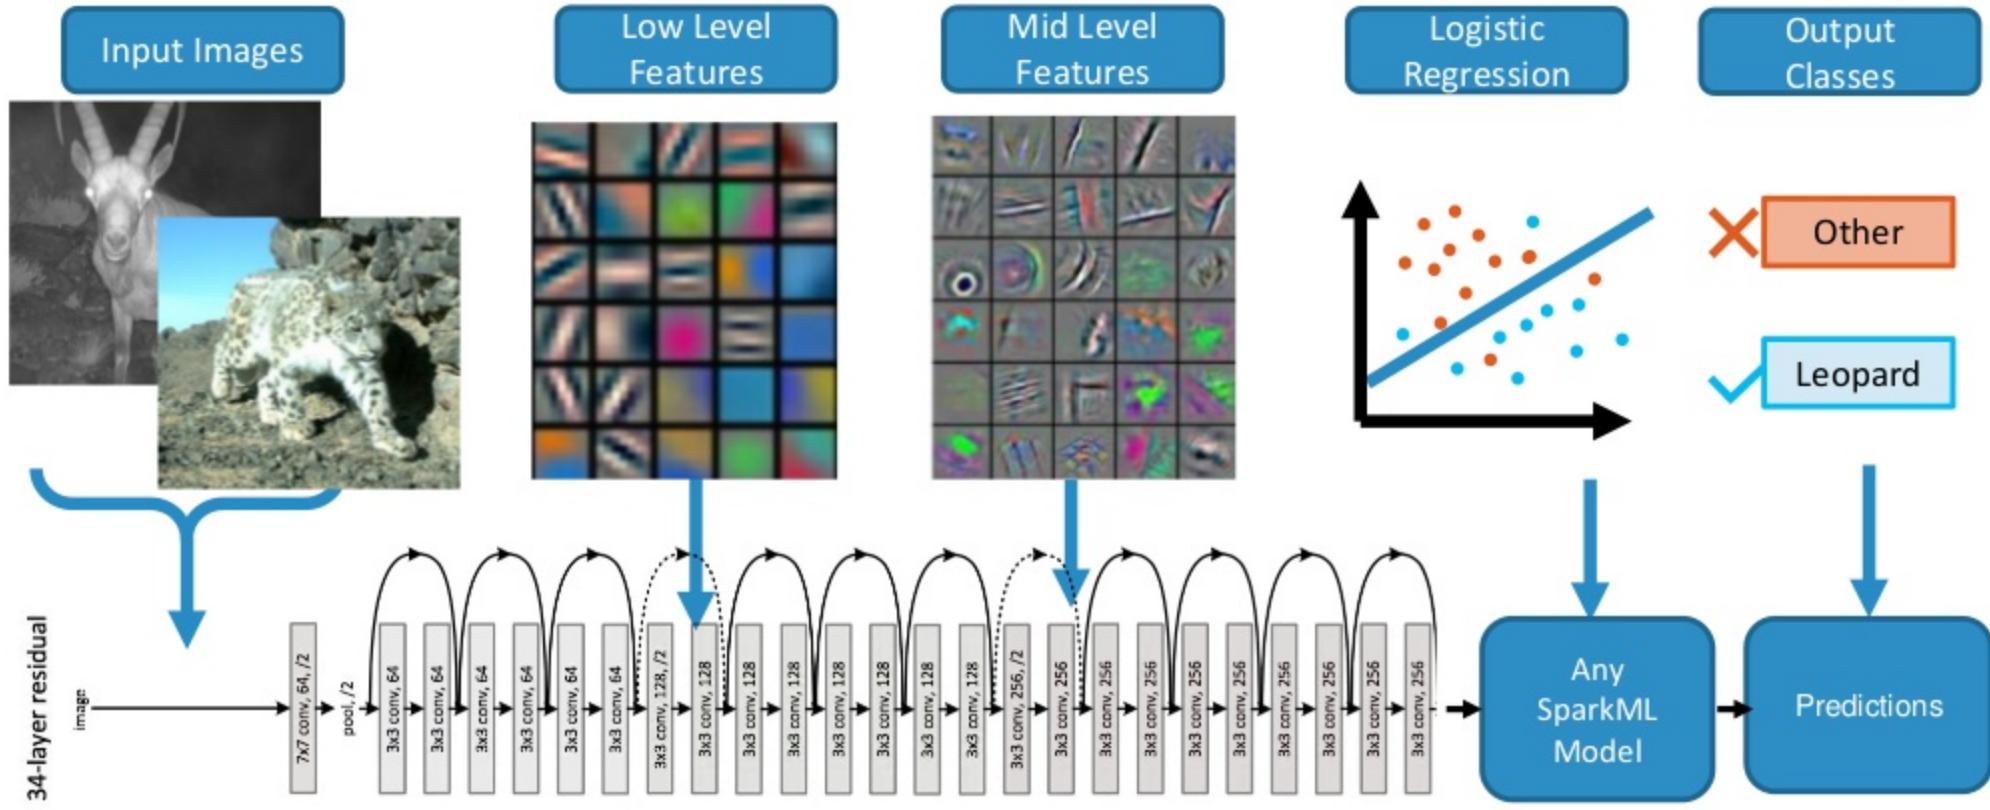


Deep Learning

- Brain-like algorithms trained with gradient descent
- Has become a recent favorite because:
 - Spectacular performance in many domains
 - Quick training
 - Low memory
 - Large space of possible model architectures
 - Automatic differentiation software makes it very easy





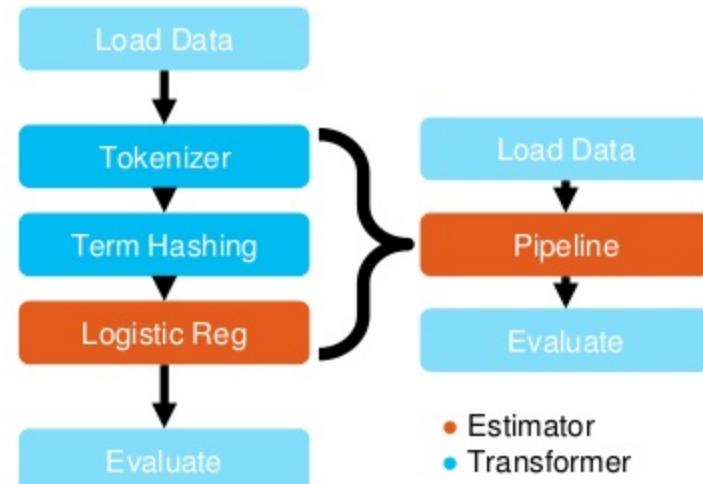


Filters from Zeiler + Fergus 2013

APACHE Spark™ ML

- High level library for distributed machine learning
- More general than SciKit-Learn
- All models have a uniform interface
 - Can compose models into complex pipelines
 - Can save, load, and transport models

```
data = spark.read.csv("hdfs://...")  
train, test = data.randomSplit([.5,.5])  
model = LogisticRegression().fit(train)  
predictions = model.transform(test)
```



Distributed Computing



- ▶ Large scale parallelism
- ▶ Fault tolerance
- ▶ Auto-scaling / elasticity
- ▶ High throughput streaming
- ▶ Data source + orchestrator agnostic



Deep Learning



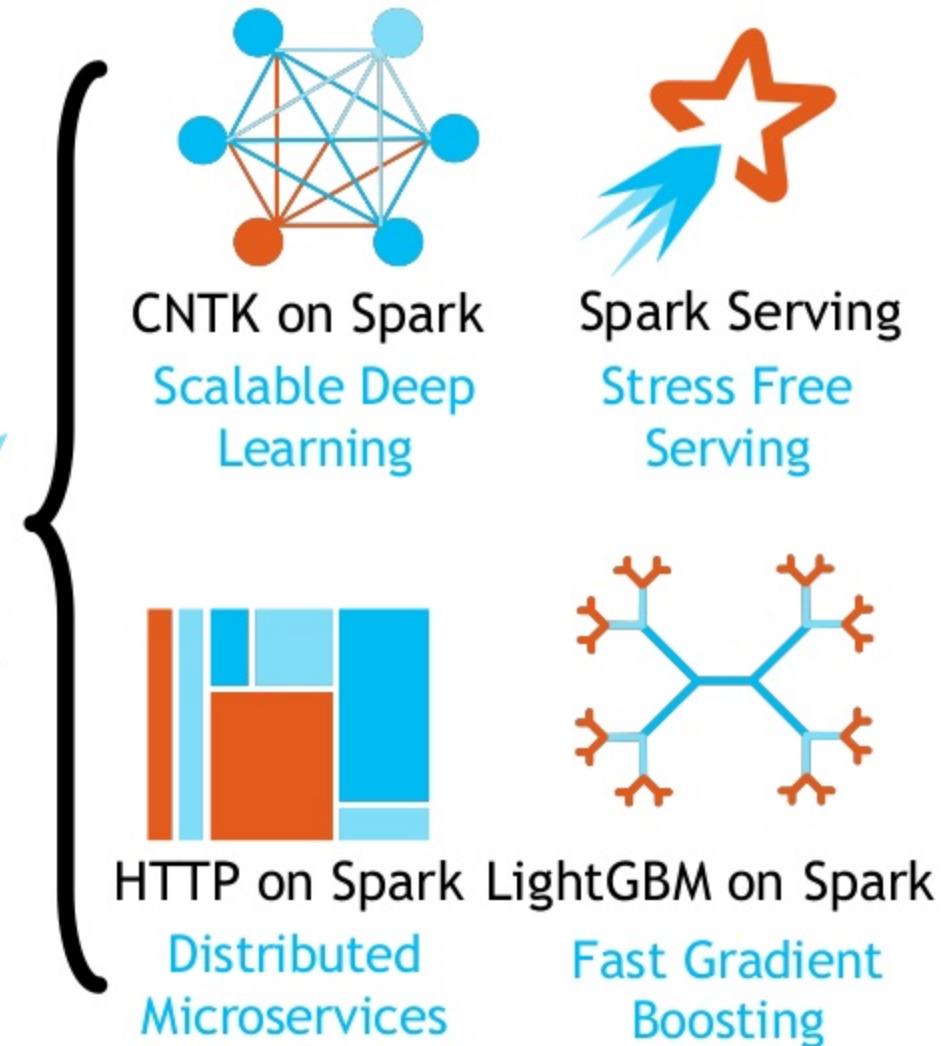
- ▶ GPU/CPU acceleration
- ▶ Automated gradient calculations
- ▶ Flexible language for defining huge space of models
- ▶ State of the art performance and accuracy
- ▶ ONNX Runtime



Microsoft ML for Apache Spark

www.aka.ms/spark

- **Open Source** Distributed ML library
- Unifies major computing paradigms and Microsoft tools into a single library and API
 - Python, Java, Scala, R
 - Batch, Streaming, Serving
 - Local, Static Cluster, Dynamic Cluster, Serverless
 - Databricks, HDI, AZTK, Kubernetes, ...



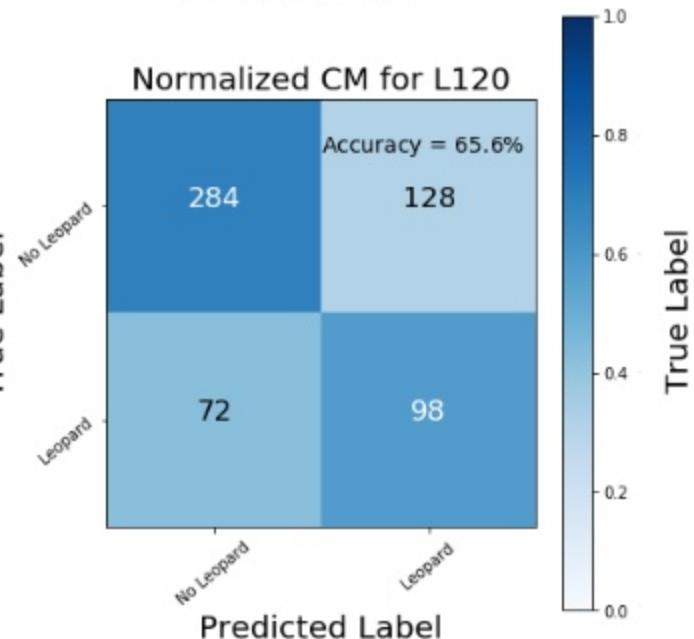
Training a Deep Transfer Learner:

```
1 import mmlspark as mml
2
3 network = mml.ModelDownloader(...).downloadByName("ResNet50")
4
5 model = Pipeline(stages = [
6     mml.ImageSetAugmenter(flipHorizontal=True),
7     mml.ImageFeaturizer(cutOutputLayers=2).setModel(network),
8     LogisticRegression()
9     mml.EnsembleByKey(keys=["group","location"], cols=["prob"])))
10
11 train, test = spark.read.parquet(...).split(...)
12
13 fit_model = model.fit(train)
14 results = fit_model.transform(test)
```

Performance

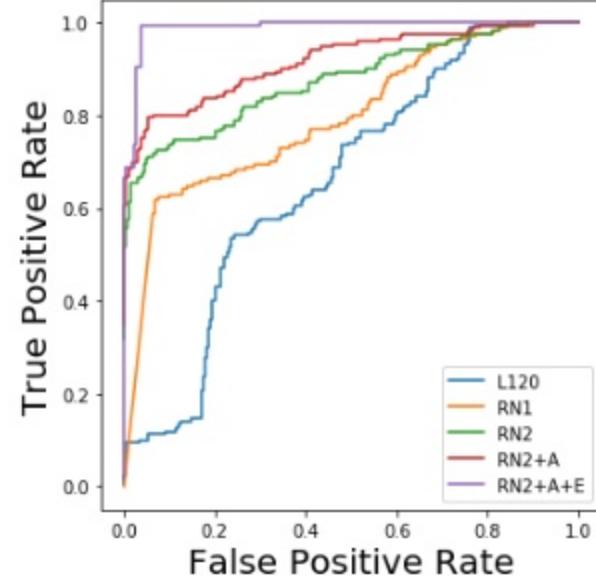
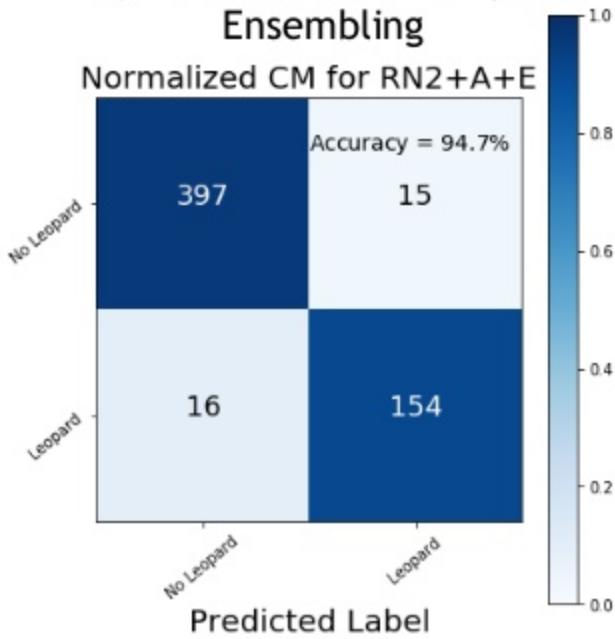
Without Deep
Featurization

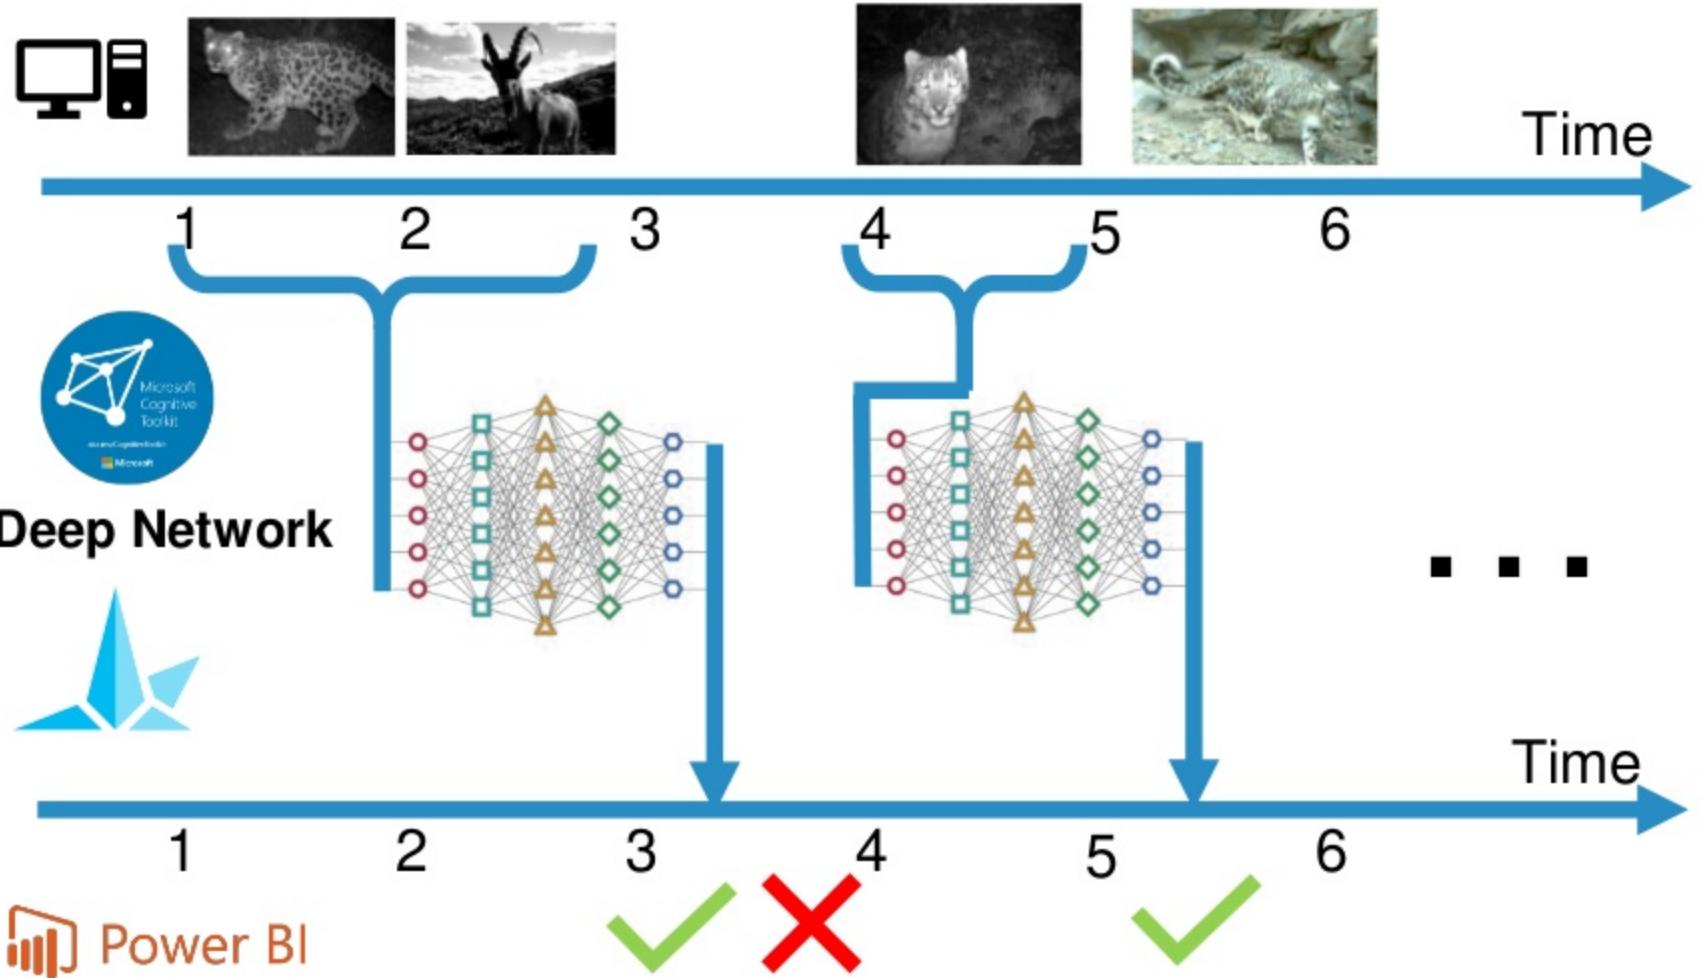
Normalized CM for L120



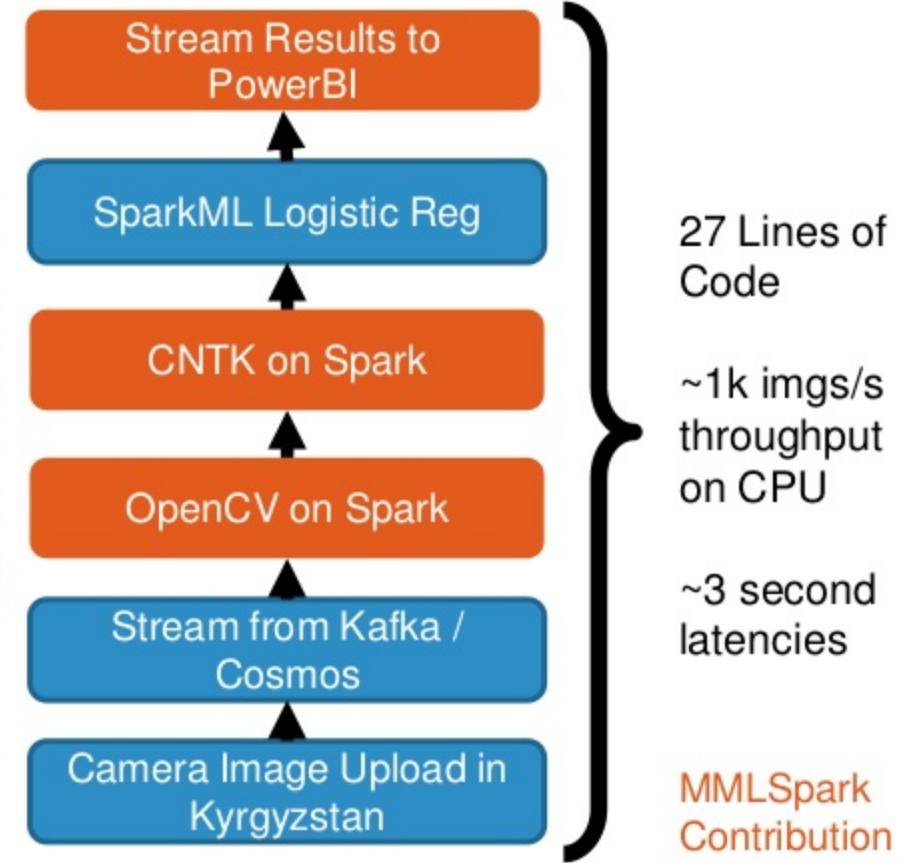
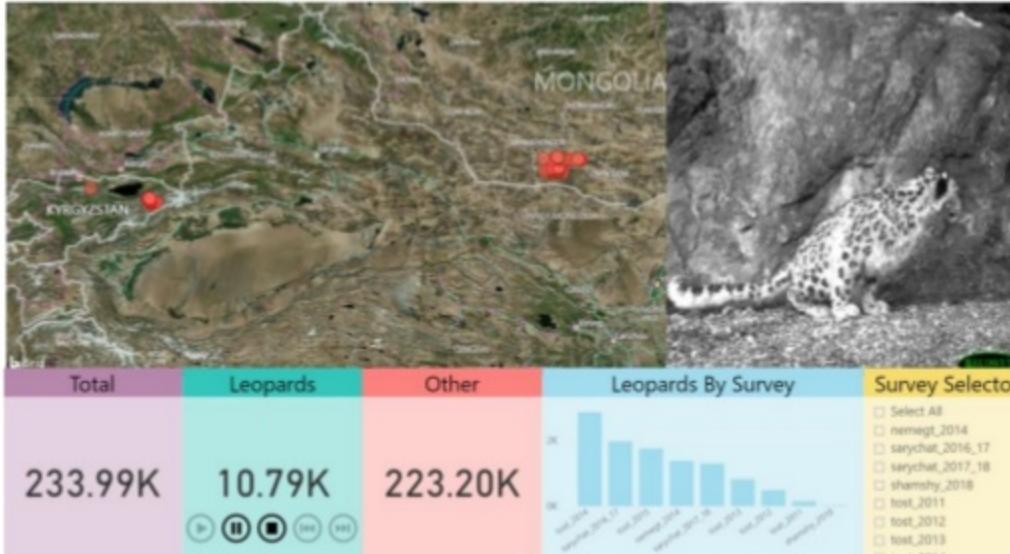
With Deep Featurization,
Augmentation, and Temporal
Ensembling

Normalized CM for RN2+A+E





Real Time monitoring with PowerBI



Simple APIs matter

```
16 streaming_df = spark.readStream.file(...)  
17 fit_model.transform(streaming_df)  
18 stream = PowerBIWriter.stream(results, powerbi_url).start()
```

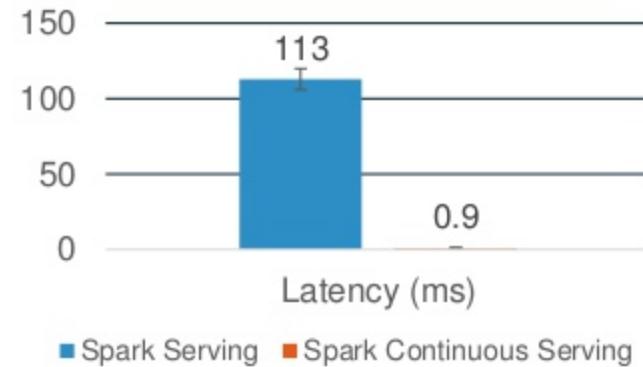
APACHE Spark™ Serving

Lightning Fast Web Services on **Any** Spark Cluster

- Sub-millisecond latencies
- Fully Distributed
- Spins up in seconds
- Same API as Batch and Streaming
- Scala, Python, R and Java
- Fully Open Source



Announcing: **100x Latency Reduction**
with MMLSpark v0.14



www.aka.ms/spark
JIRA: SPARK-25350

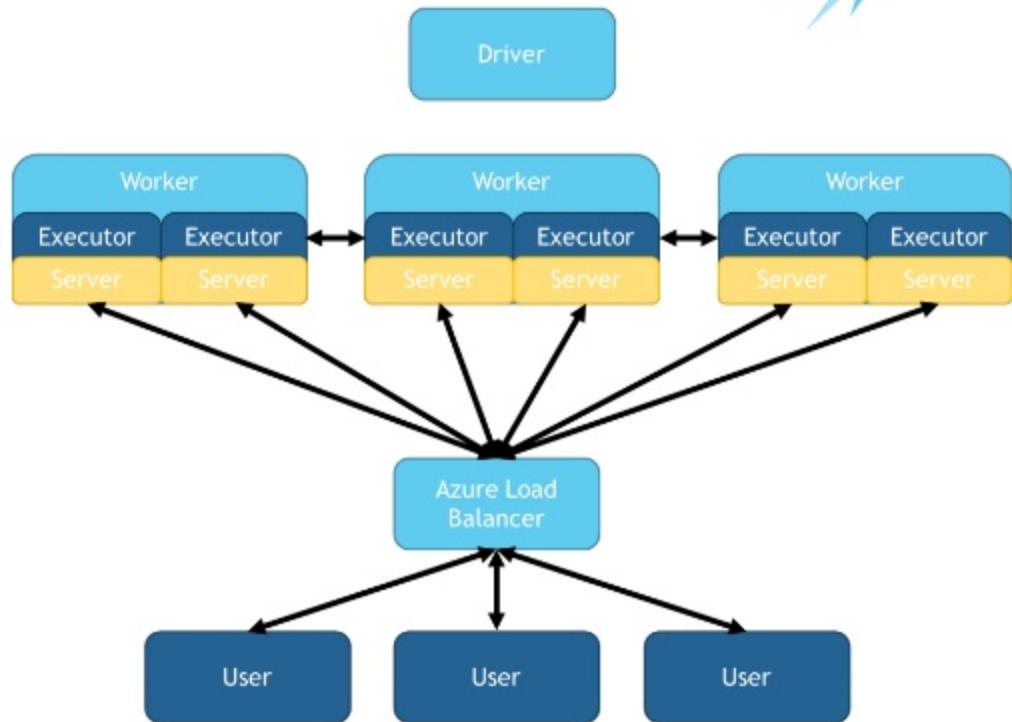
Web Serving with the Same API:

```
4 serving_df = spark.readStream.server()  
5     .address("0.0.0.0", 8888, "leopard_detector")  
6     .load()  
7     .parseRequest(BinaryType())  
8  
9 results = fitModel.transform(serving_df)  
10  
11 results.makeReply("probability")  
12     .writeStream  
13     .server()  
14     .replyTo("leopard_detector")  
15     .start()  
16
```

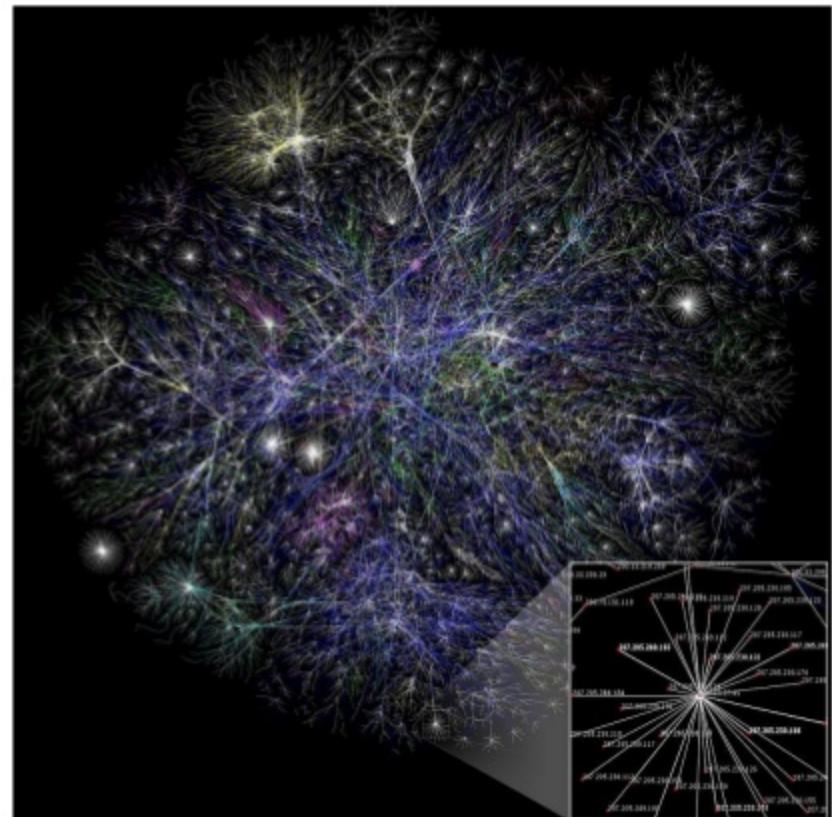
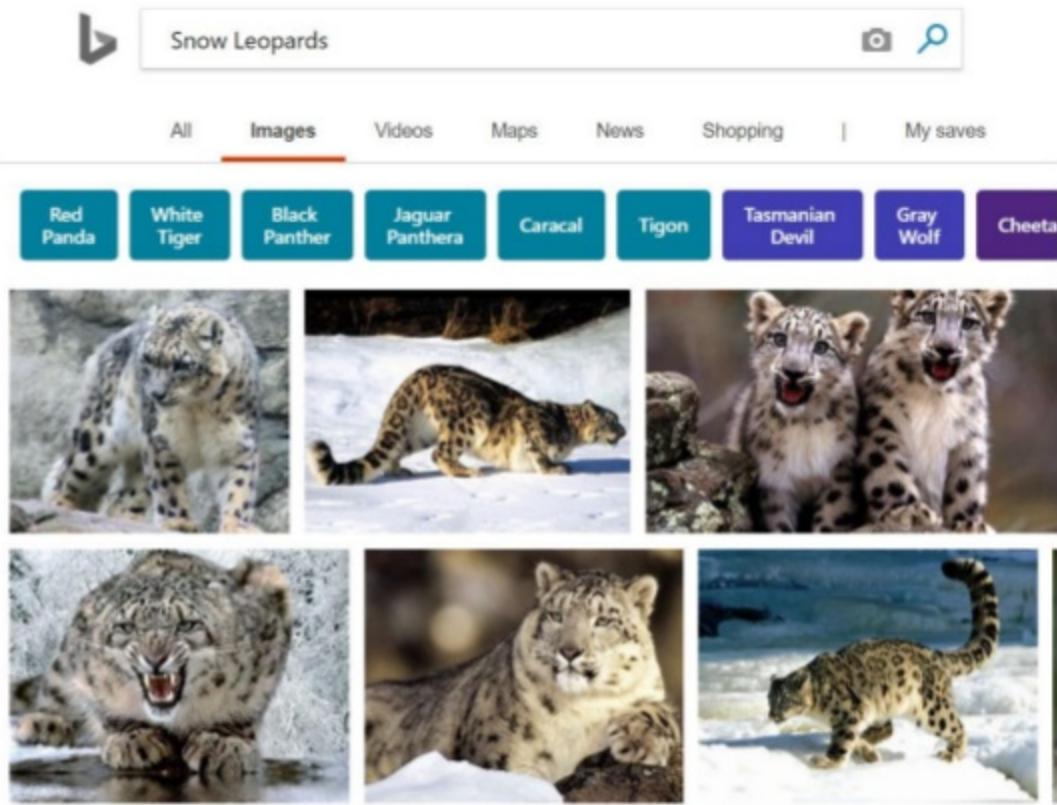
Spark Serving Architecture



- 3 main modes:
 - server: 1 service on the head node
 - distributedServer: 1 service per executor
 - continuousServer: 1 service per partition
- Built on top of Spark Streaming
- Each worker keeps a running service, and a routing table
- Use HTTP on Spark objects to represent requests and responses in Spark SQL

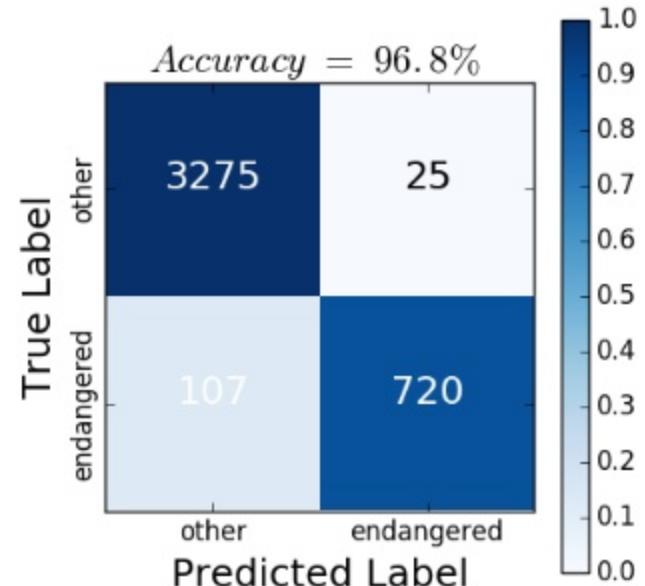
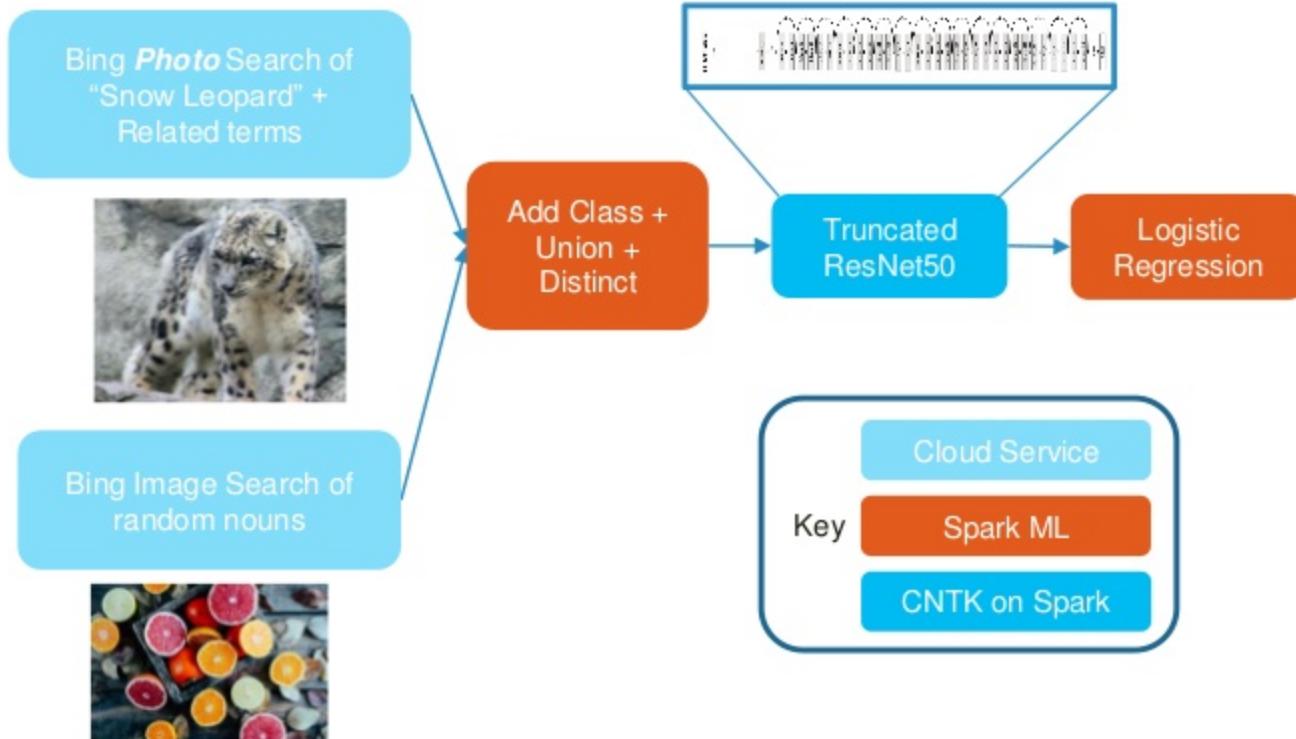


Going Further: No Labels, No Problem



Hamilton and Raman, #SAISExp4

Bing Powered Snow Leopard Detection Pipeline



Bing Image Search API

```
endangeredLinks = requestParameters \  
    .transform(BingImageSearch()) # Apply Bing Image Search  
    .setSubscriptionKey(os.environ['API_KEY'])  
    .setOffsetCol("offsets")  
    .setQueryCol("queries")  
    .setCount(10)  
    .setImageType("photo")  
    .setOutputCol("images")) \  
    .transform(GetBingURLs(inputCol="images", outputCol="urls")) \  
    .withColumn("labels", lit("endangered")) # Add the labels
```

```
randomLinks = spark.read.table("random_words") \  
    .transform(BingImageSearch())  
    .setSubscriptionKey(os.environ['API_KEY'])  
    .setCount(10)  
    .setQueryCol("words")  
    .setOutputCol("images")) \  
    .transform(GetBingURLs(inputCol="images", outputCol="urls")) \  
    .withColumn("label", lit("other"))
```

urls	labels	urls	label
	endangered		other
	endangered		other
	endangered		other
	endangered		other



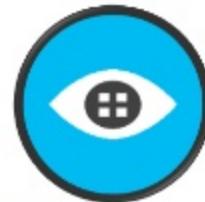
Announcing: Intelligent and Distributed
Microservices in SparkML



Bing



Text



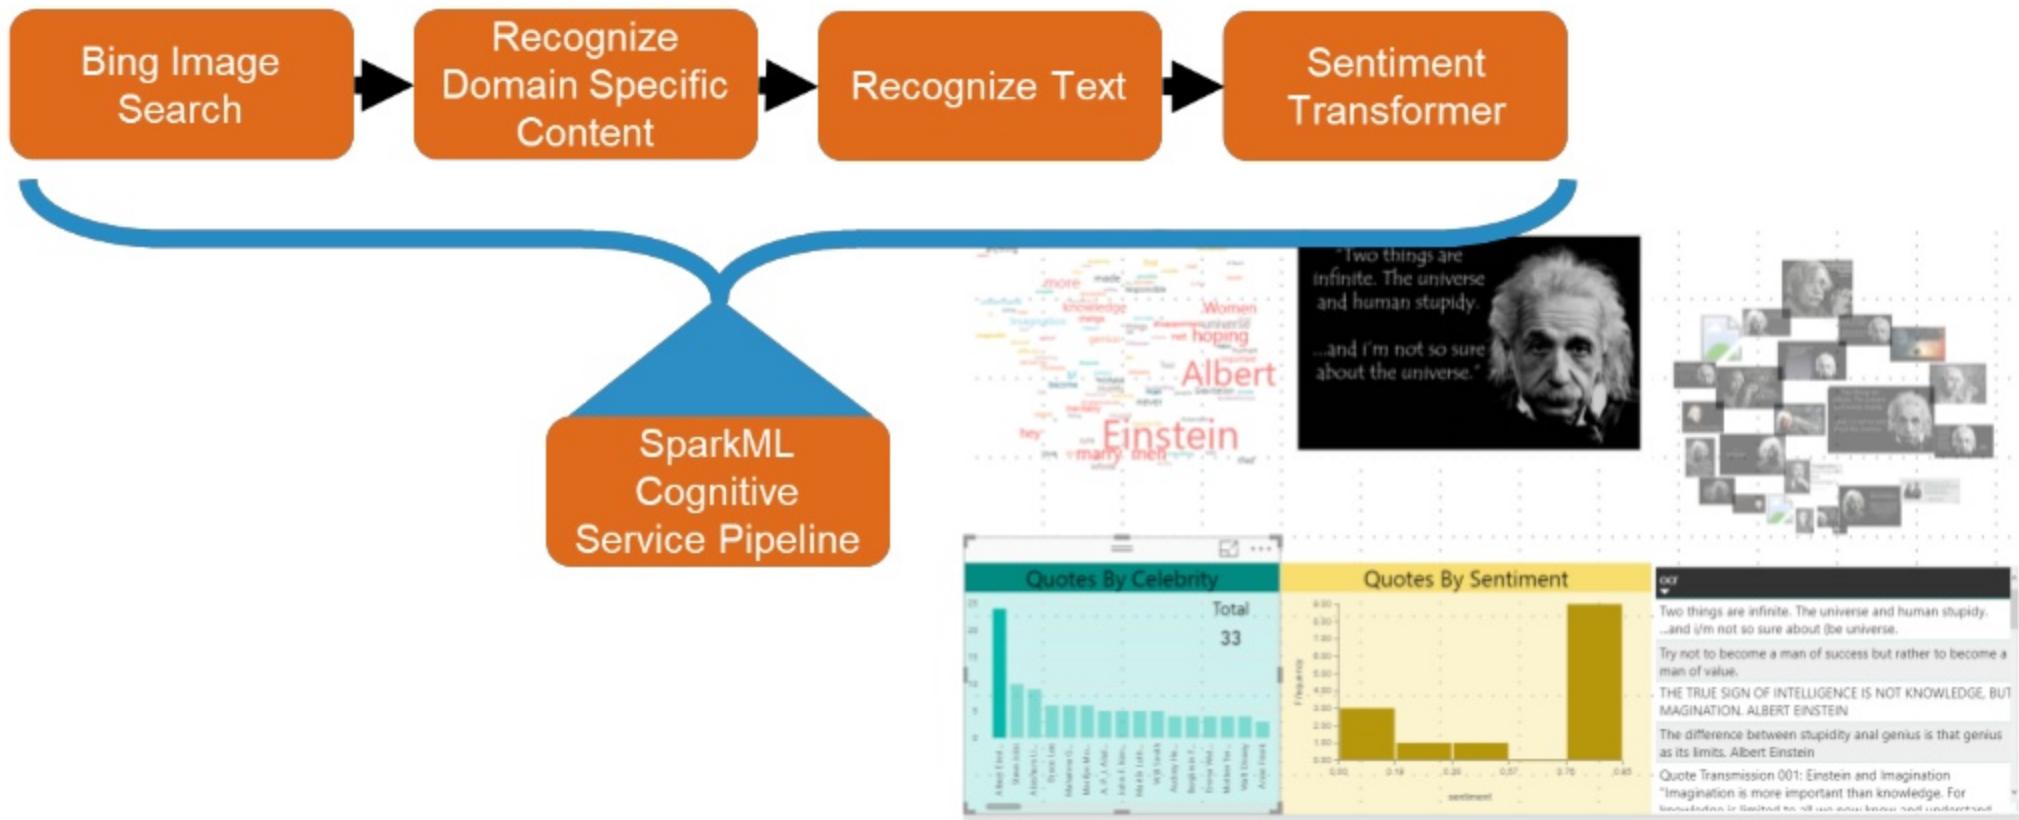
Vision

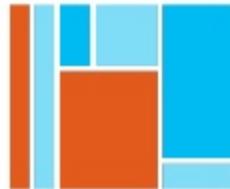


Face

www.aka.ms/spark

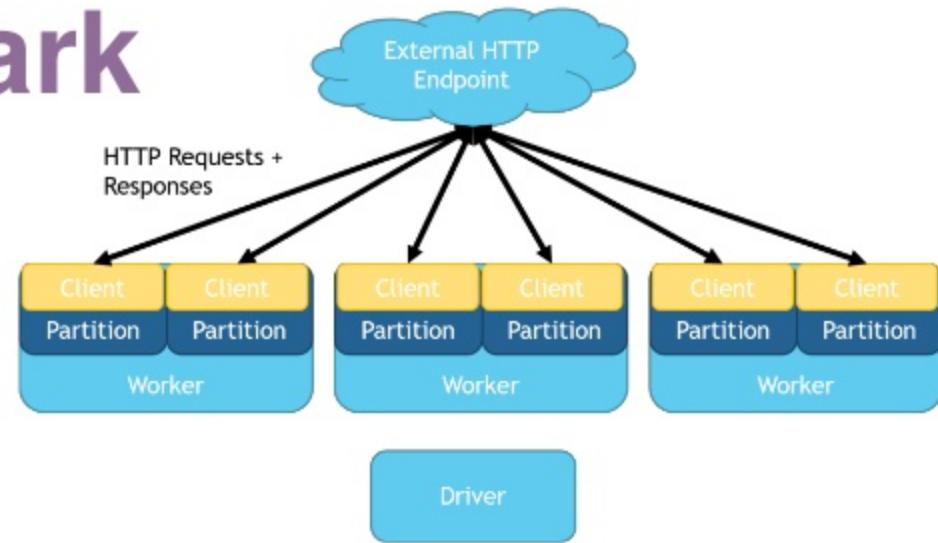
Celebrity Quote Analysis





HTTP on Spark

- Allows Spark users to call into *any* HTTP service
- Leverage parallel networking of cluster
- Allows Spark to work with microservice architectures
- Mini-batching, asynchrony, and buffering supported
- Statically typed and usable from Spark, PySpark, and SparklyR



```
replies = SimpleHTTPTransformer()  
    .setOutputParser(new JSONOutputParser().setDataType(...)) \  
    .setUrl("http://my_favorite_webservice.com") \  
    .setInputCol("data") \  
    .setOutputCol("results") \  
    .setConcurrency(3) \  
    .transform(df)
```

Model Interpretability with LIME



(a) Original Image



(b) Explaining *Electric guitar*



(c) Explaining *Acoustic guitar*



(d) Explaining *Labrador*

Figure 4: Explaining an image classification prediction made by Google’s Inception network, highlighting positive pixels. The top 3 classes predicted are “Electric Guitar” ($p = 0.32$), “Acoustic guitar” ($p = 0.24$) and “Labrador” ($p = 0.21$)

Source: [Marco Tulio Ribeiro et al.](#)

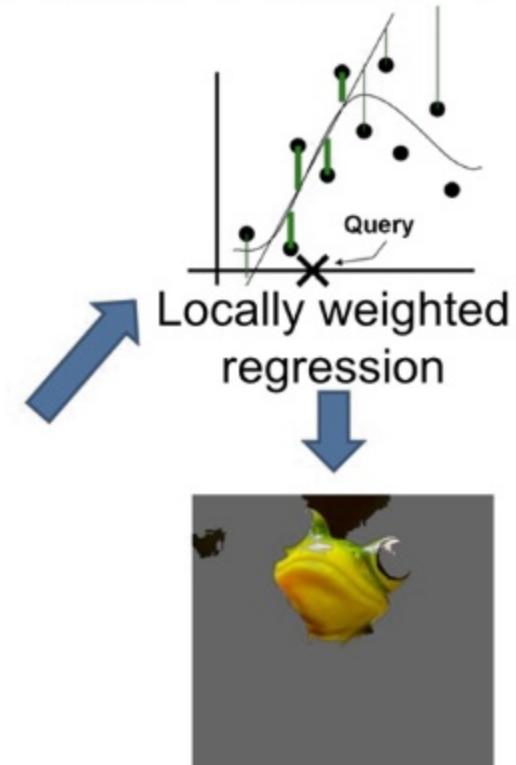
Model Interpretability with LIME



Original Image
 $P(\text{tree frog}) = 0.54$



Perturbed Instances	$P(\text{tree frog})$
	0.85
	0.00001
	0.52

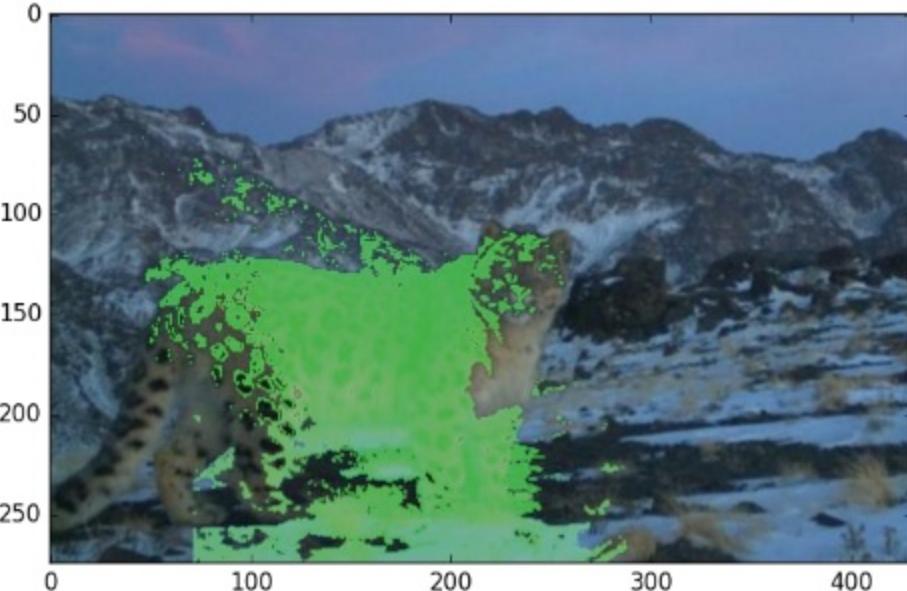
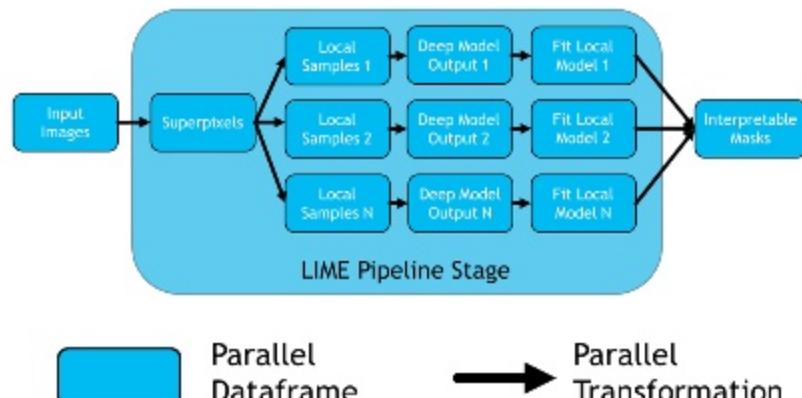


Explanation

Source: [Marco Tulio Ribeiro et al.](#)

Announcing: LIME on Spark

- Elastic and Distributed
- Works with *any* image classification model
- Example notebook instructions at:
 - www.aka.ms/spark



Using LIME to generate bounding boxes

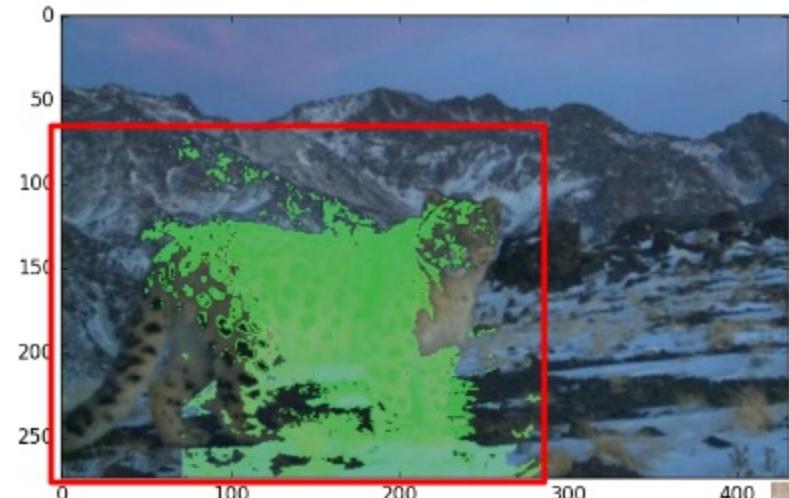
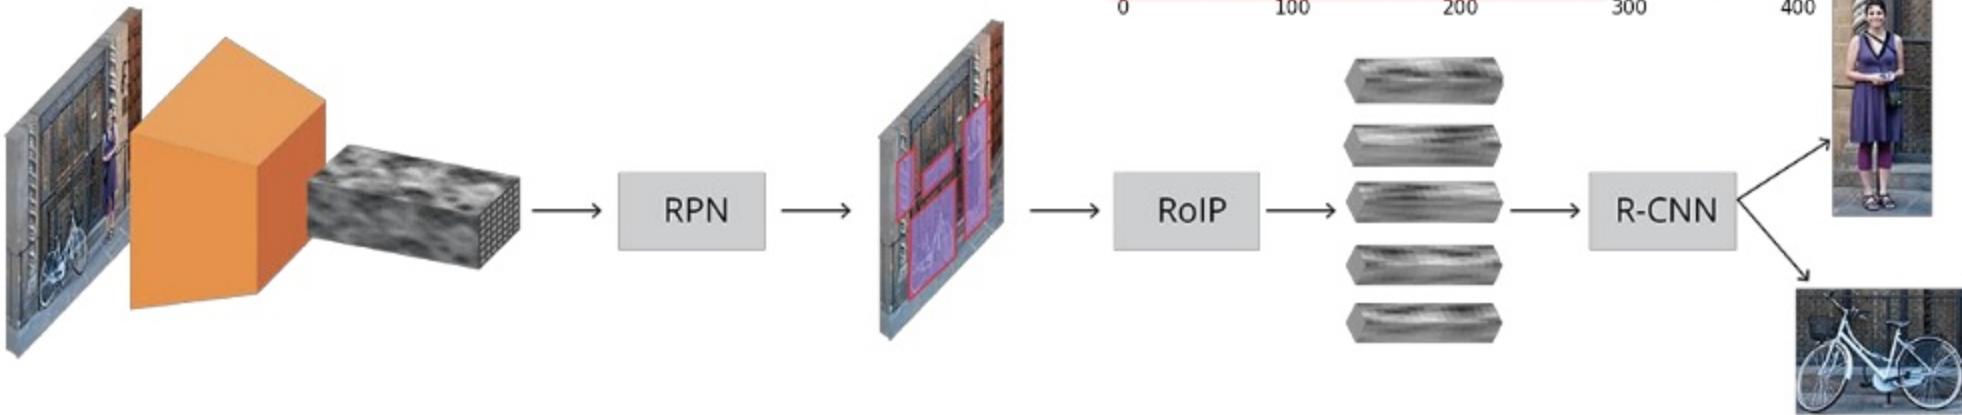


Transferring LIME's Knowledge to FasterRCNN

$mask = \text{union}(\text{superpixels} \mid \text{top } 70\% \text{ of weights})$

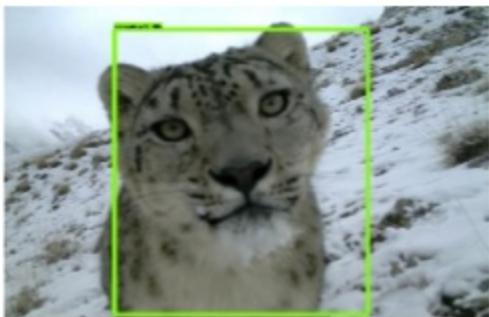
$x1, y1 = \min(mask_x), \min(mask_y)$

$x2, y2 = \max(mask_x), \max(mask_y)$

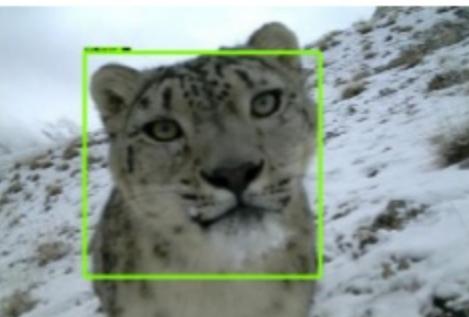


Results

Human Labels



Unsupervised
FRCNN Outputs

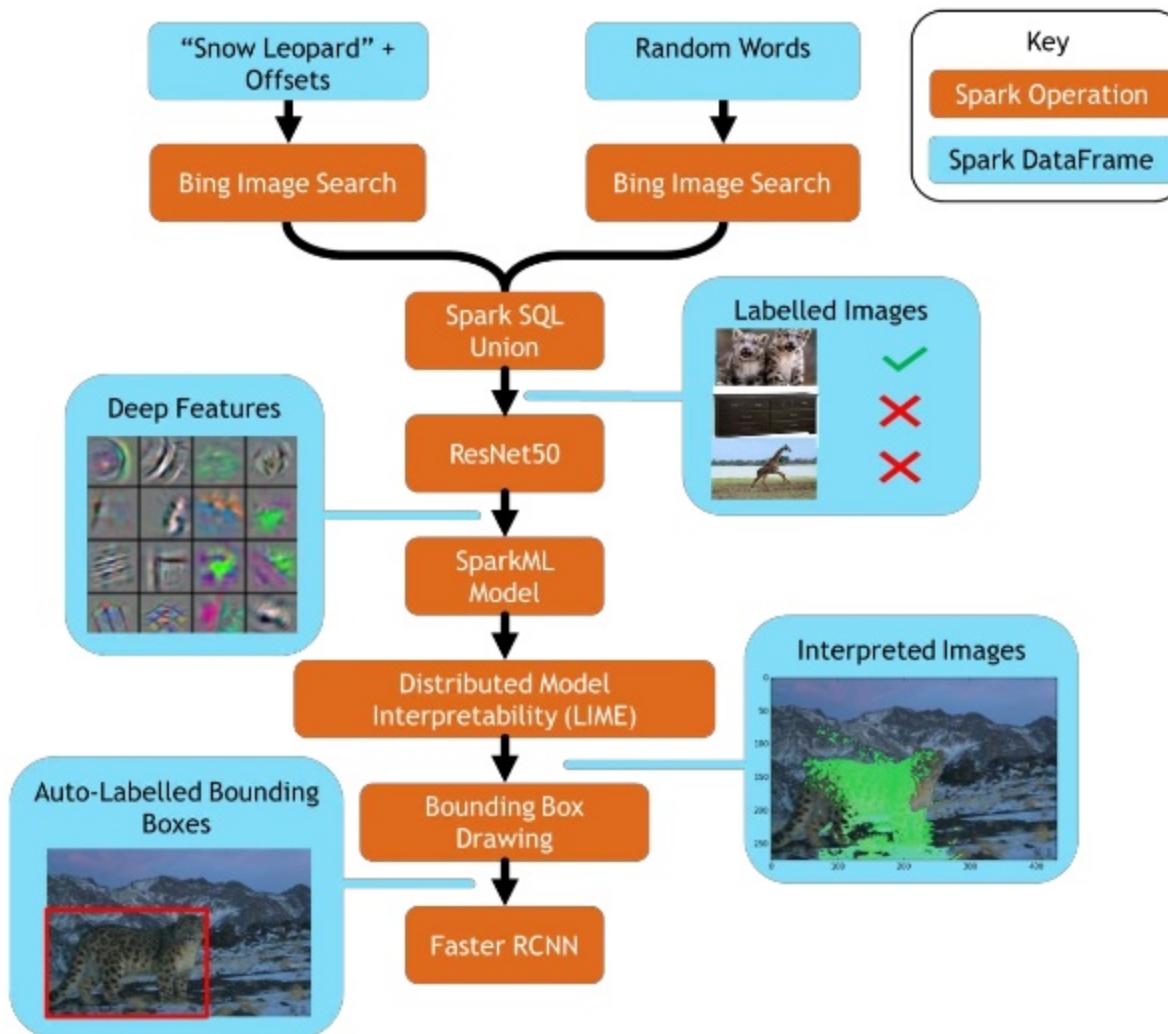


Human Labels

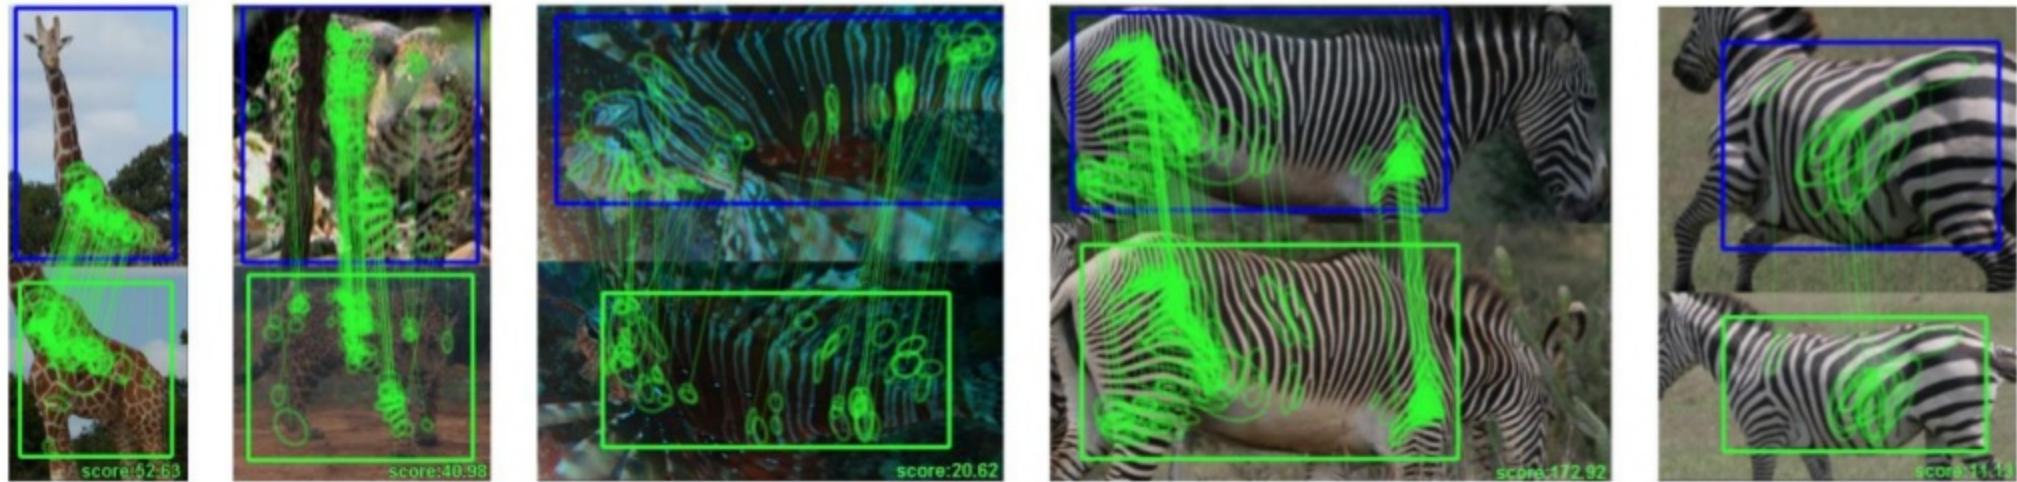


Unsupervised
FRCNN Outputs





Next Steps: Hotspotter



Source: HotSpotter - Patterned Species Instance Recognition

In Conclusion

- Big Releases in v0.14:
 - **Cognitive Services on Spark**
 - **Sub-Millisecond latency distributed web services**
 - **Distributed Model Interpretability**
- Easy batch, streaming apps, PowerBI dashboards, or RESTful web services
- We aim to give all work back to the community!
- Easy to Get Started on Databricks:
 - 16 Jupiter notebook guide



www.aka.ms/spark

Contributions Welcome!
github.com/Azure/mmlspark

Thanks to

- You all!
- **Microsoft AI Development Acceleration Program: Abhiram Eswaran, Ari Green, Courtney Cochrane, Janhavi Suresh Mahajan, Karthik Rajendran, Minsoo Thigpen, Casey Hong, Soundar Srinivasan**
- MMLSpark Team: Sudarshan Raghunathan, Ilya Matiach, Eli Barzilay, Tong Wen, Ben Brodsky
- The Snow Leopard Trust: Rhetick Sengupta, Koustubh Sharma, Jeff Brown, Michael Despines
- Microsoft: Joseph Sirosh, Lucas Joppa, WeeHyong Tok

MMLSpark Website: aka.ms/spark

Get in touch: marhamil@microsoft.com