

Lessons from the Field, Episode II:

**Applying Best Practices to Your Apache Spark™
Applications**

Silvio Fiorito
@granturing

#SAIExp9

About Me

Resident Solutions Architect @ Databricks

Developer, application security engineer,
consultant, instructor

Using Spark since ~2012

Work with Databricks customers in banking,
manufacturing, info-sec, new & experienced Spark users

SSE17: [Lessons from the Field](#)



To build more reliable,
& better performing Spark
applications we should
understand...

File Formats & Data Types

Partitioning, Shuffling, & Parallelization

Using Query Plans to Debug & Tune

CSV, JSON, Parquet, Avro, ORC,...

File formats have different
benefits and tradeoffs that
we need to be aware of...

Reading CSV Files

Raw text

```
1000,user1,"2010-01-23",123.34,.....  
2002,user2,"2008-02-14",349.02,.....
```

Read each
line as String

```
"1000,user1,2010-01-23,123.34,.....",  
"2002,user2,2008-02-14,349.02,.....",
```

Tokenize

```
Array("1000","user1","2010-01-23","123.34",.....),  
Array("2002","user2","2008-02-14","349.02",.....),
```

Convert to types
& create rows

| | | | | |
|------|-------|-------|--------|--|
| 1000 | user1 | 14632 | 123.34 | |
| 2002 | user2 | 13923 | 349.02 | |

Common Issues with CSV

- Slow startup on large datasets
- Poor parallelization due to non-splittable files
- Be mindful of “wide” files
 - GC issues
 - Spark 2.4 column pruning ([SPARK-24244](#))

| GC Time | Min | 25th percentile | Median | 75th percentile | Max |
|---------------------------|-------|-----------------|--------|-----------------|-------|
| Column Pruning Off | 0.5 s | 0.5 s | 0.6 s | 0.6 s | 0.6 s |
| Column Pruning On | 9 ms | 31 ms | 61 ms | 64 ms | 70 ms |

Selecting 2 out of 200 columns
from 3.6GB compressed dataset

Schema Inference

- CSV
 - `inferSchema` causes full scan of data
 - Specify schema or use tables
 - Be precise, use the right types!
- JSON
 - Full scan of data
 - Specify schema or use tables

| | Shuffle Write |
|----------|---------------|
| strings | 217.0 MB |
| doubles | 159.0 MB |
| decimals | 145.3 MB |

Aggregate query on
2 columns & 10 million rows

Schema Inference

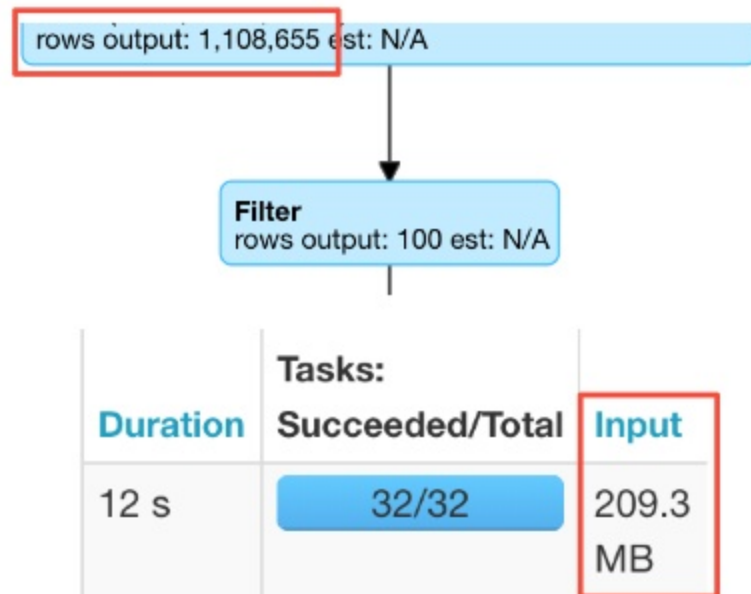
- Parquet
 - If no `_metadata`, use schema from first file
 - If `mergeSchema` read footers from all files
 - Specify schema or use tables
- ORC
 - Reads first file for schema
 - Specify schema or use tables

Optimizing Parquet Storage

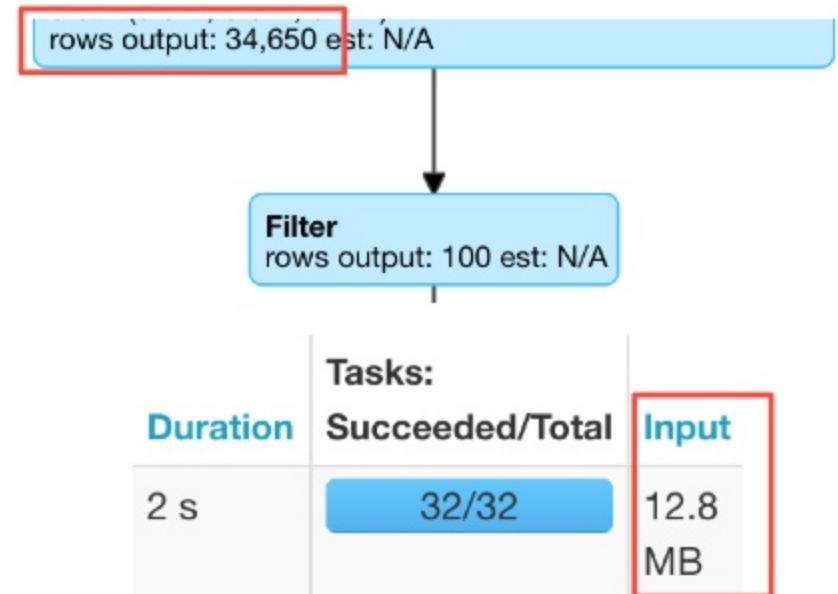
- Embedded column statistics
 - min, max, count
 - Useful for filtering
- Most effective on sorted data
 - Order by filter column when writing
 - Reduce data reads by skipping files

```
select * from mytable  
where id between 1 and 100
```

UNSORTED



SORTED



Challenges with Parquet Data Skipping

- Really most effective on single sort key
- Need to be careful with data skew
- How to maintain performance as new data appended?
- Data Skipping and ZORDER clustering
 - [Processing Petabytes of Data in Seconds with Databricks Delta](#)

Parallelization with File Formats

- Spark 2.0 default split 128MB
 - `spark.sql.files.maxPartitionBytes`
- Bin packing of smaller files
 - Auto-compaction, avoids coalesce
- Based on file-size
 - High compression ratio may result in smaller files
 - Data is decoded & uncompressed when reading

Parallelization & File Sizes

- Controlling file sizes during ETL
- Generating more rows in queries (`explode`, `flatMap`, broadcast joins)
- Solution - `repartition`?
 - Adds shuffle boundary, memory & I/O constraints
- Consider instead
 - `maxPartitionBytes` (for splittable input)
 - `maxRecordsPerFile` (for controlling output)

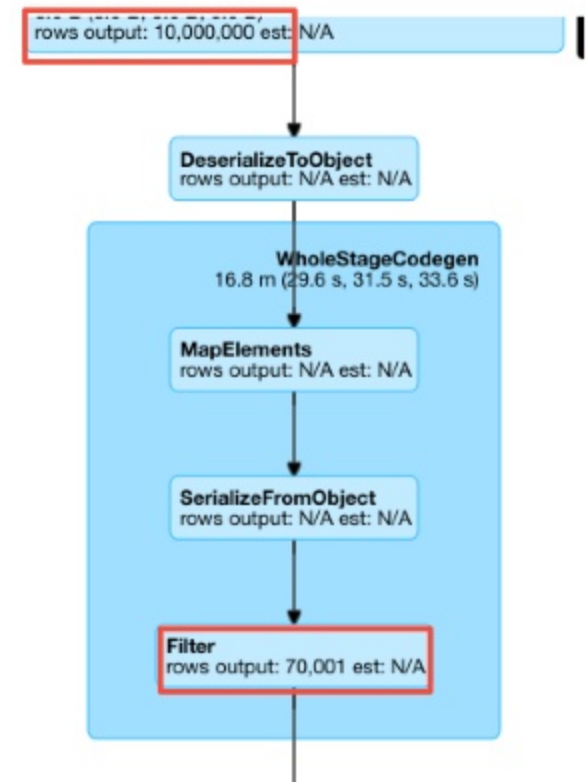
Using Spark UI & Query Plans

- SQL UI is invaluable to understanding execution
- Identify stage boundaries (shuffles)
- Use metrics for diagnosing your query
 - Are filters and column pruning applied?
 - How large are your shuffles?
 - How much memory usage and spilling?

Datasets & Order of Operations

```
case class Input(id: Long, col1: String)
case class Output(id: Long, col1: String, col2: String)
def transformToOutput(input: Input): Output = {
  ...
}

val query = basedata.as[Input]
  .map(i => transformToOutput(i))
  .filter('id <= 700000)
```

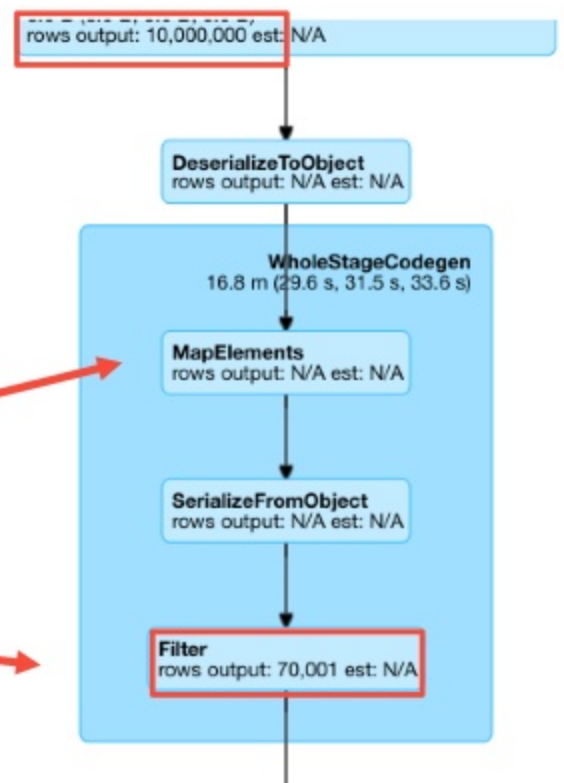


Datasets & Order of Operations

```
case class Input(id: Long, col1: String, col2: String)
case class Output(id: Long, col1: String, col2: String)
def transformToOutput(input: Input): Output = {
  ...
}
```

```
val query = basedata.as[Input]
  .map(i => transformToOutput(i))
  .filter('id <= 700000)
```

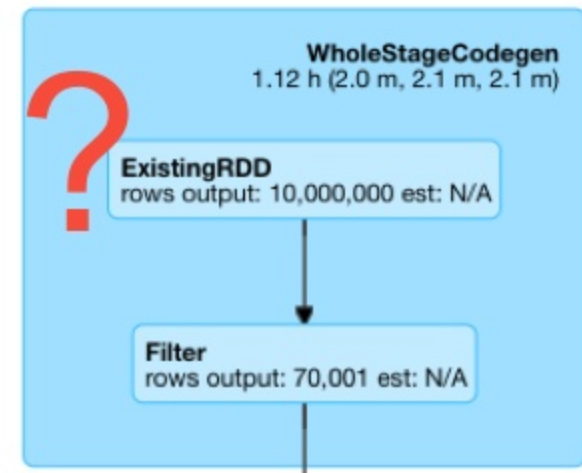
Dataset `map` breaks
lineage of "id" column,
no predicate pushdown!



Moving Between DataFrames & RDDs

```
output = df.rdd.map(myPythonFunc)

outputDF = spark.createDataFrame(
    output,
    ["id", "col1", "col2"]
).where("id <= 700000")
```

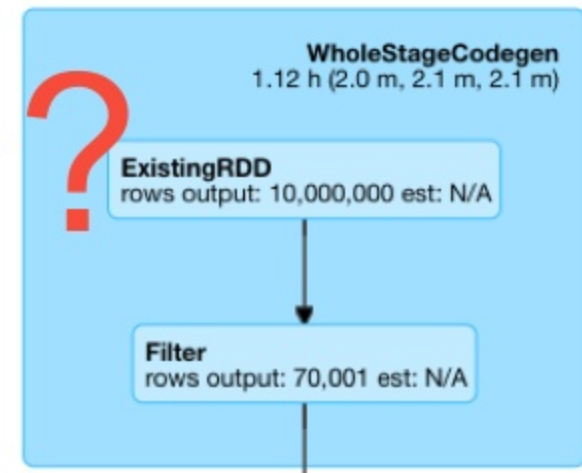


Moving Between DataFrames & RDDs

```
output = df.rdd.map(myPythonFunc)

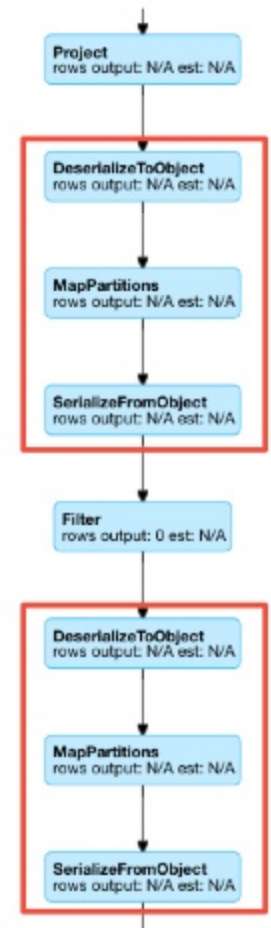
outputDF = spark.createDataFrame(
    output,
    ["id", "col1", "col2"]
).where("id <= 700000")
```

Converting to RDD breaks
DataFrame lineage, no predicate
pushdown, no column pruning, no SQL plan,
and less efficient PySpark RDD transformations...



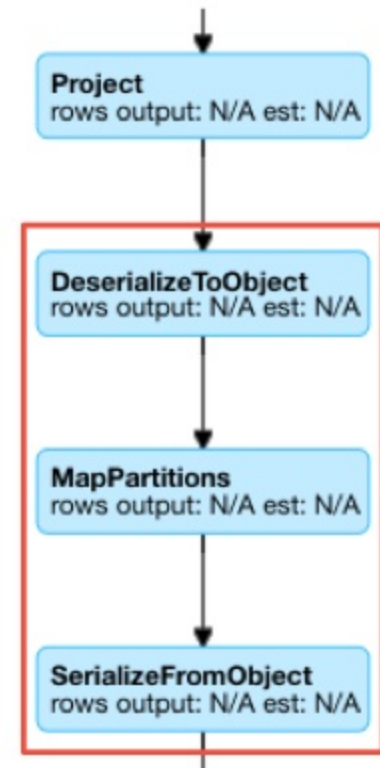
Using SparkR UDFs

```
step1 = dapply(input, firstUDF, step1Schema) # UDF  
step2 = filter(step1, step1$id2 >= 5000000) # Spark filter  
step3 = dapply(step2, secondUDF, step3Schema) # UDF
```



Using SparkR UDFs

```
bothUDFs <- function(input) {  
  step1 <- firstUDF(input)  
  step2 <- subset(step1, id2 >= 5000000)  
  step3 <- secondUDF(step2)  
  
  return(step3)  
}  
  
singleStep = dapply(input, bothUDFs, step3Schema)
```

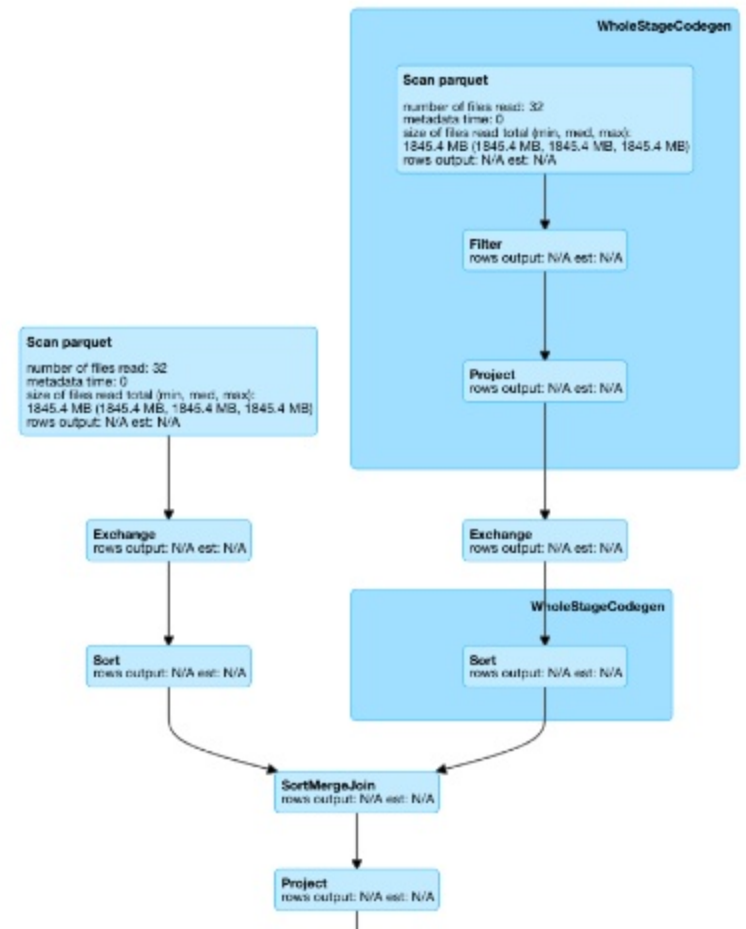


Self-Joins

```
val basedata_ex = basedata  
  .withColumn("id", ...)   
  .withColumn("mycol", ...)
```

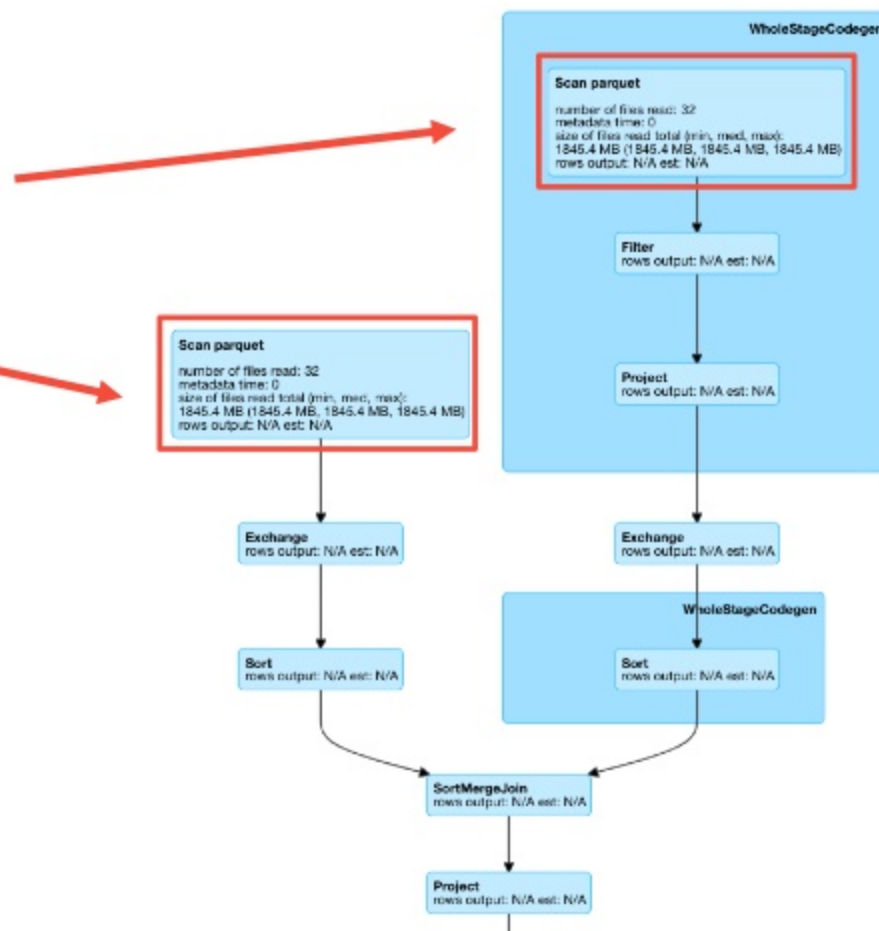
```
val joined = basedata  
  .join(basedata_ex, Seq("id"), "left")
```

- Two DataFrames referring to same data
- Typically, some transformation applied to original DF
- Join second DF to original DF (look ahead/behind, shifting rows, self reference, etc)



Self-Joins

Reading the SAME data twice!



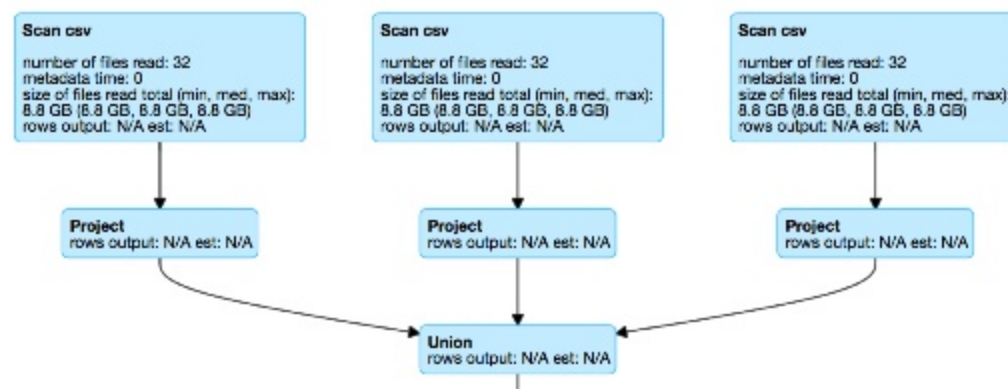
Instead consider window functions,
`flatMap`, `explode`

If absolutely necessary then
try caching, checkpointing, or
[local checkpointing](#)

Repetitive Unions

```
for config in configs:  
    temp = computedDF(input, config)  
    unioned_dfs = unioned_dfs.union(temp)
```

- Typically for different transformations to same record
- Sometimes used for Python distributed Cross Validation



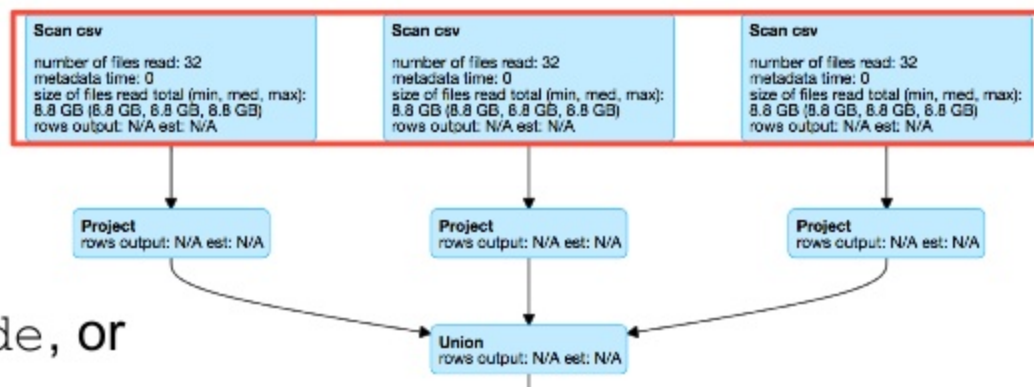
Repetitive Unions

Reading the SAME data!

Increases number of tasks
and Catalyst complexity
with each union

Instead consider `flatMap`, `explode`, or
broadcast join (dep. on data size)

If absolutely necessary then
try caching, checkpointing, or
[local checkpointing](#)



UDFs and Datasets

- PLEASE – Learn the built-in [functions](#)!
 - Prefer [Pandas UDFs](#) to PySpark UDFs
- Filter and project early
 - UDFs, lambdas are opaque to Catalyst
- Prefer DataFrames & Datasets
- SparkR – [100x Faster UDFs](#)

How Can You Learn More?

- Spend the time to understand your query plan
 - Is Spark doing what you expect?
- Understand how different transformations, UDFs, and configs impact execution
- If you're not sure...go to the source!
 - [Apache Spark GitHub](#)

QUESTIONS?