

# How to Avoid Drowning in Logs

Streaming 180 Billion Events/Day and  
Batching 150 TB/Hour

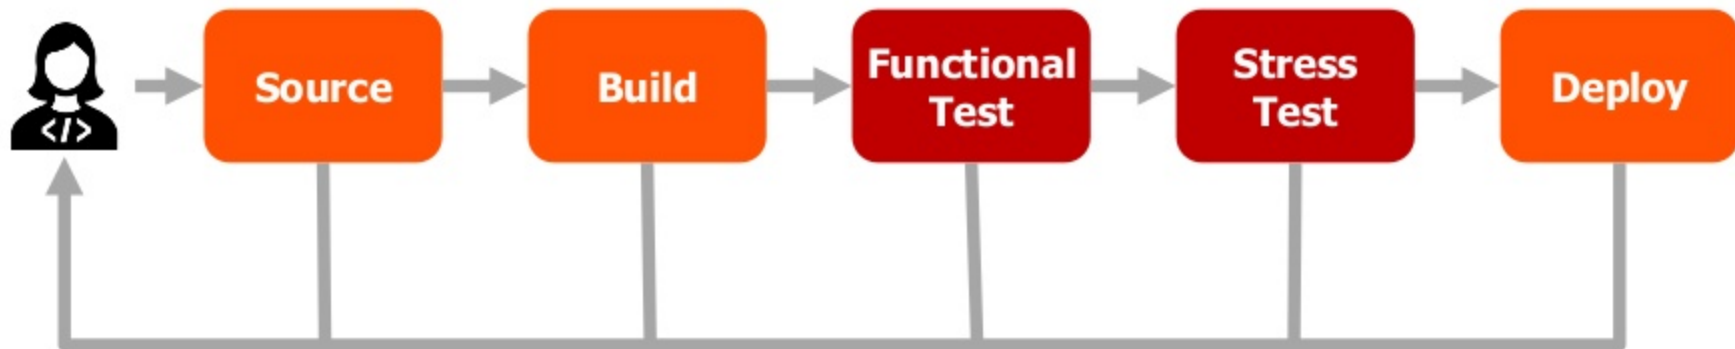
**Joshua Robinson**

Founding Engineer, FlashBlade


# Log Analytics Pipeline in Numbers

- ✓ **2M** events / second
- ✓ **5** seconds SLA
- ✓ **0.5 - 1 PB** of data / day

# Continuous Integration & Continuous Deployment



# CI/CD works!




# Scale Problems



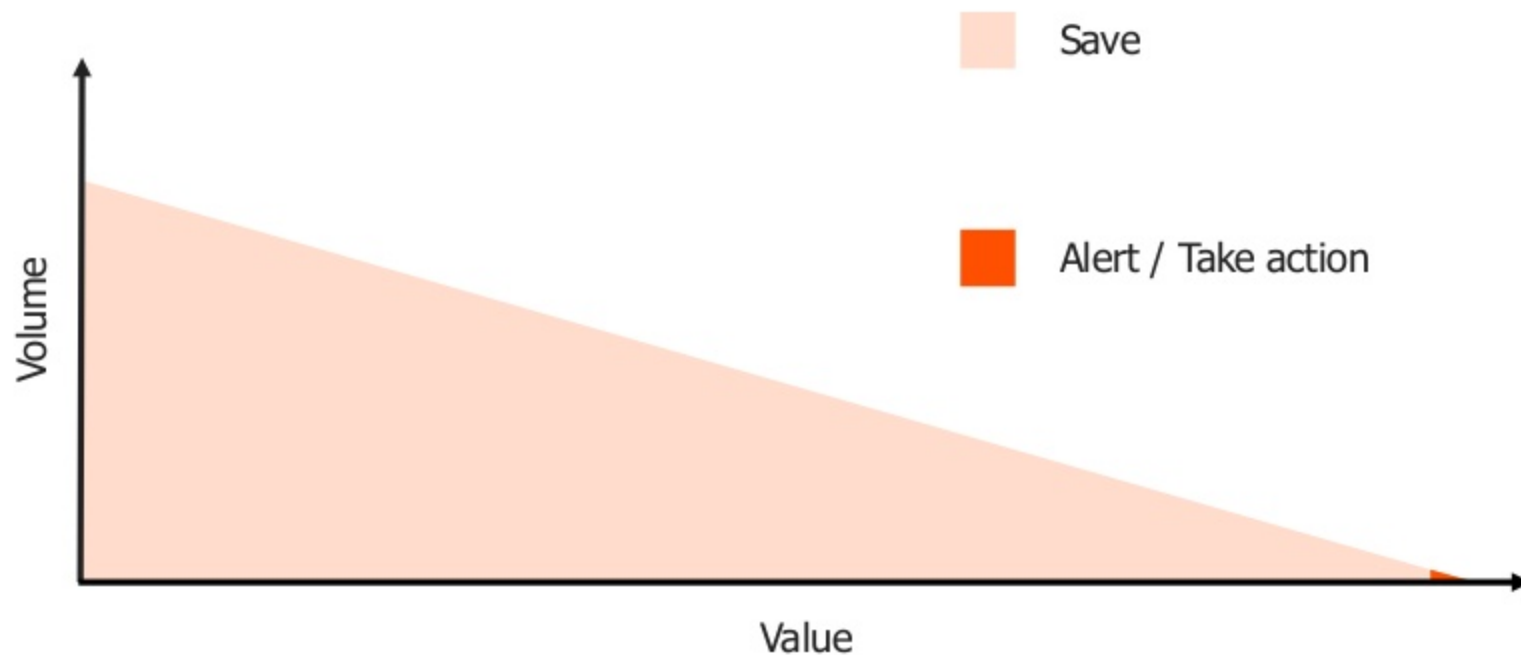
# Log Analysis Dream

1. Automate triaging of failures
2. Extract performance metrics
3. Save our logs for future use
4. Do all of this in a scalable system
5. Real-time results!

# Log Analysis

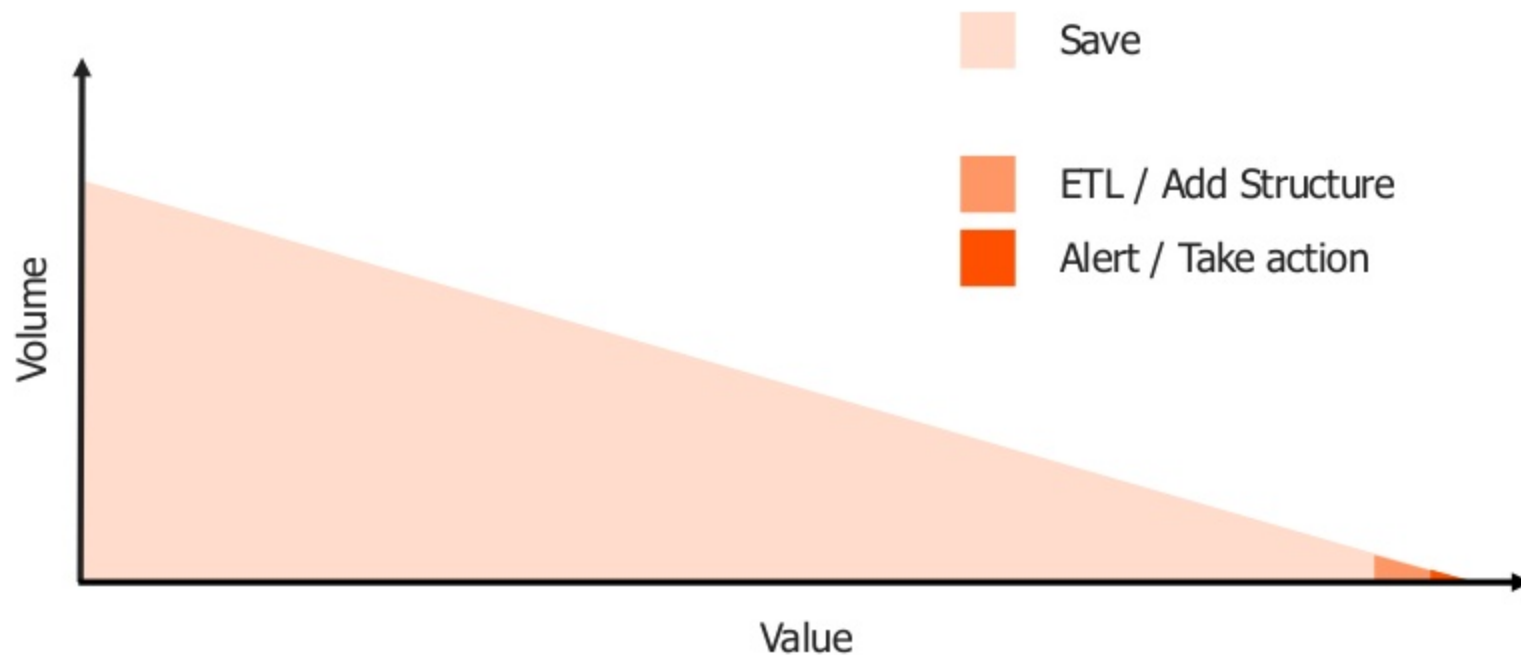


# Log Analysis v1







# Log Analysis v2



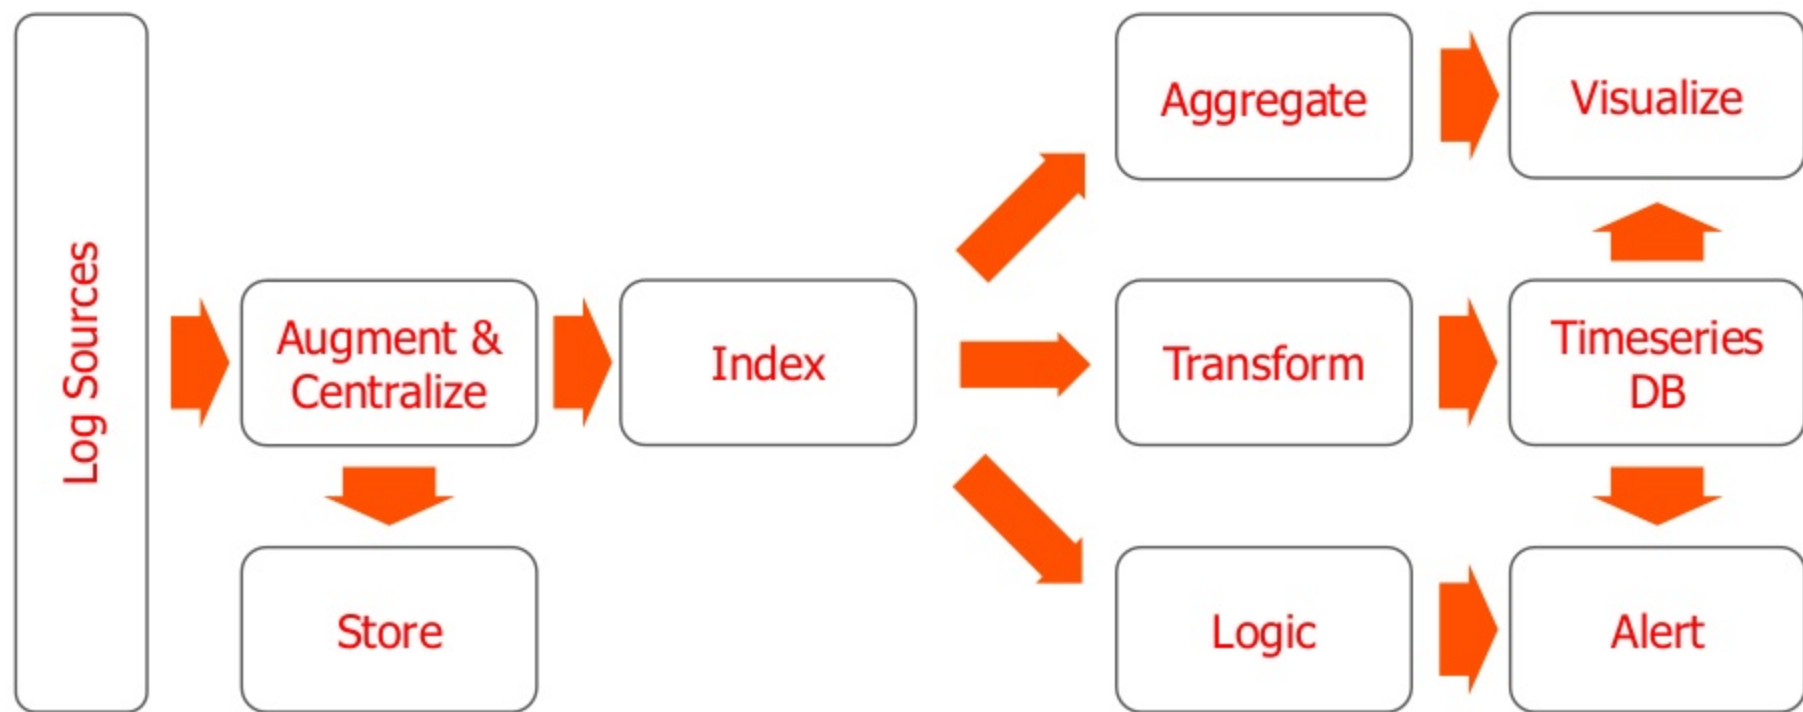
# Log Analysis v3



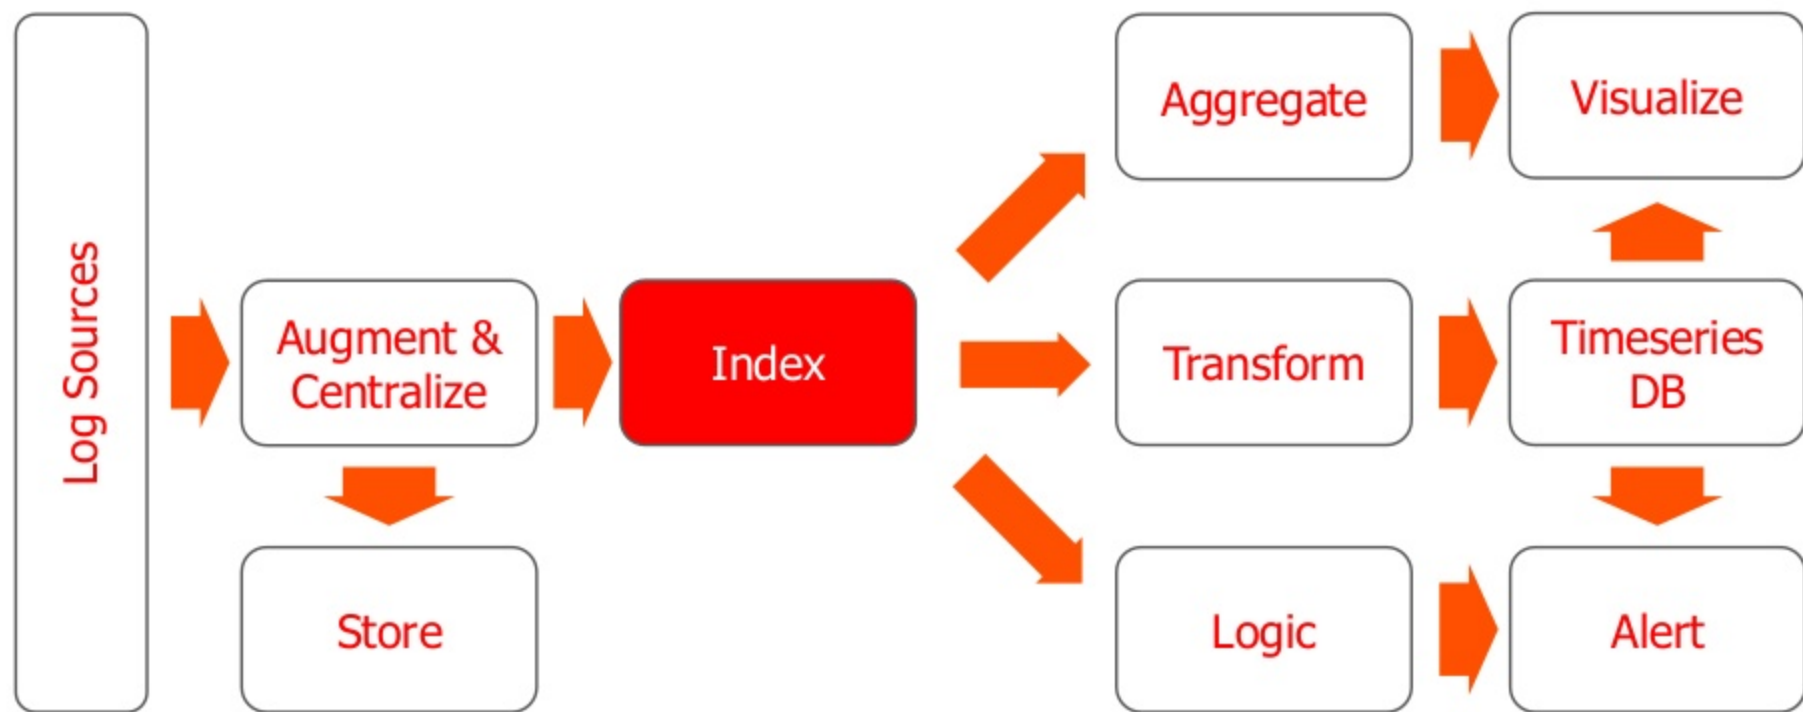
# Log Analysis v10



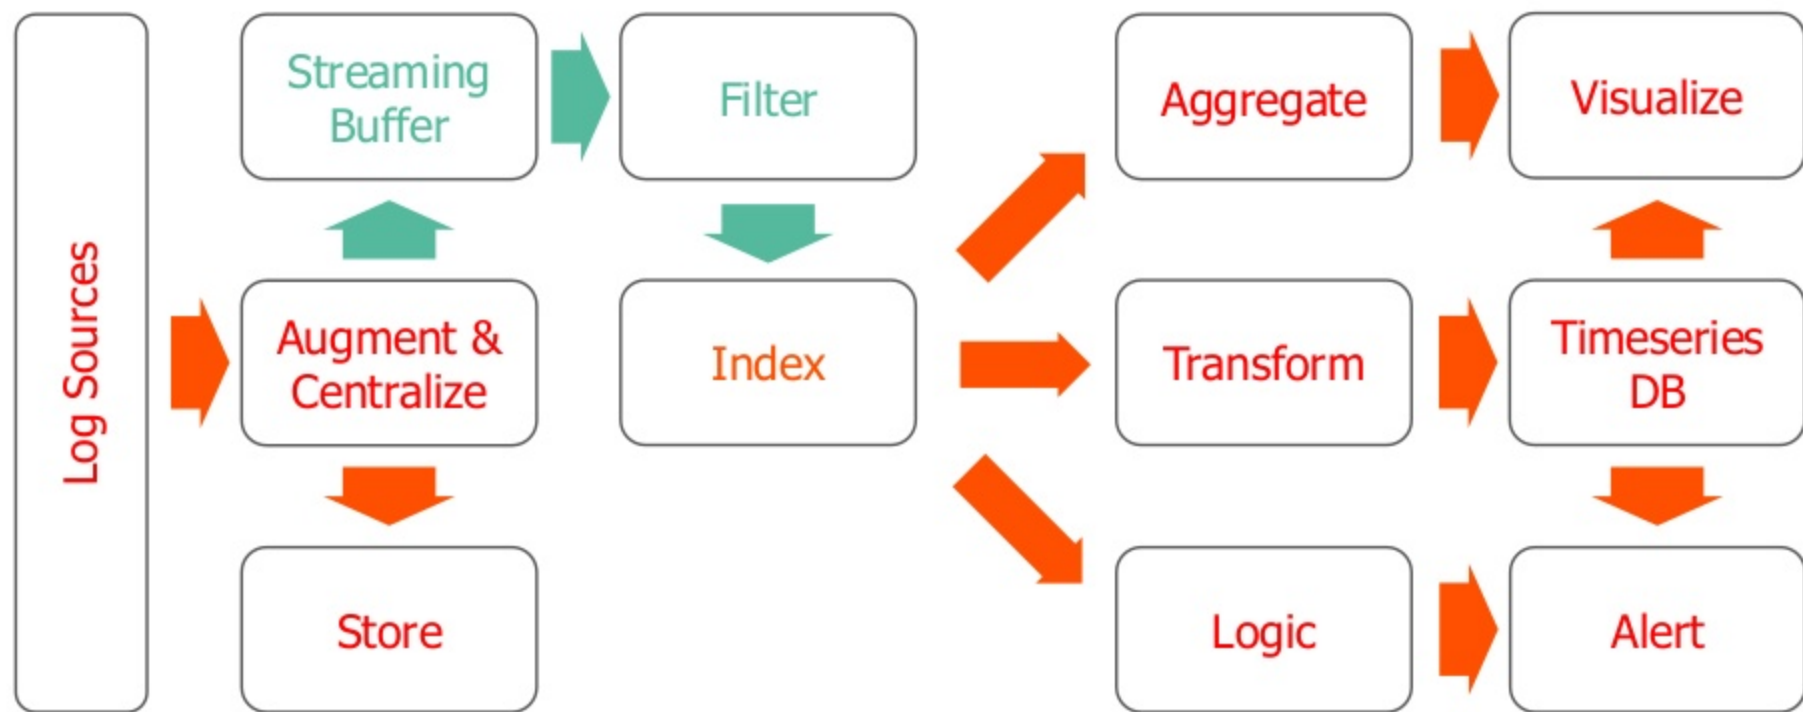
# Log Analysis Pipeline



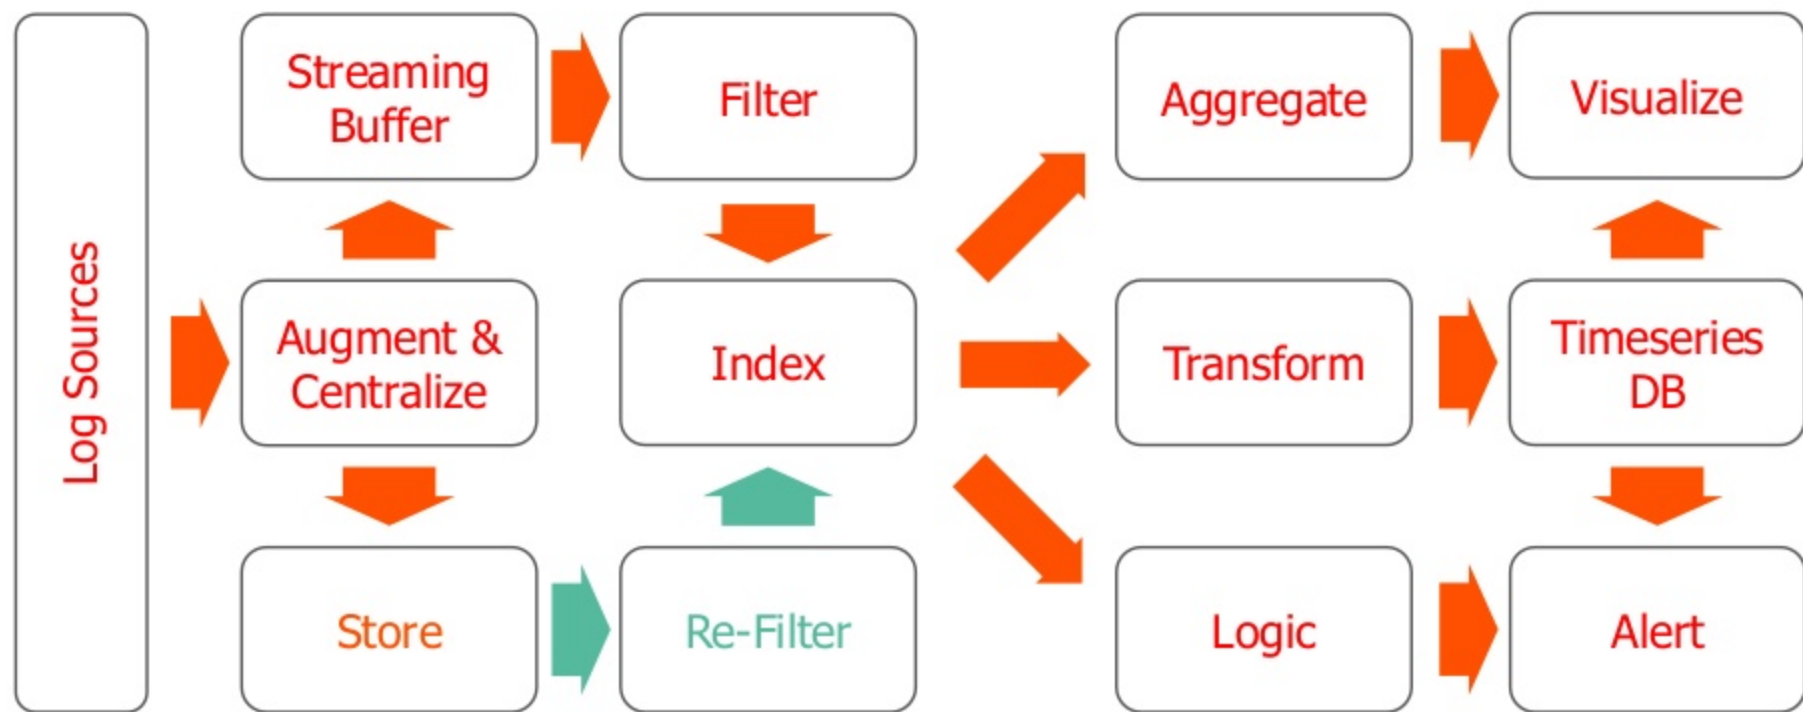
# Log Analysis Pipeline



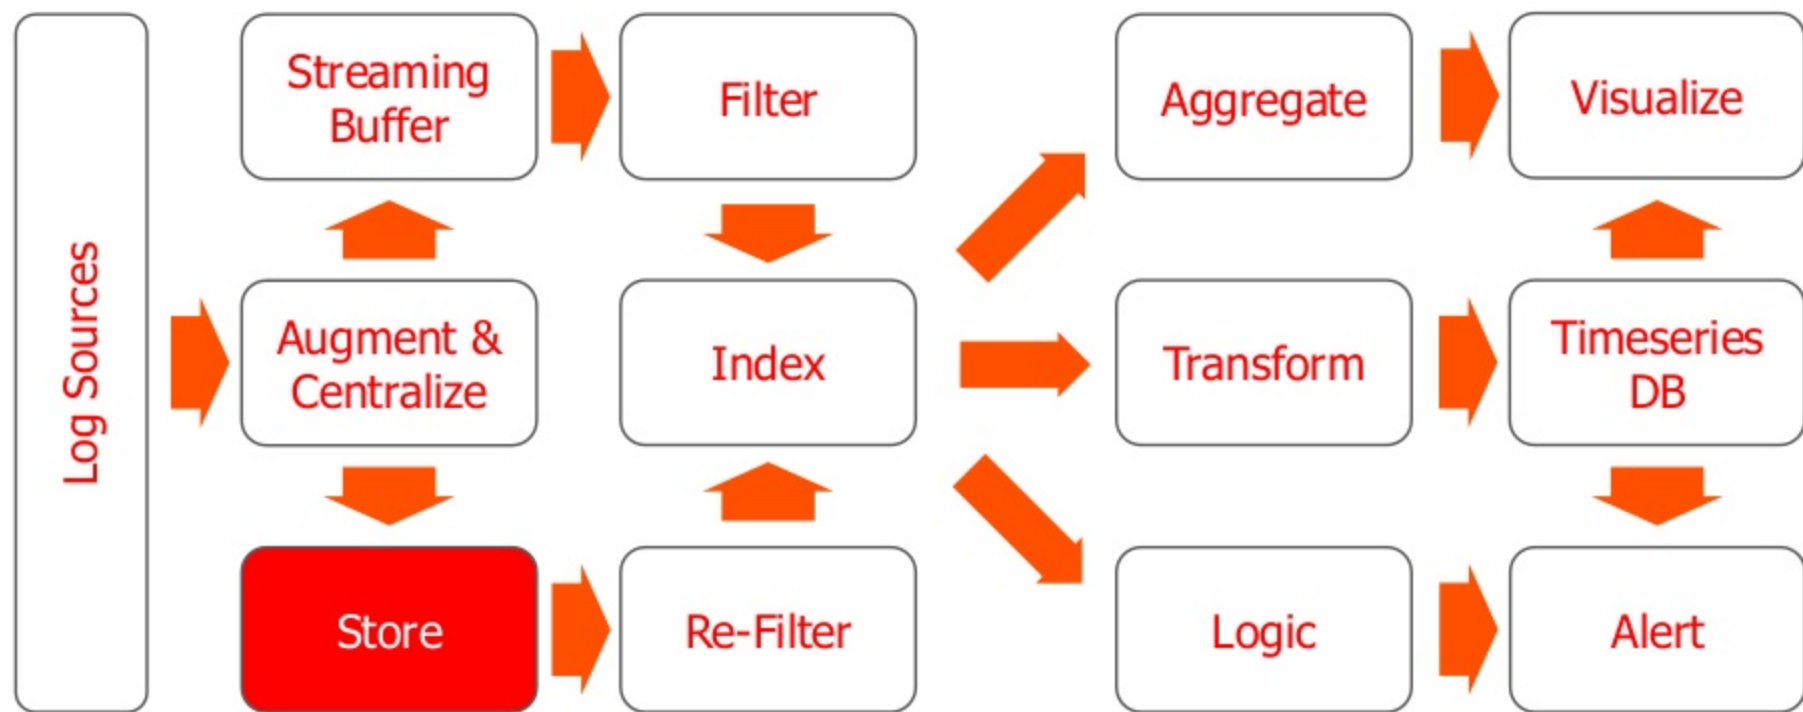
# Log Analysis Pipeline



# Log Analysis Pipeline

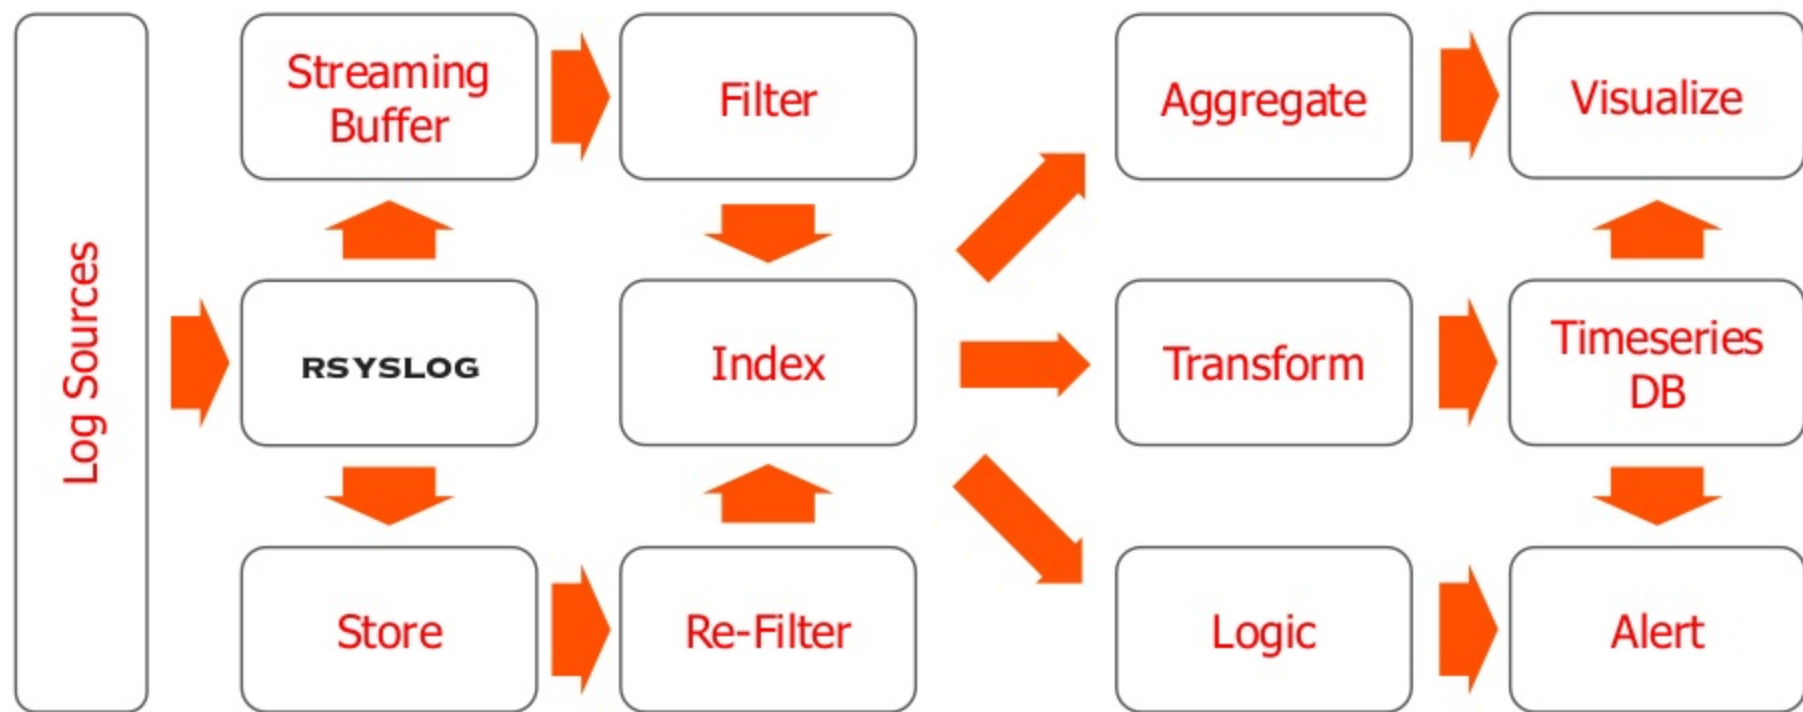


# Log Analysis Pipeline

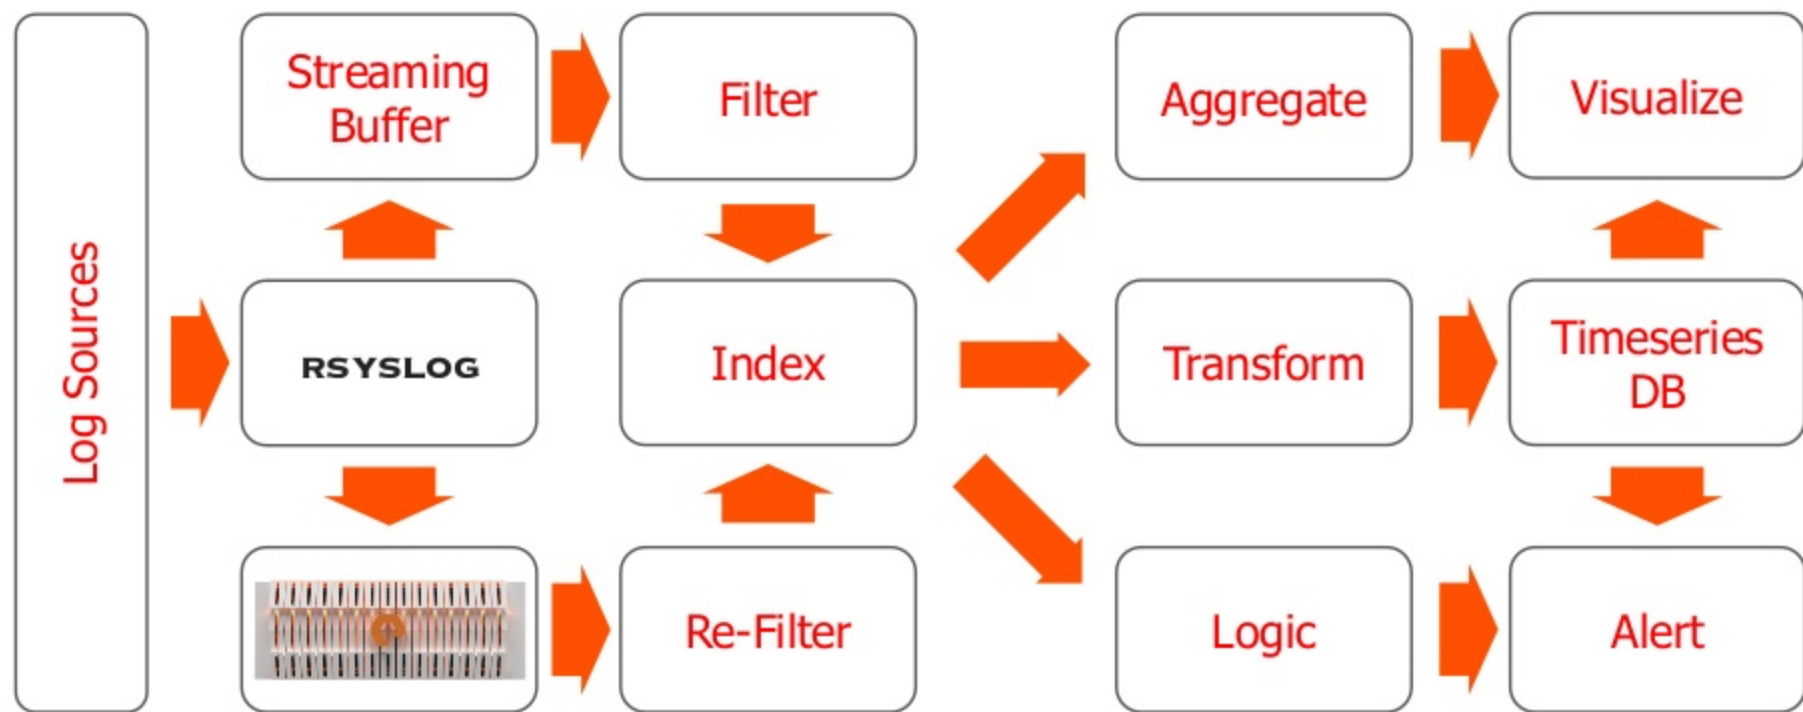




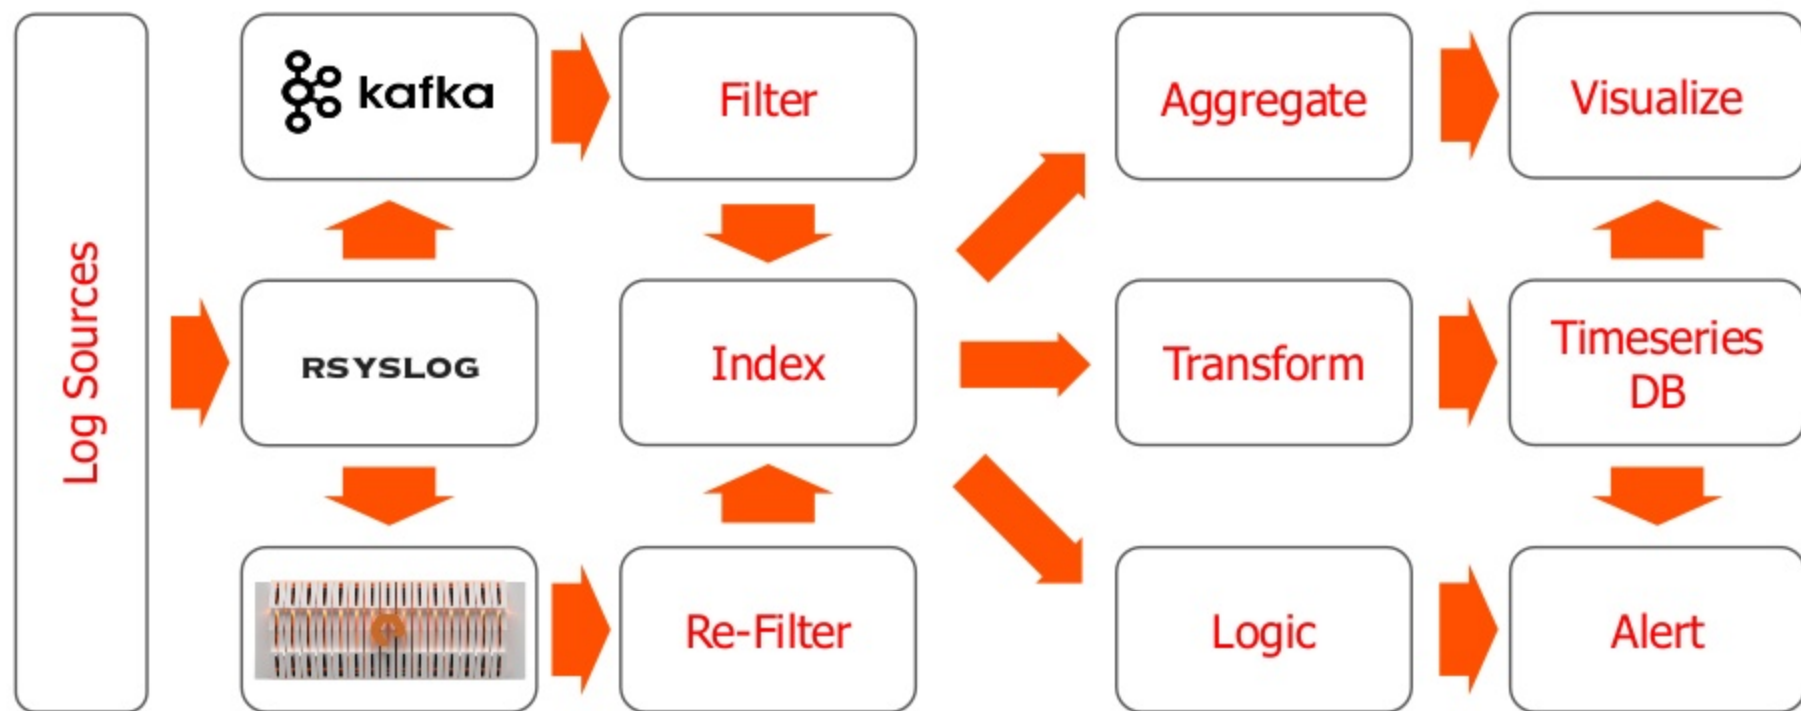
# Log Analysis Pipeline



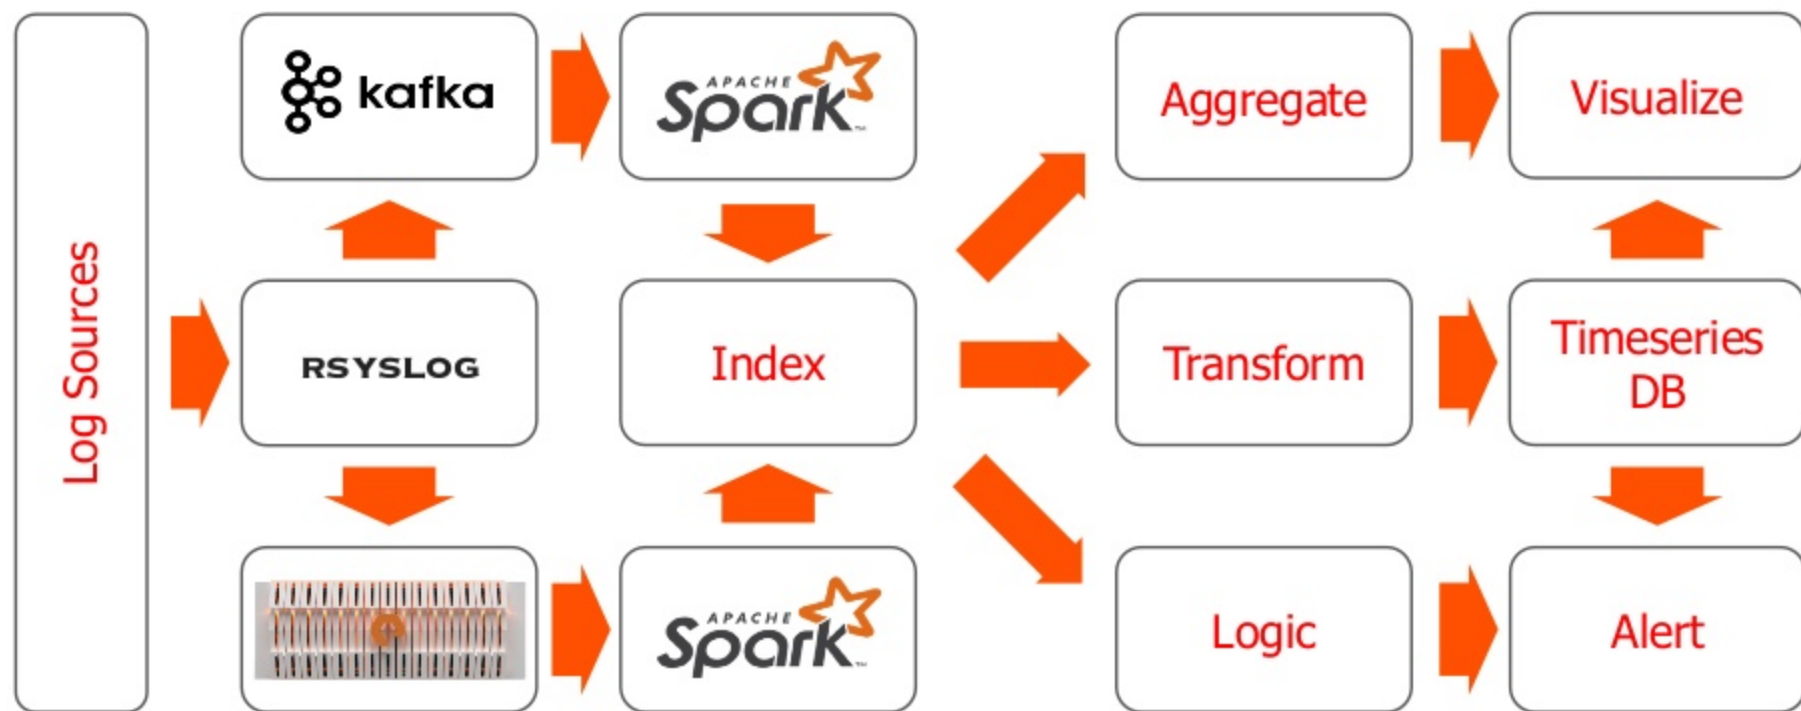
# Log Analysis Pipeline



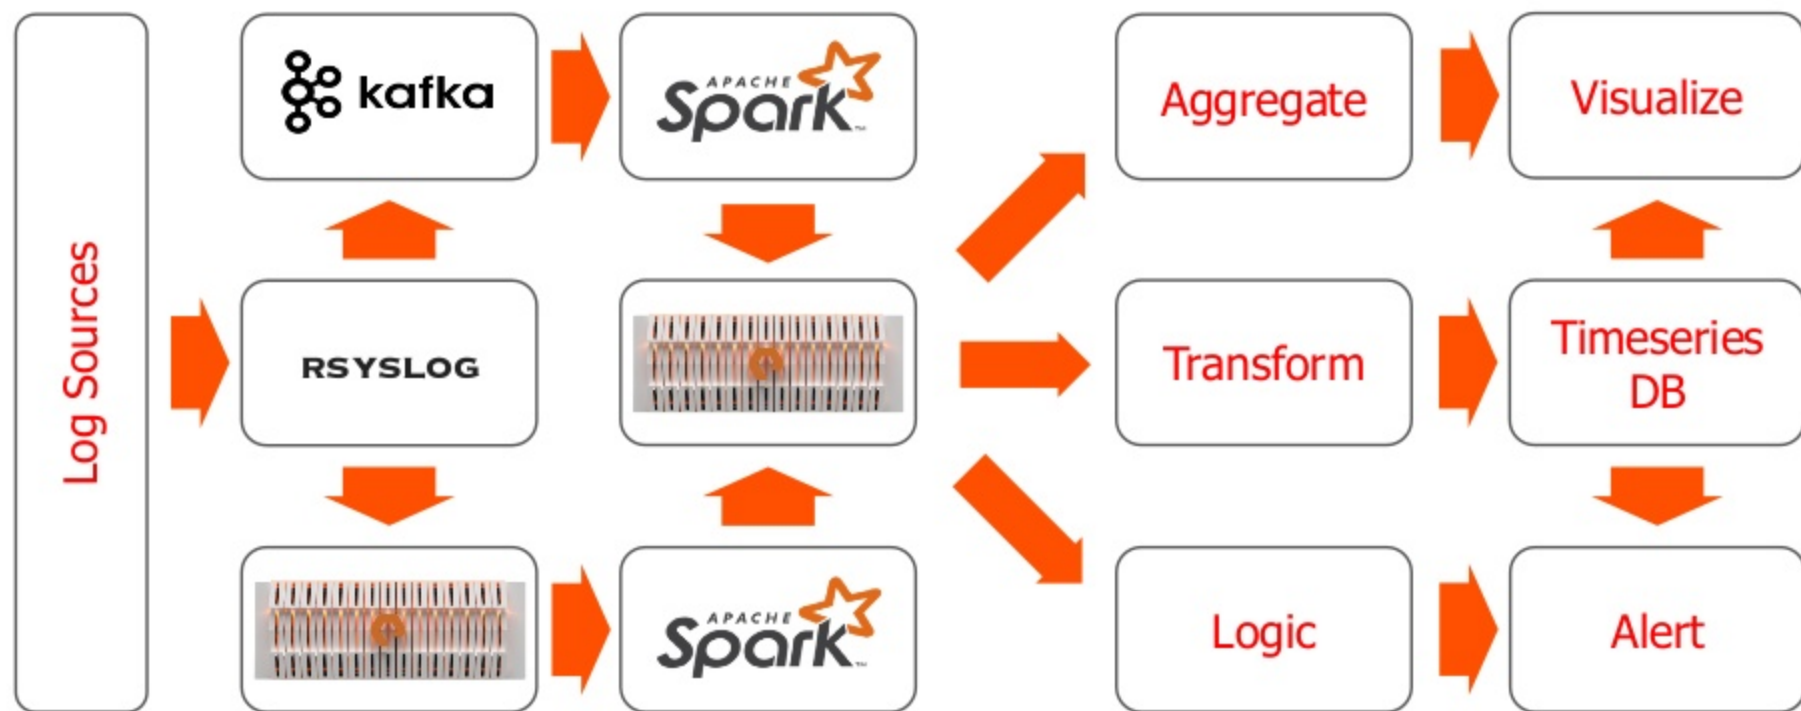
# Log Analysis Pipeline



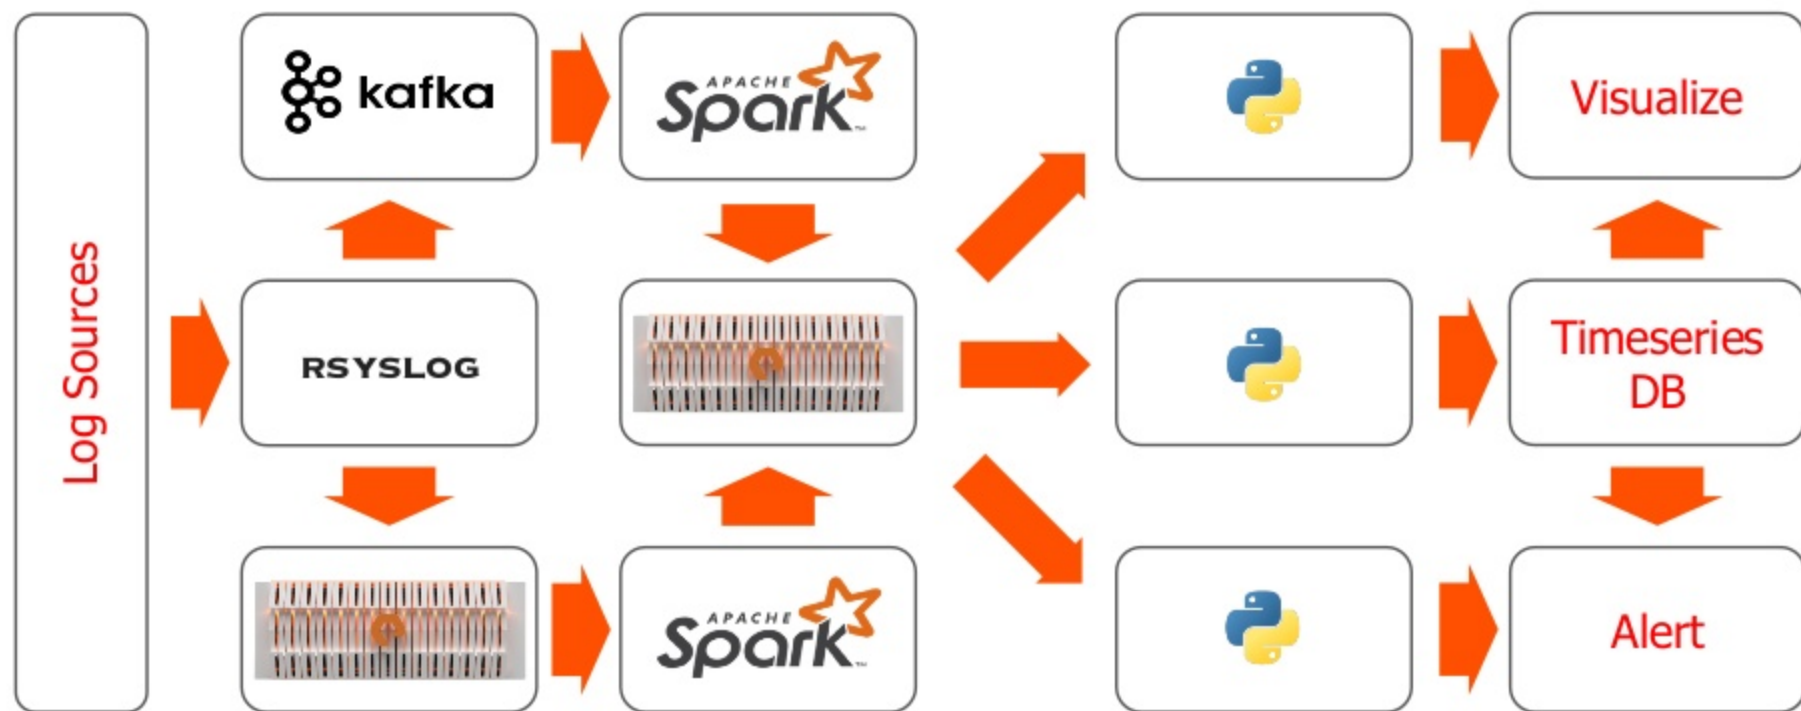
# Log Analysis Pipeline



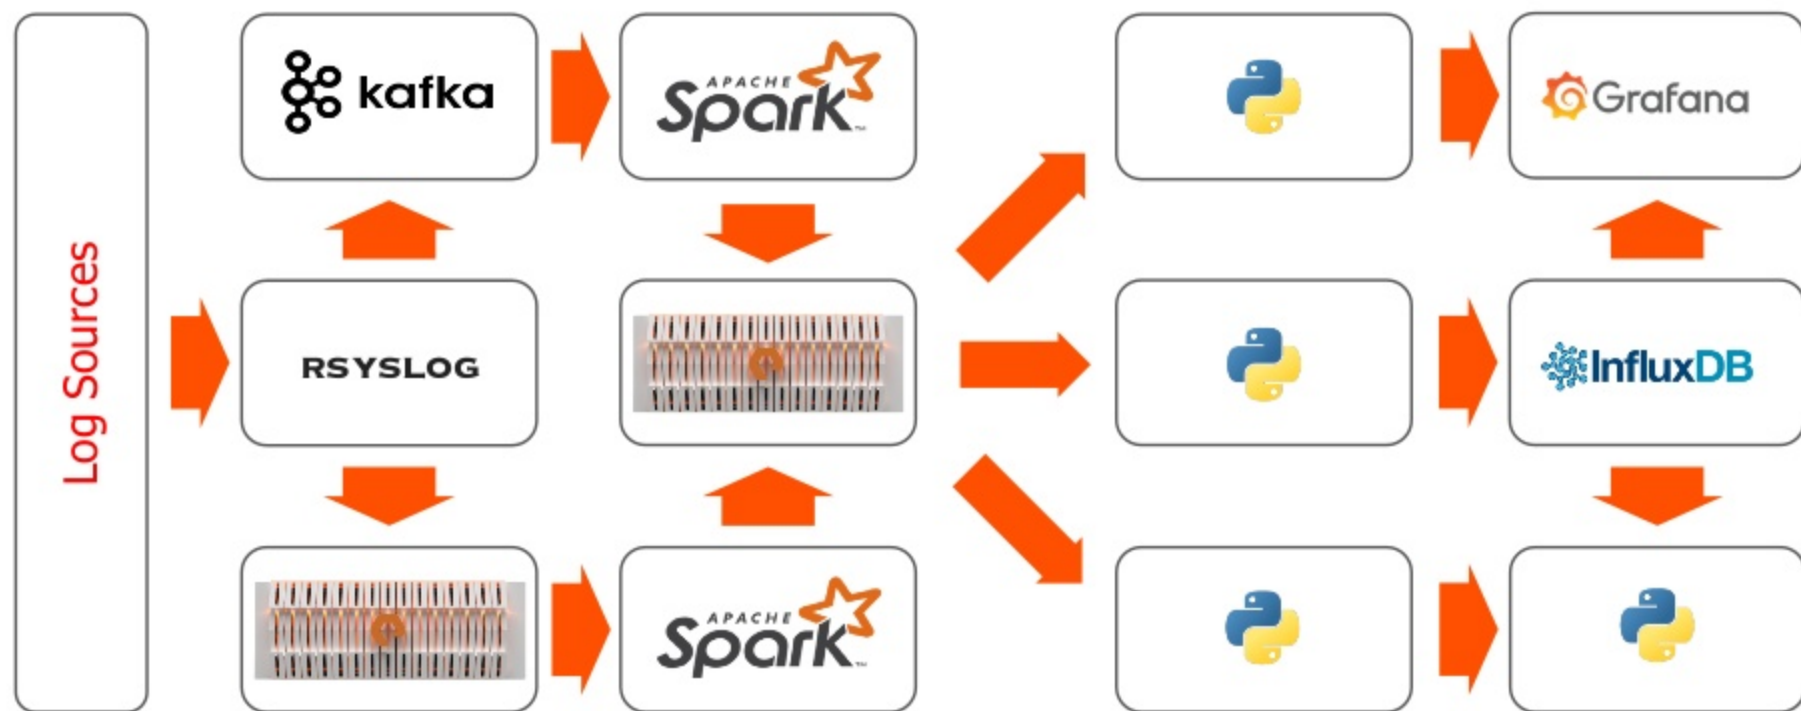
# Log Analysis Pipeline



# Log Analysis Pipeline

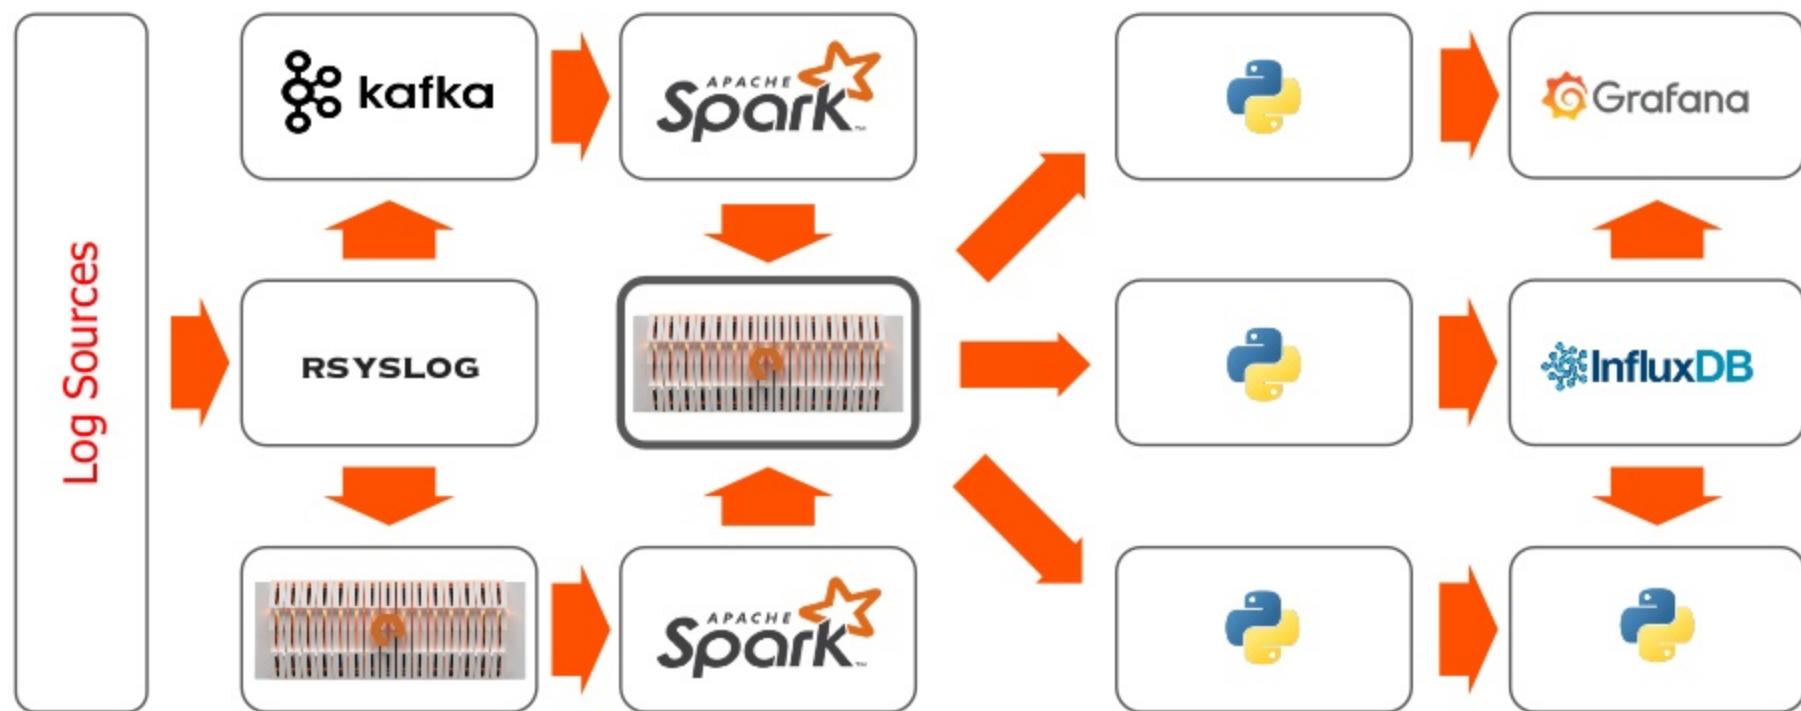


# Log Analysis Pipeline





# Log Analysis Pipeline





# Indexing

Use filesystem directory structure to encode metadata

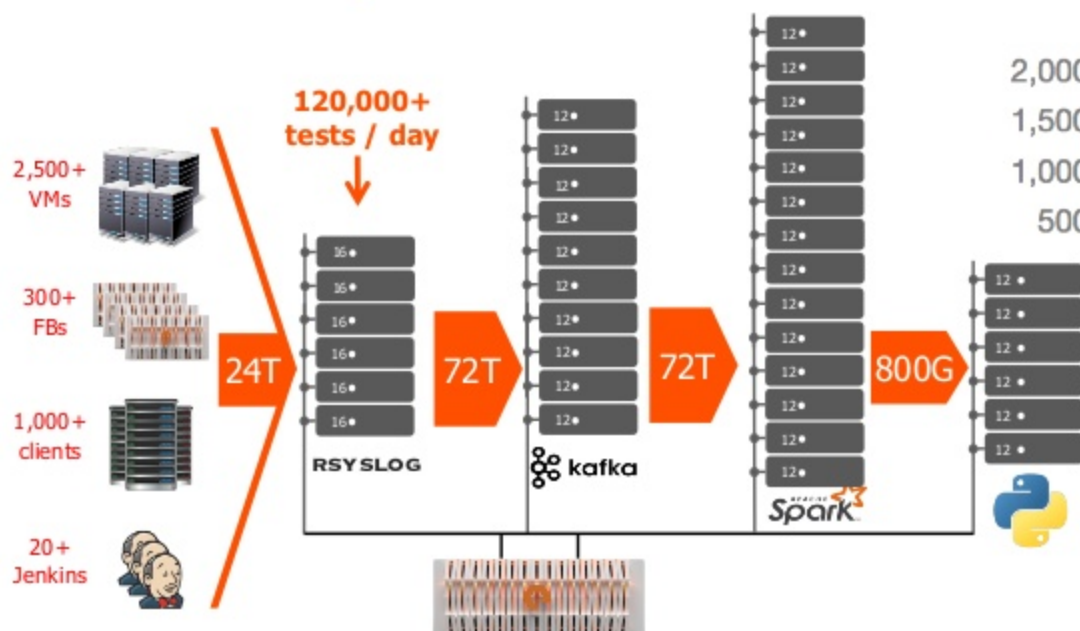
- **Raw data:** <host>/<year>/<month>/<day>/<flat files>
  - Producer: Rsyslog
  - Consumer: Spark batch (re-filter or custom lookbacks)
- **Indexed data:** <pattern>/<year>/<month>/<day>/<hour>/<host>/<flat files>
  - Producer: Spark streaming (filter)
  - Consumer: Python services (e.g. ETL, alert, searchability)

# Querying

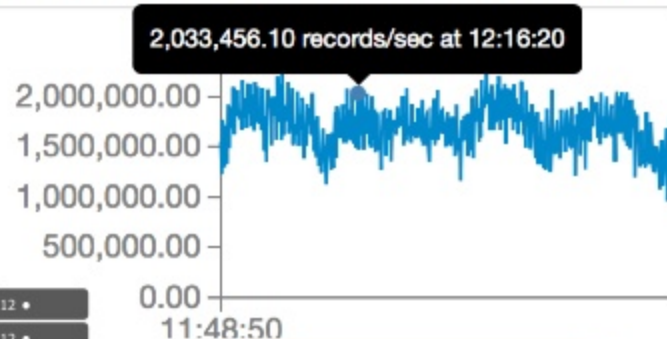
Find and load data

- FlashBlade NFS protocol. < 1ms latency
- **Listing**
  - "ls -aR" is still SLOW
  - NFS client in kernel sequentially discovers filesystem structure.
  - Solution: Skip the kernel. Use libnfs to create our own parallelized discovery. 1000x faster for 1M files
- **Reading**
  - Buffering: Create input pipeline to optimize for throughput and hide latency away

# Full Pipeline



Timelines (Last 1000 batches, 0 active, 1000 co



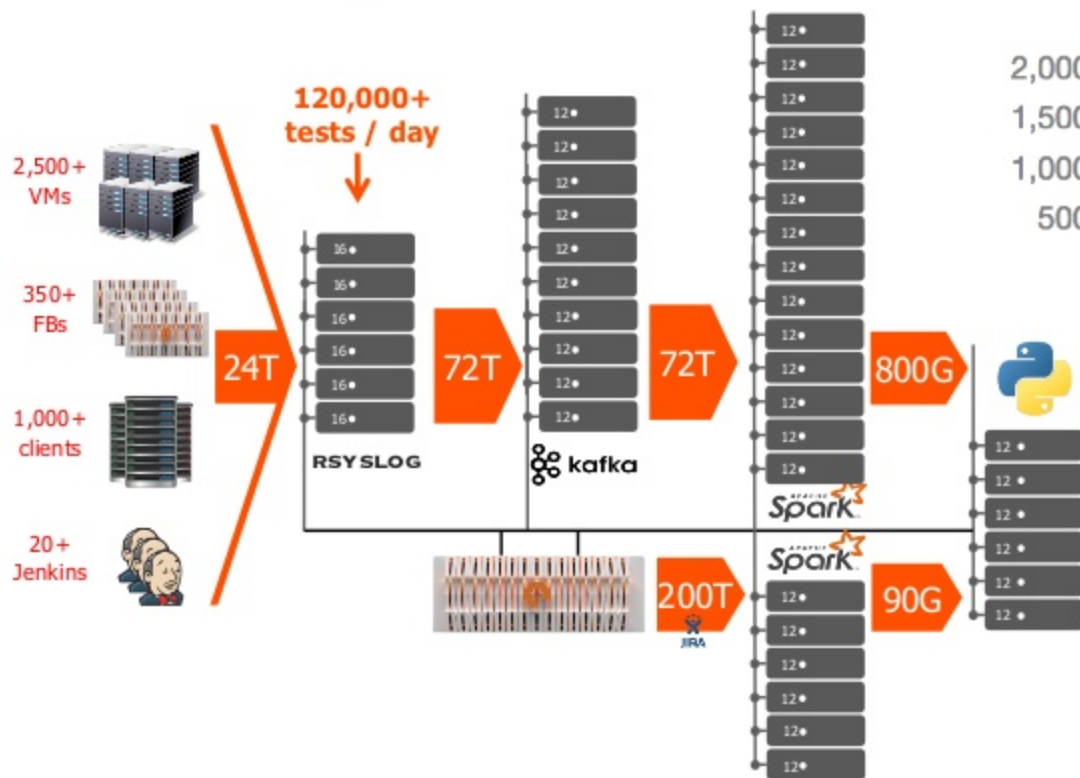
JIRA




Grafana

- ✓ Duplicate bug
- ✓ Infrastructure failure
- ✓ Performance regression

# Full Pipeline



Timelines (Last 1000 batches, 0 active, 1000 co



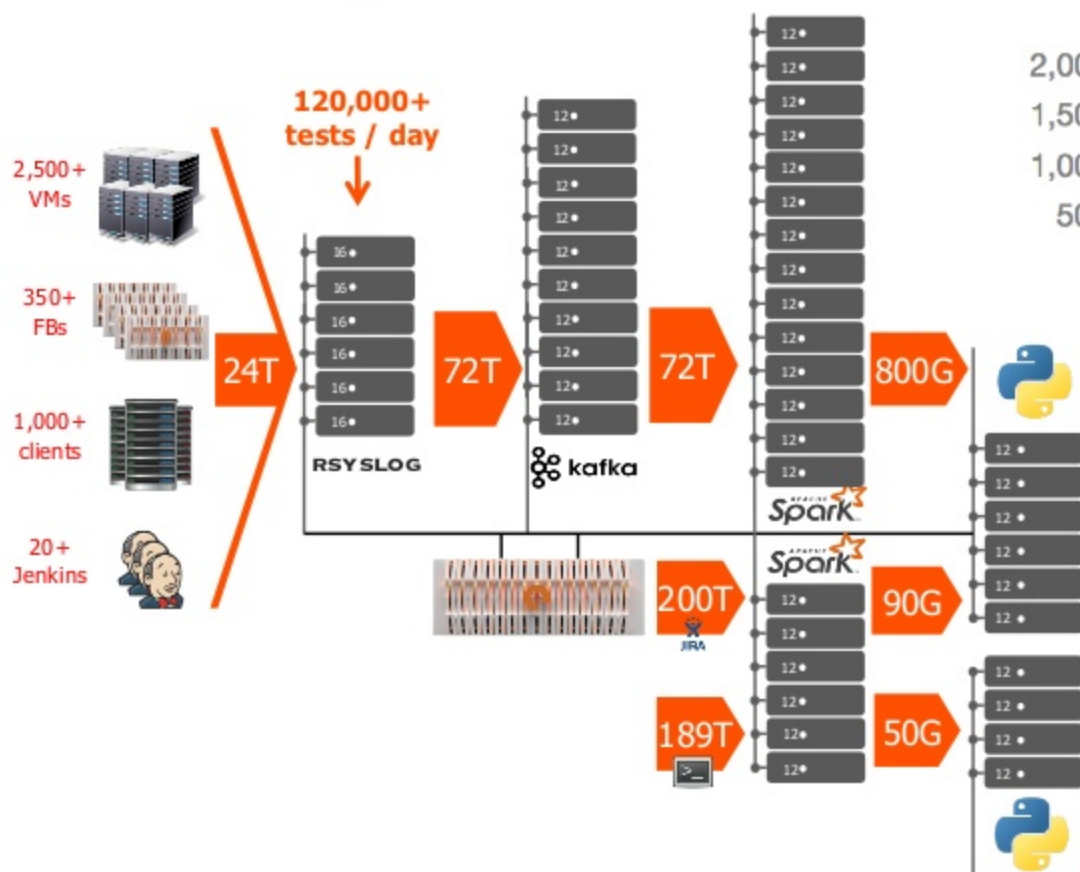
JIRA




Grafana

- ✓ Duplicate bug
- ✓ Infrastructure failure
- ✓ Performance regression

# Full Pipeline



Timelines (Last 1000 batches, 0 active, 1000 c)



JIRA



Grafana

- ✓ Duplicate bug
- ✓ Infrastructure failure
- ✓ Performance regression

InfluxDB



Grafana

- ✓ Low level details
- ✓ Easy to read graphs

# Takeaways

- ✓ **Index only what you need, store the rest**  
(in a storage layer that scales in throughput and to billions of files/objects)
- ✓ **Optimize for throughput and not latency**
- ✓ **Disaggregation of compute and storage for scalability of subsystems**





**QUESTIONS?**